

Statistical Methods in Natural Language Processing

3. Statistical language modeling

Pavel Pecina, Jindřich Helcl

21 October, 2025

Course Segments

1. Introduction, probability, essential information theory
2. Statistical language modelling (n-gram)
3. Statistical properties of words
4. Word embeddings
5. Hidden Markov models, Tagging

Recap from Last Week

Entropy

- Let $p_X(x)$ be a distribution of random variable X
- Basic outcomes (alphabet) Ω

$$H(X) = - \sum_{x \in \Omega} p(x) \log_2 p(x)$$

- Unit: bits (\log_e : nats)
- Notation: $H(X) = H_p(X) = H(p) = H_X(p) = H(p_X)$

Joint Entropy and Conditional Entropy

- Two random variables: X (space Ω), Y (Ψ)
- Joint entropy:
 - no big deal: (X, Y) considered a single event:

$$H(X, Y) = - \sum_{x \in \Omega} \sum_{y \in \Psi} p(x, y) \log_2 p(x, y)$$

- Conditional entropy:
$$H(Y|X) = - \sum_{x \in \Omega} \sum_{y \in \Psi} \underline{p(x, y)} \log_2 p(y|x)$$
 - recall that $H(X) = E(\log_2(1/p_X(x)))$
 - (weighted “average”, and weights are not conditional)

Kullback-Leibler Distance (Relative Entropy)

- Remember:
 - long series of experiments ... c_i/T_i oscillates around some number... we can only estimate it ... to get a distribution q .
- So we get a distribution q ; (sample space Ω , r.v. X)
- The true distribution is, however, p (same Ω , X)
⇒ how big error are we making?
- $D(p||q)$ (the Kullback-Leibler distance):

$$D(p||q) = \sum_{x \in \Omega} p(x) \log_2 (p(x)/q(x)) = E_p \log_2 (p(x)/q(x))$$

Mutual Information

- Rewrite the definition:
 - recall: $D(r||s) = \sum_{v \in \Omega} r(v) \log_2 (r(v)/s(v));$
 - substitute: $r(v) = p(x,y), s(v) = p(x)p(y); \langle v \rangle \sim \langle x, y \rangle$

$$\begin{aligned} I(X,Y) &= D(p(x,y) || p(x)p(y)) = \\ &= \sum_{x \in \Omega} \sum_{y \in \Psi} p(x,y) \log_2 (p(x,y)/p(x)p(y)) \end{aligned}$$

- Measured in bits (what else? :-)

Information Inequality

$$D(p||q) \geq 0$$

Proof:

$$0 = -\log 1 = -\log \sum_{x \in \Omega} q(x) = -\log \sum_{x \in \Omega} p(x)(q(x)/p(x)) \leq$$

...apply Jensen's inequality here (-log is convex)...

$$\leq \sum_{x \in \Omega} p(x)(-\log(q(x)/p(x))) = \sum_{x \in \Omega} p(x)\log(p(x)/q(x)) = D(p||q)$$

Cross Entropy: The Formula

- $H_{p'}(\tilde{p}) = H(p') + D(p' \parallel \tilde{p})$

$$H_{p'}(\tilde{p}) = - \sum_{x \in \Omega} p'(x) \log_2 \tilde{p}(x)$$

- p' is certainly not the true p , but we can consider it the “real world” distribution against which we test
- note on notation (confusing): $p/p' \leftrightarrow \tilde{p}$, also $H_T(p)$
- (Cross) Perplexity:

$$G_{p'}(\tilde{p}) = G_T(\tilde{p}) = 2^{H_{p'}(\tilde{p})}$$

Sample Space vs. Data

- In practice, it is often inconvenient to sum over the sample space(s) Ψ, Ω (especially for cross entropy!)
- Use the following formula:

$$H_{p'}(p) = - \sum_{y \in \Psi, x \in \Omega} p'(y, x) \log_2 p(y|x) = \\ - 1/|T'| \sum_{i=1 \dots |T'|} \log_2 p(y_i|x_i)$$


- This is the normalized log probability of the “test” data:

$$H_{p'}(p) = - 1/|T'| \log_2 \prod_{i=1 \dots |T'|} p(y_i|x_i)$$

Noisy Channel Model


Noisy Channel

Prototypical case:



Noisy Channel


Prototypical case:



- Model: probability of error (noise)
- Example: $p(0|1) = 0.25$, $p(1|1) = 0.75$, $p(1|0) = 0.5$, $p(0|0) = 0.5$
- The Task:
 - known: the noisy output; want to know: the input (decoding)

Noisy Channel

Prototypical case:



- Model: probability of error (noise)
- Example: $p(0|1) = 0.25$, $p(1|1) = 0.75$, $p(1|0) = 0.5$, $p(0|0) = 0.5$
- The Task:
 - known: the noisy output; want to know: the input (decoding)

Noisy Channel Applications

- Optical Character Recognition (OCR)
 - text → *print (adds noise)*, scan → image
- Handwriting recognition (HR)
 - text → *neurons, muscles ("noise")*, scan/digitize → image
- Speech recognition (ASR)
 - text → *conversion to acoustic signal ("noise")* → acoustic waves
- Machine Translation (MT)
 - text in target language → *translation ("noise")* → text in source language
- Also: Part of Speech Tagging
 - sequence of tags → *selection of word forms* → text

Noisy Channel: The Golden Rule of ... OCR, HR, ASR, MT, ...

- Recall:

$$A_{\text{best}} = \operatorname{argmax}_A p(A|B)$$

$$A_{\text{best}} = \operatorname{argmax}_A p(B|A) \ p(A) / \cancel{p(B)}$$

(Bayes Formula)

$$A_{\text{best}} = \operatorname{argmax}_A p(B|A) \ p(A)$$

(The Golden Rule)

- Where:

- $p(B|A)$: the acoustic/image/translation/lexical model
 - application-specific name
 - will explore later
- $p(A)$: **the language model**

The Perfect Language Model

- Sequence of word forms
- Notation: $A \sim W = (w_1, w_2, w_3, \dots, w_d)$
- The big (modeling) question:

$$p(W) = ?$$

- Well, we know (Bayes/chain rule \rightarrow):

$$p(W) = p(w_1, w_2, w_3, \dots, w_d) =$$


$$= p(w_1) \times p(w_2 | w_1) \times p(w_3 | w_1, w_2) \times \dots \times p(w_d | w_1, w_2, \dots, w_{d-1})$$

- Not practical (even short $W \rightarrow$ too many parameters)

Markov Chain

- Unlimited memory (cf. previous slide):
 - for w_i , we know all its predecessors $w_1, w_2, w_3, \dots, w_{i-1}$
- Limited memory:
 - we disregard “too old” predecessors
 - remember only k previous words: $w_{i-k}, w_{i-k+1}, \dots, w_{i-1}$
 - called “ k^{th} order Markov approximation”
- + stationary character (no change over time) → Markov Chain:

$$p(W) \cong \prod_{i=1..d} p(w_i | w_{i-k}, w_{i-k+1}, \dots, w_{i-1}), d = |W|$$



n-gram Language Models

n-gram Language Model

- $(n-1)^{\text{th}}$ order Markov approximation $\rightarrow n\text{-gram Language Model}:$

$$p(W) =^{\text{df}} \prod_{i=1..d} p(w_i | w_{i-n+1}, w_{i-n+2}, \dots, w_{i-1})$$

prediction history



n-gram Language Model

- $(n-1)^{\text{th}}$ order Markov approximation $\rightarrow n\text{-gram Language Model}:$

$$p(W) =^{\text{df}} \prod_{i=1..d} p(w_i | w_{i-n+1}, w_{i-n+2}, \dots, w_{i-1})$$

prediction history

- In particular (assume vocabulary $|V| = 60k$):
 - 0-gram LM: uniform model, $p(w) = 1/|V|$, 1 parameter
 - 1-gram LM: unigram model, $p(w)$, 6×10^4 parameters
 - 2-gram LM: bigram model, $p(w_i | w_{i-1})$, 3.6×10^9 parameters
 - 3-gram LM: trigram model, $p(w_i | w_{i-2}, w_{i-1})$, 2.16×10^{14} parameters

LM: Observations

- How large n ?
 - Nothing is enough (theoretically)
 - But anyway: as much as possible (\rightarrow close to “perfect” model)
 - Neural models allow context/history of thousands of words
 - n -gram models: 3–7
 - parameter estimation? (*reliability, data availability, storage space, ...*)
 - 4 is too much: $|V|=60k \rightarrow 1.296 \times 10^{19}$ parameters
 - but: 6–7 would be (almost) ideal (having enough data)
 - in fact: one can recover the original text sequence from 7-grams!
- Reliability $\sim (1 / \text{Detail})$ (\rightarrow need compromise)

The Length Issue

- $\forall n: \sum_{w \in \Omega^n} p(w) = 1 \Rightarrow \sum_{n=1..\infty} \sum_{w \in \Omega^n} p(w) \gg 1 (\rightarrow \infty)$
- We want to model all sequences of words
 - for “fixed” length tasks: no problem - n fixed, sum is 1
 - *tagging, OCR/handwriting (if words identified ahead of time)*
 - for “variable” length tasks: have to account for
 - *discount shorter sentences*

The Length Issue

- $\forall n: \sum_{w \in \Omega^n} p(w) = 1 \Rightarrow \sum_{n=1.. \infty} \sum_{w \in \Omega^n} p(w) >> 1 (\rightarrow \infty)$
- We want to model all sequences of words
 - for “fixed” length tasks: no problem - n fixed, sum is 1
 - *tagging, OCR/handwriting (if words identified ahead of time)*
 - for “variable” length tasks: have to account for
 - *discount shorter sentences*
- General model:
 - for each sequence of words of length n , define:
 $p'(w) = \lambda_n p(w)$ such that $\sum_{n=1.. \infty} \lambda_n = 1 \Rightarrow \sum_{n=1.. \infty} \sum_{w \in \Omega^n} p'(w) = 1$
 - e.g., estimate λ_n from data; or use normal or other distribution

Parameter Estimation

- Parameter: numerical value needed to compute $p(w|h)$
- From data (how else?)
- Data preparation:
 - get rid of formatting etc. (“text cleaning”)
 - define words (separate but include punctuation, call it “word”)
 - define sentence boundaries (insert “words” <s> and </s>)
 - letter case: keep, discard, or be smart:
 - name recognition
 - number type identification
 - [these are huge problems per se!]
 - numbers: keep, replace by <num>, or be smart (form ~ pronunciation)

Maximum Likelihood Estimate

- MLE: Relative Frequency ...
 - ...best predicts the data at hand (the “training data”)
- Trigrams from Training Data T:
 - count sequences of three words in T: $c_3(w_{i-2}, w_{i-1}, w_i)$
(notation: just saying that the three words follow each other)
 - count sequences of two words in T: $c_2(w_{i-1}, w_i)$:
 - either use $c_2(y, z) = \sum_w c_3(y, z, w)$
 - or count differently at the beginning (& end) of data!

$$p(w_i | w_{i-2}, w_{i-1}) = \text{est. } c_3(w_{i-2}, w_{i-1}, w_i) / c_2(w_{i-2}, w_{i-1})$$

Character Language Model

- Use individual characters instead of words:

$$p(W) = \prod_{i=1..d} p(c_i | c_{i-n+1}, c_{i-n+2}, \dots, c_{i-1})$$

- Same formulas etc.
- Might consider 4-grams, 5-grams or even more
- Good for language comparison (how?)
- Transform cross-entropy between letter- and word-based models:

$$H_S(p_{\text{char}}) = H_S(p_{\text{word}}) / \text{avg. \# of characters per word in } S$$

LM: an Example

- Training data: <s> <s> He can buy the can of soda .

LM: an Example

- Training data: $\langle s \rangle \langle s \rangle$ He can buy the can of soda .
- Unigram model (n=8):
 - $p_1(\text{He}) = p_1(\text{buy}) = p_1(\text{the}) = p_1(\text{of}) = p_1(\text{soda}) = p_1(.) = 0.125,$
 $p_1(\text{can}) = 0.25$

LM: an Example

- Training data: $\langle s \rangle \langle s \rangle$ He can buy the can of soda .
- Unigram model (n=8):
 - $p_1(\text{He}) = p_1(\text{buy}) = p_1(\text{the}) = p_1(\text{of}) = p_1(\text{soda}) = p_1(.) = 0.125,$
 $p_1(\text{can}) = 0.25$
- Bigram model:
 - $p_2(\text{He}|\langle s \rangle) = 1, p_2(\text{can}|\text{He}) = 1, p_2(\text{buy}|\text{can}) = 0.5, p_2(\text{of}|\text{can}) = 0.5,$
 $p_2(\text{the}|\text{buy}) = 1, \dots$

LM: an Example

- Training data: $\langle s \rangle \langle s \rangle \text{ He can buy the can of soda .}$
- Unigram model (n=8):
 - $p_1(\text{He}) = p_1(\text{buy}) = p_1(\text{the}) = p_1(\text{of}) = p_1(\text{soda}) = p_1(.) = 0.125,$
 $p_1(\text{can}) = 0.25$
- Bigram model:
 - $p_2(\text{He}|\langle s \rangle) = 1, p_2(\text{can}|\text{He}) = 1, p_2(\text{buy}|\text{can}) = 0.5, p_2(\text{of}|\text{can}) = 0.5,$
 $p_2(\text{the}|\text{buy}) = 1, \dots$
- Trigram model:
 - $p_3(\text{He}|\langle s \rangle, \langle s \rangle) = 1, p_3(\text{can}|\langle s \rangle, \text{He}) = 1, p_3(\text{buy}|\text{He}, \text{can}) = 1,$
 $p_3(\text{of}|\text{the}, \text{can}) = 1, \dots, p_3(\cdot|\text{of}, \text{soda}) = 1$

LM: an Example

- Training data: $\langle s \rangle \langle s \rangle$ He can buy the can of soda .
- Unigram model (n=8):
 - $p_1(\text{He}) = p_1(\text{buy}) = p_1(\text{the}) = p_1(\text{of}) = p_1(\text{soda}) = p_1(.) = 0.125,$
 $p_1(\text{can}) = 0.25$
- Bigram model:
 - $p_2(\text{He}|\langle s \rangle) = 1, p_2(\text{can}|\text{He}) = 1, p_2(\text{buy}|\text{can}) = 0.5, p_2(\text{of}|\text{can}) = 0.5,$
 $p_2(\text{the}|\text{buy}) = 1, \dots$
- Trigram model:
 - $p_3(\text{He}|\langle s \rangle, \langle s \rangle) = 1, p_3(\text{can}|\langle s \rangle, \text{He}) = 1, p_3(\text{buy}|\text{He}, \text{can}) = 1,$
 $p_3(\text{of}|\text{the}, \text{can}) = 1, \dots, p_3(\cdot|\text{of}, \text{soda}) = 1$
- Entropy: $H(p_1) = 2.75, H(p_2) = 0.25, H(p_3) = 0 \quad \leftarrow \underline{\text{Great?}}$

LM: an Example (The Problem)

- Test data:

$S = \langle s \rangle \langle s \rangle$ It was the greatest buy of all.

- Cross entropy:


- Even $H_S(p_1)$ fails ($= H_S(p_2) = H_S(p_3) = \infty$), because:
 - all unigrams but $p_1(\text{the})$, $p_1(\text{buy})$, $p_1(\text{of})$ and $p_1(\cdot)$ are 0.
 - all bigram probabilities are 0.
 - all trigram probabilities are 0.

- We want:

- to make all (theoretically possible) probabilities non-zero.

Real World Situation (from Lecture 1)

- Unable to specify set of grammatical sentences today using fixed “categorical” rules (maybe never, cf. arguments in M&S)
- Use statistical “model” based on **real world data** and care about the best sentence only (disregarding the “grammaticality” issue)



Moodle Quiz

Moodle Quiz



<https://dl1.cuni.cz/course/view.php?id=18547>