

Statistical Methods in Natural Language Processing

4. Language model smoothing

Pavel Pecina, Jindřich Helcl

4 November, 2025


Course Segments

1. Introduction, probability, essential information theory
2. Statistical language modelling (n-gram)
3. Statistical properties of words
4. Word embeddings
5. Hidden Markov models, Tagging

Recap from Last Week

Noisy Channel

Prototypical case:



- Model: probability of error (noise)
- Example: $p(0|1) = 0.25$, $p(1|1) = 0.75$, $p(1|0) = 0.5$, $p(0|0) = 0.5$
- The Task:
 - known: the noisy output; want to know: the input (decoding)

Noisy Channel: The Golden Rule of ... OCR, HR, ASR, MT, ...

- Recall:

$$A_{\text{best}} = \operatorname{argmax}_A p(A|B)$$

$$A_{\text{best}} = \operatorname{argmax}_A p(B|A) \ p(A) / \cancel{p(B)}$$

(Bayes Formula)

$$A_{\text{best}} = \operatorname{argmax}_A p(B|A) \ p(A)$$

(The Golden Rule)


- Where:

- $p(B|A)$: the acoustic/image/translation/lexical model
 - application-specific name
 - will explore later
- $p(A)$: **the language model**

Markov Chain

- Unlimited memory (cf. previous slide):
 - for w_i , we know all its predecessors $w_1, w_2, w_3, \dots, w_{i-1}$
- Limited memory:
 - we disregard “too old” predecessors
 - remember only k previous words: $w_{i-k}, w_{i-k+1}, \dots, w_{i-1}$
 - called “ k^{th} order Markov approximation”
- + stationary character (no change over time) → Markov Chain:

$$p(W) \cong \prod_{i=1..d} p(w_i | w_{i-k}, w_{i-k+1}, \dots, w_{i-1}), d = |W|$$



n-gram Language Model

- $(n-1)^{\text{th}}$ order Markov approximation $\rightarrow n\text{-gram Language Model}:$

$$p(W) =^{\text{df}} \prod_{i=1..d} p(w_i | w_{i-n+1}, w_{i-n+2}, \dots, w_{i-1})$$

prediction history

```
graph TD; subgraph Box [ ]; p_W["p(W) =df \prod_{i=1..d} p(wi | wi-n+1, wi-n+2, ..., wi-1)"]; end; prediction[prediction] --> p_W; history[history] --> w_hist[wi-n+1, wi-n+2, ..., wi-1];
```

- In particular (assume vocabulary $|V| = 60k$):
 - 0-gram LM: uniform model, $p(w) = 1/|V|$, 1 parameter
 - 1-gram LM: unigram model, $p(w)$, 6×10^4 parameters
 - 2-gram LM: bigram model, $p(w_i | w_{i-1})$, 3.6×10^9 parameters
 - 3-gram LM: trigram model, $p(w_i | w_{i-2}, w_{i-1})$, 2.16×10^{14} parameters

Maximum Likelihood Estimate

- MLE: Relative Frequency ...
 - ...best predicts the data at hand (the “training data”)
- Trigrams from Training Data T:
 - count sequences of three words in T: $c_3(w_{i-2}, w_{i-1}, w_i)$
(notation: just saying that the three words follow each other)
 - count sequences of two words in T: $c_2(w_{i-1}, w_i)$:
 - either use $c_2(y, z) = \sum_w c_3(y, z, w)$
 - or count differently at the beginning (& end) of data!

$$p(w_i | w_{i-2}, w_{i-1}) = \text{est. } c_3(w_{i-2}, w_{i-1}, w_i) / c_2(w_{i-2}, w_{i-1})$$

LM: an Example

- Training data: $\langle s \rangle \langle s \rangle$ He can buy the can of soda .
- Unigram model (n=8):
 - $p_1(\text{He}) = p_1(\text{buy}) = p_1(\text{the}) = p_1(\text{of}) = p_1(\text{soda}) = p_1(.) = 0.125,$
 $p_1(\text{can}) = 0.25$
- Bigram model:
 - $p_2(\text{He}|\langle s \rangle) = 1, p_2(\text{can}|\text{He}) = 1, p_2(\text{buy}|\text{can}) = 0.5, p_2(\text{of}|\text{can}) = 0.5,$
 $p_2(\text{the}|\text{buy}) = 1, \dots$
- Trigram model:
 - $p_3(\text{He}|\langle s \rangle, \langle s \rangle) = 1, p_3(\text{can}|\langle s \rangle, \text{He}) = 1, p_3(\text{buy}|\text{He}, \text{can}) = 1,$
 $p_3(\text{of}|\text{the}, \text{can}) = 1, \dots, p_3(\cdot|\text{of}, \text{soda}) = 1$
- Entropy: $H(p_1) = 2.75, H(p_2) = 0.25, H(p_3) = 0 \quad \leftarrow \underline{\text{Great?}}$

LM: an Example (The Problem)

- Test data:

$S = \langle s \rangle \langle s \rangle$ It was the greatest buy of all.

- Cross entropy:

- Even $H_S(p_1)$ fails ($= H_S(p_2) = H_S(p_3) = \infty$), because:
 - all unigrams but $p_1(\text{the})$, $p_1(\text{buy})$, $p_1(\text{of})$ and $p_1(\cdot)$ are 0.
 - all bigram probabilities are 0.
 - all trigram probabilities are 0.

- We want:

- to make all (theoretically possible) probabilities non-zero.

Language Model Smoothing

The Zero Problem

- “Raw” n-gram language model estimate:
 - necessarily, some zeros
 - !many: trigram model $\rightarrow 2.16 \times 10^{14}$ parameters, data $\sim 10^9$ words
 - which are true 0?
 - optimal situation: even the least frequent trigram would be seen several times, in order to distinguish its probability vs. other trigrams
 - optimal situation cannot happen, unfortunately (open question: how many data would we need?)
 - \rightarrow we don't know
 - we must eliminate the zeros
- Two kinds of zeros: $p(w|h) = 0$, or even $p(h) = 0$!

Why do we need Nonzero Probs?

- To avoid infinite Cross Entropy:
 - happens when an event is found in test data which has not been seen in training data

$H(p) = \infty$: prevents comparing data with > 0 “errors”

- To make the system more robust
 - low count estimates:
 - they typically happen for “detailed” but relatively rare appearances
 - high count estimates:
 - reliable but less “detailed”

Eliminating the Zero Probabilities: Smoothing

- Get new $p'(w)$ (same Ω): almost $p(w)$ but no zeros
- Discount w for (some) $p(w) > 0$: new $p'(w) < p(w)$

$$\sum_{w \in \text{discounted}} (p(w) - p'(w)) = D$$

- Distribute D to all w; $p(w) = 0$: new $p'(w) > p(w)$
 - possibly also to other w with low $p(w)$
- For some w (possibly): $p'(w) = p(w)$
- Make sure $\sum_{w \in \Omega} p'(w) = 1$
- There are many ways of **smoothing**

Smoothing Example

- We often want to make predictions from sparse statistics

$P(w|denied\ the)$


3 allegations

2 reports

1 claims

1 request

7 total



- Smoothing flattens spiky distributions so they generalize better

$P(w|denied\ the)$

2.5 allegations


1.5 reports

0.5 claims

0.5 request

2 other

7 total



Smoothing by Adding 1 (Laplace Smoothing)

- Simplest but not really usable:
 - Predicting words w from a vocabulary V , training data T :
$$p'(w|h) = (c(h,w) + 1) / (c(h) + |V|)$$
 - for non-conditional distributions: $p'(w) = (c(w) + 1) / (|T| + |V|)$
 - Problem if $|V| > c(h)$ (as is often the case; even $\gg c(h)!$)

Smoothing by Adding 1 (Laplace Smoothing)

- Simplest but not really usable:
 - Predicting words w from a vocabulary V , training data T :

$$p'(w|h) = (c(h,w) + 1) / (c(h) + |V|)$$

- for non-conditional distributions: $p'(w) = (c(w) + 1) / (|T| + |V|)$

- Problem if $|V| > c(h)$ (as is often the case; even $\gg c(h)!$)

- Example:

- Training data T : <s> what is it what is small ? $|T| = 8$
- Vocabulary $V = \{ \text{what, is, it, small, ?, } <\text{s}>, \text{flying, birds, are, a, bird, .} \}$ $|V| = 12$
- $p(\text{it}) = 0.125$, $p(\text{what}) = 0.25$, $p(\cdot) = 0$
 $p(\text{what is it?}) = 0.25^2 \times 0.125^2 \approx 0.001$, $p(\text{it is flying.}) = 0.125 \times 0.25 \times 0^2 = 0$
- $p'(\text{it}) = 0.1$, $p'(\text{what}) = 0.15$, $p'(\cdot) = 0.05$
 $p'(\text{what is it?}) = 0.15^2 \times 0.1^2 \approx 0.0002$, $p'(\text{it is flying.}) = 0.1 \times 0.15 \times 0.05^2 \approx 0.00004$

Adding *less* than 1

- Equally simple:
 - Predicting words w from a vocabulary V , training data T :
$$p'(w|h) = (c(h,w) + \lambda) / (c(h) + \lambda|V|), \lambda < 1$$
 - for non-conditional distributions: $p'(w) = (c(w) + \lambda) / (|T| + \lambda|V|)$

Adding *less* than 1

- Equally simple:
 - Predicting words w from a vocabulary V , training data T :
$$p'(w|h) = (c(h,w) + \lambda) / (c(h) + \lambda|V|), \lambda < 1$$
 - for non-conditional distributions: $p'(w) = (c(w) + \lambda) / (|T| + \lambda|V|)$
- Example:
 - Training data T : <s> what is it what is small ? $|T| = 8$
 - Vocabulary $V = \{ \text{what, is, it, small, ?, } <\text{s}\>, \text{flying, birds, are, a, bird, .} \}$ $|V| = 12$
 - $p(\text{it}) = 0.125, p(\text{what}) = 0.25, p(\cdot) = 0$
 $p(\text{what is it?}) = 0.25^2 \times 0.125^2 \approx 0.001, p(\text{it is flying.}) = 0.125 \times 0.25 \times 0^2 = 0$
 - **Use $\lambda = 0.1$:** $p'(\text{it}) \approx 0.12, p'(\text{what}) \approx 0.23, p'(\cdot) \approx 0.01$
 $p'(\text{what is it?}) = 0.23^2 \times 0.12^2 \approx 0.0007, p'(\text{it is flying.}) = 0.12 \times 0.23 \times 0.01^2 \approx 0.00003$

Good-Turing Smoothing

- Suitable for estimation from large data
 - Estimate probability of things that occur c times with the probability of things that occur $c+1$ times:
$$c' = (c+1) \times N(c + 1) / N(c)$$
 - Full formula:
$$p_r(w) = (c(w) + 1) \times N(c(w) + 1) / (|T| \times N(c(w))),$$
where $N(c)$ is the count of words with count c (count-of-counts)
specifically, for $c(w) = 0$ (unseen words), $p_r(w) = N(1) / (|T| \times N(0))$
- good for small counts (< 5-10, where $N(c)$ is high)
- variants (see M&S)
- normalization! (so that we have $\sum_w p'(w) = 1$)

Good-Turing: An Example

- Good-Turing formula: $p_r(w) = (c(w) + 1) \times N(c(w) + 1) / (|T| \times N(c(w)))$
- Training data T: <s> what is it what is small ? $|T| = 8$
- Vocabulary V = { what, is, it, small, ?, <s>, flying, birds, are, a, bird, . } $|V| = 12$
- $p(it) = 0.125, p(what) = 0.25, p(.) = 0$
 $p(what\ is\ it?) = 0.25^2 \times 0.125^2 \approx 0.001, p(it\ is\ flying.) = 0.125 \times 0.25 \times 0^2 = 0$


Good-Turing: An Example

- Good-Turing formula: $p_r(w) = (c(w) + 1) \times N(c(w) + 1) / (|T| \times N(c(w)))$
- Training data T: <s> what is it what is small ? $|T| = 8$
- Vocabulary V = { what, is, it, small, ?, <s>, flying, birds, are, a, bird, . } $|V| = 12$
- $p(it) = 0.125, p(what) = 0.25, p(.) = 0$
 $p(\text{what is it?}) = 0.25^2 \times 0.125^2 \approx 0.001, p(\text{it is flying.}) = 0.125 \times 0.25 \times 0^2 = 0$
- Raw reestimation ($N(0) = 6, N(1) = 4, N(2) = 2, N(i) = 0$ for $i > 2$):
 $p_r(it) = (1+1) \times N(1+1) / (8 \times N(1)) = 2 \times 2 / (8 \times 4) = 0.125$
 $p_r(\text{what}) = (2+1) \times N(2+1) / (8 \times N(2)) = 3 \times 0 / (8 \times 2) = 0$: keep orig. $p(\text{what})$
 $p_r(.) = (0+1) \times N(0+1) / (8 \times N(0)) = 1 \times 4 / (8 \times 6) \approx 0.083$

Good-Turing: An Example

- Good-Turing formula: $p_r(w) = (c(w) + 1) \times N(c(w) + 1) / (|T| \times N(c(w)))$
- Training data T: <s> what is it what is small ? $|T| = 8$
- Vocabulary V = { what, is, it, small, ?, <s>, flying, birds, are, a, bird, . } $|V| = 12$
- $p(it) = 0.125, p(what) = 0.25, p(.) = 0$
 $p(what is it?) = 0.25^2 \times 0.125^2 \approx 0.001, p(it is flying.) = 0.125 \times 0.25 \times 0^2 = 0$
- Raw reestimation ($N(0) = 6, N(1) = 4, N(2) = 2, N(i) = 0$ for $i > 2$):
 $p_r(it) = (1+1) \times N(1+1) / (8 \times N(1)) = 2 \times 2 / (8 \times 4) = 0.125$
 $p_r(what) = (2+1) \times N(2+1) / (8 \times N(2)) = 3 \times 0 / (8 \times 2) = 0$: keep orig. $p(what)$
 $p_r(.) = (0+1) \times N(0+1) / (8 \times N(0)) = 1 \times 4 / (8 \times 6) \approx 0.083$
- Normalize (divide by $1.5 = \sum_{w \in |V|} p_r(w)$) and compute:
 $p'(it) \approx 0.08, p'(what) \approx 0.17, p'(.) \approx 0.06$
 $p'(what is it?) = 0.17^2 \times 0.08^2 \approx 0.0002, p'(it is flying.) = 0.08 \times 0.17 \times 0.06^2 \approx 0.00004$

Smoothing by Combination: Linear Interpolation

- Combine what?
 - distributions of various level of detail vs. reliability
- n-gram models:
 - use (n-1)gram, (n-2)gram, ..., uniform

A horizontal double-headed arrow pointing from left to right, with the word "reliability" above it and the word "detail" below it.
- Simplest possible combination:
 - sum of probabilities, normalize:
 - bigram: $p(0|0) = 0.8, p(1|0) = 0.2, p(0|1) = 1, p(1|1) = 0$
 - unigram: $p(0) = 0.4, p(1) = 0.6$:
 - combined: $p'(0|0) = 0.6, p'(1|0) = 0.4, p'(0|1) = 0.7, p'(1|1) = 0.3$

Language Model Interpolation

- Weight in less detailed distributions using $\lambda = (\lambda_0, \lambda_1, \lambda_2, \lambda_3)$:

$$p'_{\lambda}(w_i | w_{i-2}, w_{i-1}) = \lambda_3 p_3(w_i | w_{i-2}, w_{i-1}) + \lambda_2 p_2(w_i | w_{i-1}) + \lambda_1 p_1(w_i) + \lambda_0 / |V|$$

- Normalize:

$$\lambda_i > 0, \sum_{i=0..n} \lambda_i = 1 \quad (\lambda_0 = 1 - \sum_{i=1..n} \lambda_i) \quad (n=3)$$

- Estimation using MLE:

1. Fix the p_3, p_2, p_1 and $|V|$ parameters as estimated from training data
2. Find such $\{\lambda_i\}$ which minimizes the cross entropy (maximizes prob. of **data**):

$$-(1/|D|) \sum_{i=1..|D|} \log_2(p'_{\lambda}(w_i | h_i))$$

What Data?

- Training data T? But we will always get $\lambda_3 = 1$!
 - Why? (let p_{iT} be an i-gram distribution ML-estimated from T)
 - minimizing $H_T(p'_{\lambda})$ over a vector λ , $p'_{\lambda} = \lambda_3 p_{3T} + \lambda_2 p_{2T} + \lambda_1 p_{1T} + \lambda_0 / |V|$
 - $H_T(p'_{\lambda}) = H(p_{3T}) + D(p_{3T} \| p'_{\lambda})$; p_{3T} fixed $\rightarrow H(p_{3T})$ fixed, and it is the best
 - which p'_{λ} minimizes $H_T(p'_{\lambda})$?
 - ... a p'_{λ} for which $D(p_{3T} \| p'_{\lambda}) = 0$
 - ... and that's p_{3T} (because $D(p \| p) = 0$, as we know).
 - ... and certainly $p'_{\lambda} = p_{3T}$ if $\lambda_3 = 1$ (maybe in some other cases, too).
 $(p'_{\lambda} = 1 \times p_{3T} + 0 \times p_{2T} + 0 \times p_{1T} + 0 / |V|)$

What Data?

- Training data T? But we will always get $\lambda_3 = 1$!
 - Why? (let p_{iT} be an i-gram distribution ML-estimated from T)
 - minimizing $H_T(p'_{\lambda})$ over a vector λ , $p'_{\lambda} = \lambda_3 p_{3T} + \lambda_2 p_{2T} + \lambda_1 p_{1T} + \lambda_0 / |V|$
 - $H_T(p'_{\lambda}) = H(p_{3T}) + D(p_{3T} \| p'_{\lambda})$; p_{3T} fixed $\rightarrow H(p_{3T})$ fixed, and it is the best
 - which p'_{λ} minimizes $H_T(p'_{\lambda})$?
 - ... a p'_{λ} for which $D(p_{3T} \| p'_{\lambda}) = 0$
 - ... and that's p_{3T} (because $D(p \| p) = 0$, as we know).
 - ... and certainly $p'_{\lambda} = p_{3T}$ if $\lambda_3 = 1$ (maybe in some other cases, too).
 $(p'_{\lambda} = 1 \times p_{3T} + 0 \times p_{2T} + 0 \times p_{1T} + 0 / |V|)$
- Thus: do not use the training data for estimation of λ !
 - hold out part of the training data (heldout H); remaining data is the true training data T, the test data S still must be some different data!

The Formulas

- Minimizing $-(1/|H|) \sum_{i=1..|H|} \log_2(p'_{\lambda}(w_i|h_i))$ over λ :

$$\begin{aligned} p'_{\lambda}(w_i | h_i) &= p'_{\lambda}(w_i | w_{i-2}, w_{i-1}) = \\ &= \lambda_3 p_3(w_i | w_{i-2}, w_{i-1}) + \lambda_2 p_2(w_i | w_{i-1}) + \lambda_1 p_1(w_i) + \lambda_0 / |V| \end{aligned}$$

- "Expected Counts (of lambdas)": $j = 0,..,3$

$$c(\lambda_j) = \sum_{i=1..|H|} (\lambda_j p_j(w_i|h_i) / p'_{\lambda}(w_i|h_i))$$

- "Next λ ": $j = 0,..,3$

$$\lambda_{j,\text{next}} = c(\lambda_j) / \sum_{k=0..3} (c(\lambda_k))$$

The (Smoothing) EM Algorithm

1. Start with some λ , such that $\lambda_j > 0$ for all $j \in 0, \dots, 3$.
 2. Compute “Expected Counts” for each λ_j .
 3. Compute new set of λ_j , using the “Next λ ” formula.
 4. Start over at step 2, unless a termination condition is met.
-
- Termination condition: convergence of λ
 - Simply set an ε , and finish if $|\lambda_{j,\text{old}} - \lambda_{j,\text{new}}| < \varepsilon$ for each j (step 3).
 - Guaranteed to converge:
 - Follows from Jensen’s inequality, plus a technical proof

Simple Example

- Raw distribution (unigram only; smooth with uniform): $|V|=26$
 $p(a)=0.25, p(b)=0.5, p(\alpha)=1/64$ for $\alpha \in \{c..r\}$, = 0 for the rest: s,t,u,v,w,x,y,z
- Heldout data: baby; use one set of λ (λ_1 : unigram, λ_0 : uniform)

Simple Example

- Raw distribution (unigram only; smooth with uniform): $|V|=26$
 $p(a)=0.25, p(b)=0.5, p(\alpha)=1/64$ for $\alpha \in \{c..r\}$, = 0 for the rest: s,t,u,v,w,x,y,z
- Heldout data: baby; use one set of λ (λ_1 : unigram, λ_0 : uniform)
- Start with $\lambda_1 = 0.5, \lambda_0 = 0.5$;
 - $p'_{\lambda}(b) = 0.5 \times 0.5 + 0.5 \times 1/26 = 0.27$

Simple Example

- Raw distribution (unigram only; smooth with uniform): $|V|=26$
 $p(a)=0.25, p(b)=0.5, p(\alpha)=1/64$ for $\alpha \in \{c..r\}$, = 0 for the rest: s,t,u,v,w,x,y,z
- Heldout data: baby; use one set of λ (λ_1 : unigram, λ_0 : uniform)
- Start with $\lambda_1 = 0.5, \lambda_0 = 0.5$;
 - $p'_{\lambda}(b) = 0.5 \times 0.5 + 0.5 \times 1/26 = 0.27$

Simple Example

- Raw distribution (unigram only; smooth with uniform): $|V|=26$
 $p(a)=0.25, p(b)=0.5, p(\alpha)=1/64$ for $\alpha \in \{c..r\}$, = 0 for the rest: s,t,u,v,w,x,y,z
- Heldout data: baby; use one set of λ (λ_1 : unigram, λ_0 : uniform)
- Start with $\lambda_1 = 0.5, \lambda_0 = 0.5$;
 - $p'_{\lambda}(b) = 0.5 \times 0.5 + 0.5 \times 1/26 = 0.27$
 - $p'_{\lambda}(a) = 0.5 \times 0.25 + 0.5 \times 1/26 = 0.14$
 - $p'_{\lambda}(y) = 0.5 \times 0 + 0.5 \times 1/26 = 0.02$

Simple Example

- Raw distribution (unigram only; smooth with uniform): $|V|=26$
 $p(a)=0.25, p(b)=0.5, p(\alpha)=1/64$ for $\alpha \in \{c..r\}$, = 0 for the rest: s,t,u,v,w,x,y,z
- Heldout data: baby; use one set of λ (λ_1 : unigram, λ_0 : uniform)
- Start with $\lambda_1 = 0.5, \lambda_0 = 0.5$;
 - $p'_{\lambda}(b) = 0.5 \times 0.5 + 0.5 \times 1/26 = 0.27$
 - $p'_{\lambda}(a) = 0.5 \times 0.25 + 0.5 \times 1/26 = 0.14$
 - $p'_{\lambda}(y) = 0.5 \times 0 + 0.5 \times 1/26 = 0.02$

Simple Example

- Raw distribution (unigram only; smooth with uniform): $|V|=26$
 $p(a)=0.25, p(b)=0.5, p(\alpha)=1/64$ for $\alpha \in \{c..r\}$, = 0 for the rest: s,t,u,v,w,x,y,z
- Heldout data: baby; use one set of λ (λ_1 : unigram, λ_0 : uniform)
- Start with $\lambda_1 = 0.5, \lambda_0 = 0.5$;
 - $p'_{\lambda}(b) = 0.5 \times 0.5 + 0.5 \times 1/26 = 0.27$
 - $p'_{\lambda}(a) = 0.5 \times 0.25 + 0.5 \times 1/26 = 0.14$
 - $p'_{\lambda}(y) = 0.5 \times 0 + 0.5 \times 1/26 = 0.02$
 - $c(\lambda_1) = 0.5 \times 0.5 / 0.27 + 0.5 \times 0.25 / 0.14 + 0.5 \times 0.5 / 0.27 + 0.5 \times 0 / 0.02 = 2.72$

Simple Example

- Raw distribution (unigram only; smooth with uniform): $|V|=26$
 $p(a)=0.25, p(b)=0.5, p(\alpha)=1/64$ for $\alpha \in \{c..r\}$, = 0 for the rest: s,t,u,v,w,x,y,z
- Heldout data: baby; use one set of λ (λ_1 : unigram, λ_0 : uniform)
- Start with $\lambda_1 = 0.5, \lambda_0 = 0.5$;
 - $p'_{\lambda}(b) = 0.5 \times 0.5 + 0.5 \times 1/26 = 0.27$
 - $p'_{\lambda}(a) = 0.5 \times 0.25 + 0.5 \times 1/26 = 0.14$
 - $p'_{\lambda}(y) = 0.5 \times 0 + 0.5 \times 1/26 = 0.02$
 - $c(\lambda_1) = 0.5 \times 0.5 / 0.27 + 0.5 \times 0.25 / 0.14 + 0.5 \times 0.5 / 0.27 + 0.5 \times 0 / 0.02 = 2.72$

Simple Example

- Raw distribution (unigram only; smooth with uniform): $|V|=26$
 $p(a)=0.25, p(b)=0.5, p(\alpha)=1/64$ for $\alpha \in \{c..r\}$, = 0 for the rest: s,t,u,v,w,x,y,z
- Heldout data: baby; use one set of λ (λ_1 : unigram, λ_0 : uniform)
- Start with $\lambda_1 = 0.5, \lambda_0 = 0.5$;
 - $p'_{\lambda}(b) = 0.5 \times 0.5 + 0.5 \times 1/26 = 0.27$
 - $p'_{\lambda}(a) = 0.5 \times 0.25 + 0.5 \times 1/26 = 0.14$
 - $p'_{\lambda}(y) = 0.5 \times 0 + 0.5 \times 1/26 = 0.02$
 - $c(\lambda_1) = 0.5 \times 0.5/0.27 + 0.5 \times 0.25/0.14 + 0.5 \times 0.5/0.27 + 0.5 \times 0/0.02 = 2.72$
 - $c(\lambda_0) = 0.5 \times 0.04/0.27 + 0.5 \times 0.04/0.14 + 0.5 \times 0.04/0.27 + 0.5 \times 0.04/0.02 = 1.28$

Simple Example

- Raw distribution (unigram only; smooth with uniform): $|V|=26$
 $p(a)=0.25, p(b)=0.5, p(\alpha)=1/64$ for $\alpha \in \{c..r\}$, = 0 for the rest: s,t,u,v,w,x,y,z
- Heldout data: baby; use one set of λ (λ_1 : unigram, λ_0 : uniform)
- Start with $\lambda_1 = 0.5, \lambda_0 = 0.5$;
 - $p'_{\lambda}(b) = 0.5 \times 0.5 + 0.5 \times 1/26 = 0.27$
 - $p'_{\lambda}(a) = 0.5 \times 0.25 + 0.5 \times 1/26 = 0.14$
 - $p'_{\lambda}(y) = 0.5 \times 0 + 0.5 \times 1/26 = 0.02$
 - $c(\lambda_1) = 0.5 \times 0.5/0.27 + 0.5 \times 0.25/0.14 + 0.5 \times 0.5/0.27 + 0.5 \times 0/0.02 = 2.72$
 - $c(\lambda_0) = 0.5 \times 0.04/0.27 + 0.5 \times 0.04/0.14 + 0.5 \times 0.04/0.27 + 0.5 \times 0.04/0.02 = 1.28$

Simple Example

- Raw distribution (unigram only; smooth with uniform): $|V|=26$
 $p(a)=0.25, p(b)=0.5, p(\alpha)=1/64$ for $\alpha \in \{c..r\}$, = 0 for the rest: s,t,u,v,w,x,y,z
- Heldout data: baby; use one set of λ (λ_1 : unigram, λ_0 : uniform)
- Start with $\lambda_1 = 0.5, \lambda_0 = 0.5$;
 - $p'_{\lambda}(b) = 0.5 \times 0.5 + 0.5 \times 1/26 = 0.27$
 - $p'_{\lambda}(a) = 0.5 \times 0.25 + 0.5 \times 1/26 = 0.14$
 - $p'_{\lambda}(y) = 0.5 \times 0 + 0.5 \times 1/26 = 0.02$
 - $c(\lambda_1) = 0.5 \times 0.5/0.27 + 0.5 \times 0.25/0.14 + 0.5 \times 0.5/0.27 + 0.5 \times 0/0.02 = 2.72$
 - $c(\lambda_0) = 0.5 \times 0.04/0.27 + 0.5 \times 0.04/0.14 + 0.5 \times 0.04/0.27 + 0.5 \times 0.04/0.02 = 1.28$
 - Normalize: $\lambda_{1,next} = 0.68, \lambda_{0,next} = 0.32$
 - Repeat from Step 2 (recompute p'_{λ} first for efficient computation, then $c(\lambda_i), \dots)$

Simple Example

- Raw distribution (unigram only; smooth with uniform): $|V|=26$
 $p(a)=0.25, p(b)=0.5, p(\alpha)=1/64$ for $\alpha \in \{c..r\}$, = 0 for the rest: s,t,u,v,w,x,y,z
- Heldout data: baby; use one set of λ (λ_1 : unigram, λ_0 : uniform)
- Start with $\lambda_1 = 0.5, \lambda_0 = 0.5$;
 - $p'_{\lambda}(b) = 0.5 \times 0.5 + 0.5 \times 1/26 = 0.27$
 - $p'_{\lambda}(a) = 0.5 \times 0.25 + 0.5 \times 1/26 = 0.14$
 - $p'_{\lambda}(y) = 0.5 \times 0 + 0.5 \times 1/26 = 0.02$
 - $c(\lambda_1) = 0.5 \times 0.5/0.27 + 0.5 \times 0.25/0.14 + 0.5 \times 0.5/0.27 + 0.5 \times 0/0.02 = 2.72$
 - $c(\lambda_0) = 0.5 \times 0.04/0.27 + 0.5 \times 0.04/0.14 + 0.5 \times 0.04/0.27 + 0.5 \times 0.04/0.02 = 1.28$
 - Normalize: $\lambda_{1,next} = 0.68, \lambda_{0,next} = 0.32$
 - Repeat from Step 2 (recompute p'_{λ} first for efficient computation, then $c(\lambda_i), \dots)$
- Finish when new λ almost equal to the old ones (say, < 0.01 difference).

Some More Technical Hints

- Set $V = \{\text{all words from training data}\}$.
 - You may also consider $V = T \cup H$, but it does not make the coding in any way simpler (in fact, harder).
 - But: you must never use the test data for your vocabulary!
- Prepend two “words” in front of all data ($\langle s \rangle$):
 - avoids beginning-of-data problems
 - call these index -1 and 0: then the formulas hold exactly
- When $c_n(w,h) = 0$:
 - Assign 0 probability to $p_n(w|h)$ where $c_{n-1}(h) > 0$, but a uniform probability $(1/|V|)$ to those $p_n(w|h)$ where $c_{n-1}(h) = 0$ [this must be done both when working on the heldout data during EM, as well as when computing cross-entropy on the test data!]

Bucketed Smoothing

- Use several λ vectors instead of one, based on (frequency of) history: $\lambda(h)$
- e.g. for $h=(\text{micrograms}, \text{per})$ we have $\lambda(h) = (0.999, 0.0009, 0.00009, 0.00001)$
(because “cubic” is the only word to follow...)
- Actually: not a separate set for each history, but rather a set for “similar” histories (“bucket”):
$$\lambda(b(h)), \text{ where } b: V^2 \rightarrow N \text{ (in the case of trigrams)}$$
- b classifies histories according to their reliability (~ frequency)

Bucketed Smoothing: The Algorithm

- First, determine the bucketing function b (use heldout!):
 - decide in advance you want e.g. 1000 buckets
 - compute the total frequency of histories in 1 bucket ($f_{\max}(b)$)
 - gradually fill your buckets from the most frequent bigrams so that the sum of frequencies does not exceed $f_{\max}(b)$ (you might end up with slightly more than 1000 buckets)
- Divide your heldout data according to buckets.
- Apply the previous algorithm to each bucket and its data.

Moodle Quiz

Moodle Quiz



<https://dl1.cuni.cz/course/view.php?id=18547>