

Statistical Methods in Natural Language Processing

6. Mutual Information and Word Classes

Pavel Pecina, Jindřich Helcl

18 November, 2025

Course Segments

1. Introduction, probability, essential information theory
2. Statistical language modelling (n-gram)
3. Statistical properties of words
4. Word representations
5. Hidden Markov models, Tagging

Recap from Last Week

Collocations

- J. R. Firth (1957):
 - “*You shall know a word by the company it keeps.*”
 - “Collocations of a given word are statements of the habitual or customary places of that word.”
- M&S (Chapter 5):
 - “A collocation is an expression consisting of two or more words that correspond to some conventional way of saying things.”
- Examples:
 - *strong tea, weapons of mass destruction*
 - *to make up, the rich and powerful*
- Valid or invalid?
 - *a stiff breeze, but not a stiff wind*
 - *broad daylight, but not bright daylight*

Properties of Collocations

- Typical properties/criterions of collocations:
 - non-compositionality
 - non-substitutability
 - non-modifiability
- Collocations usually cannot be translated word-by-word
 - e.g. *take a shower* → *osprchovat se*, but not *vzít sprchu*
- A phrase can be a collocation even if it is not consecutive
 - e.g. *knock ... door*

How to Find Collocations?

- Frequency (simplest method)
 - plain
 - filtered
- Mean and variance of the distance between focal word and collocating word
- Hypothesis testing
 - t test
 - χ^2 test
- Pointwise Mutual Information

Hypothesis Testing

- Two words can co-occur by chance
 - High frequency and low variance can be accidental
- Hypothesis Testing measures the confidence that this co-occurrence was really due to association, and not just due to chance.
- Formulate a **null hypothesis** H_0 that there is no association between the words beyond chance occurrences:
 - H_0 states what should be true if two words do not form a collocation.
 - If H_0 can be rejected, the words do not co-occur by chance, and they form a collocation
- Compute the probability p that the event would occur if H_0 were true:
 - reject H_0 if p is too low (typically beneath a significance level of $p < 0.05, 0.01, \dots$)
 - retain H_0 as possible otherwise.

t-test (Student's *t*-test)

- The test looks at the difference between the **observed** and **expected** means, scaled by the variance of the data, and tells us how likely one is to get a sample of that mean and variance, assuming that the sample is drawn from a normal distribution with mean μ .

$$t = \frac{\bar{x} - \mu}{\sqrt{\frac{s^2}{N}}}$$

- Where \bar{x} is the real data mean (**observed in data**)
- s^2 is the variance
- N is the sample size
- μ is the mean of the distribution (**expected under H_0**)

χ^2 test: An example

- Observed occurrences:

	$w_1 = \text{new}$	$w_1 \neq \text{new}$
$w_2 = \text{companies}$	8 (new companies)	4667 (e.g., old companies)
$w_2 \neq \text{companies}$	15820 (e.g., new machines)	14287181 (e.g., old machines)

- The χ^2 statistic sums differences between observed and expected values in all cells of the table, scaled by the magnitude of the expected values:

$$\chi^2 = \sum_{i,j} \frac{(O_{ij} - E_{ij})^2}{E_{ij}}$$

- where i ranges over rows of the table, j ranges over columns, O_{ij} is the observed value for cell (i, j) and E_{ij} is the expected value.



χ^2 test: A Simpler Form

- χ^2 test can be applied to tables of any size, but it has a simpler form for 2-by-2 tables (i.e. bigram collocations):

$$\chi^2 = \frac{N(O_{11}O_{22} - O_{12}O_{21})^2}{(O_{11} + O_{12})(O_{11} + O_{21})(O_{12} + O_{22})(O_{21} + O_{22})}$$

Pointwise Mutual Information

- An information-theoretically motivated measure for discovering interesting collocations is pointwise mutual information
- It is a measure of how much one word tells us about the other:

$$\text{PMI}(x, y) = \log \frac{P(x, y)}{P(x) P(y)}$$

← “observed”
← “expected (under H_0)”

- $P(x, y)$ is the joint probability of x and y occurring together,
- $P(x)$ and $P(y)$ are the individual probabilities of x and y
- PMI is NOT the MI as defined in Information Theory (MI is the average of PMI)

Mutual Information and Word Classes

The Problem

Not enough data:

- Language Modeling: we do not see “correct” n-grams
 - solution so far: smoothing
- suppose we see:
 - *short homework, short assignment, simple homework*
- but not:
 - *simple assignment*
- What happens to our (bigram) LM?
 - $p(\text{homework} \mid \text{simple}) = \text{high probability}$
 - $p(\text{assignment} \mid \text{simple}) = \text{low probability (smoothed with } p(\text{assignment}))$
- They should be much closer!

Word Classes

Observation:

- similar words behave in a similar way
- i.e. appear in similar context (the Distributional Hypothesis)
- Assuming trigram language model:
 - *a ... homework* (any attribute of homework: *short, simple, late, difficult*),
 - *... the woods* (any verb that has the woods as an object: *walk, cut, save*)
 - *a (short,long,difficult,...) (homework,assignment,task,job,...)*

Solution

Use the Word Classes as the “reliability” measure:

- Example: we see
 - *short homework, short assignment, simple homework*
- but not:
 - *simple assigment*
- Cluster into classes:
 - *(short, simple) (homework, assignment)*
 - covers “*simple assignment*”, too
- Gaining:
 - realistic estimates for unseen n-grams
- Losing:
 - accuracy (level of detail) within classes

The New Model

Rewrite the n-gram LM using classes:

- Original definition: [$k = 1 \dots n$]

$$p_k(w_i|h_i) = c(h_i, w_i) / c(h_i) \quad [\text{history: } (k-1) \text{ words}]$$

- Introduce classes:

$$p_k(w_i|h_i) = p(w_i|c_i) p_k(c_i|h_i)$$

- history: classes, too: [for trigram: $h_i = c_{i-2}, c_{i-1}$, bigram: $h_i = c_{i-1}$]
- Smoothing as usual
 - over $p_k(w_i|h_i)$, where each is defined as above (except uniform which is $1/|V|$)

Training Data

- Suppose we already have a mapping:
 - $r: V \rightarrow C$ assigning each word its class ($c_i = r(w_i)$)
- Expand the training data:
 - $T = (w_1, w_2, \dots, w_{|T|})$ into
 - $T_C = (\langle w_1, r(w_1) \rangle, \langle w_2, r(w_2) \rangle, \dots, \langle w_{|T|}, r(w_{|T|}) \rangle)$
- Effectively, we have two streams of data:
 - word stream: $w_1, w_2, \dots, w_{|T|}$
 - class stream: $c_1, c_2, \dots, c_{|T|}$ (def. as $c_i = r(w_i)$)
- Expand Heldout, Test data too

Training the New Model

- Using ML estimates (as expected):
 - $p(w_i|c_i) = p(w_i|r(w_i)) = c(w_i) / c(r(w_i)) = c(w_i) / c(c_i)$
 - !!! $c(w_i, c_i) = c(w_i)$ [since c_i determined by w_i]
 - $p_k(c_i|h_i)$:
 - $p_3(c_i|h_i) = p_3(c_i|c_{i-2}, c_{i-1}) = c(c_{i-2}, c_{i-1}, c_i) / c(c_{i-2}, c_{i-1})$
 - $p_2(c_i|h_i) = p_2(c_i|c_{i-1}) = c(c_{i-1}, c_i) / c(c_{i-1})$
 - $p_1(c_i|h_i) = p_1(c_i) = c(c_i) / |T|$
- Then smooth as usual
 - not the $p(w_i|c_i)$ nor $p_k(c_i|h_i)$ individually, but the $p_k(w_i|h_i)$

Classes: How To Get Them

- We supposed the classes are given
- Maybe there are in (human) dictionaries, but...
 - dictionaries are incomplete
 - dictionaries are unreliable
 - do not define classes as equivalence relation (overlap)
 - do not define classes suitable for LM
 - small, short... maybe; small and difficult?

→ we have to construct them from data (again...)

Creating the Word-to-Class Map (Brown's Classes)

- Consider bigram model for now.
- Bigram estimate:

$$p_2(c_i|h_i) = p_2(c_i|c_{i-1}) = c(c_{i-1}, c_i) / c(c_{i-1}) = c(r(w_{i-1}), r(w_i)) / c(r(w_{i-1}))$$

- Form of the model:
 - just raw bigram for now:
- Maximize over r (given $r \rightarrow$ fixed p, p_2):
 - define objective

$$L(r) = 1/|T| \sum_{i=1..|T|} \log(p(w_i|r(w_i)) p_2(r(w_i))|r(w_{i-1})))$$

$$r_{\text{best}} = \operatorname{argmax}_r L(r) \quad (L(r) = \text{norm. logprob of training data ... as usual})$$

Simplifying the Objective Function

- Start from $L(r) = 1/|T| \sum_{i=1..|T|} \log(p(w_i|r(w_i)) p_2(r(w_i)|r(w_{i-1})))$:
 $1/|T| \sum_{i=1..|T|} \log(p(w_i|r(w_i)) \underline{p(r(w_i))} p_2(r(w_i)|r(w_{i-1})) / \underline{p(r(w_i))}) =$
 $1/|T| \sum_{i=1..|T|} \log(\underline{p(w_i,r(w_i))} p_2(r(w_i)|r(w_{i-1})) / p(r(w_i))) =$
 $1/|T| \sum_{i=1..|T|} \log(\underline{p(w_i)}) + 1/|T| \sum_{i=1..|T|} \log(p_2(r(w_i)|r(w_{i-1})) / p(r(w_i))) =$
 $-H(W) + 1/|T| \sum_{i=1..|T|} \log(p_2(r(w_i)|r(w_{i-1})) \underline{p(r(w_{i-1}))} / (\underline{p(r(w_{i-1}))} p(r(w_i))))$
=
- $-H(W) + 1/|T| \sum_{i=1..|T|} \log(\underline{p(r(w_i),r(w_{i-1}))} / (\underline{p(r(w_{i-1}))} p(r(w_i)))) =$
 $-H(W) + \sum_{d,e \in C} p(d,e) \log(p(d,e) / (p(d) p(e))) =$
 $-H(W) + I(D,E)$ (event E picks class adjacent (to the right) to the one picked by D)
- Since W does not depend on r, we ended up with $I(D,E)$.

Maximizing Mutual Information (dependent on mapping r)

- Result from previous slide:
 - Maximizing the probability of data amounts to maximizing $I(D, E)$, the Mutual Information of the adjacent classes.
- Good:
 - We know what a MI is, and we know how to maximize.
- Bad:
 - There is no way how to maximize over so many possible partitionings: $|V|^{|V|}$ - no way to test them all.

Training or Heldout?

- Training:
 - best $I(D, E)$: all words in a class of its own
→ will not give us anything new.
- Heldout: ok, but:
 - must smooth to test any possible partitioning (unfeasible):
→ using raw model: 0 probability of heldout (almost) guaranteed
→ will not be able to compare anything
- Solution:
 - use training anyway, but only keep $I(D, E)$ as large as possible

The Greedy Algorithm

- Define merging operation on the mapping $r: V \rightarrow C$:
 - merge: $R \times C \times C \rightarrow R' \times C-1: (r,k,l) \rightarrow r',C'$ such that
 - $C^{-1} = \{C - \{k,l\} \cup \{m\}\}$ (throw out k and l, add new m $\notin C$)
 - $r'(w) = m \quad \text{for } w \in r_{INV}(\{k,l\}),$
 $r(w) \quad \text{otherwise.}$
- 1. Start with each word in its own class ($C = V$), $r = \text{id}$.
- 2. Merge two classes k,l into one, m , such that
$$(k,l) = \operatorname{argmax}_{k,l} I_{\text{merge}(r,k,l)}(D,E).$$
- 3. Set new $(r,C) = \text{merge}(r,k,l)$.
- 4. Repeat 2 and 3 until $|C|$ reaches predetermined size.

Brown's Classes: Programming Tips & Tricks

Complexity Issues

Still too complex:

- $|V|$ iterations of the steps 2 and 3.
- $|V|^2$ steps to maximize $\text{argmax}_{k,l}$ (selecting k,l freely from $|C|$, which is in the order of $|V|^2$)
- $|V|^2$ steps to compute $I(D,E)$ (sum within sum, all classes, also: includes log)
⇒ total: $|V|^5$
- i.e., for $|V| = 100$, about 10^{10} steps (several hours!)
- but $|V| \sim 50,000$ or more

Formula breakdown

- Mutual Information at k^{th} iteration (= k classes):
 - $I_k = \sum_{l,r \in C} p_k(l,r) \log(p_k(l,r) / (p_{kl}(l) p_{kr}(r)))$
- For each pair of classes at iteration k , we define:
 - $q_k(l,r) = p_k(l,r) \log(p_k(l,r) / (p_{kl}(l) p_{kr}(r)))$
- So:
 - $I_k = \sum_{l,r \in C} q_k(l,r)$
- $q_k(l,r)$ using bigram counts $c_k(l,r)$:
$$q_k(l,r) = c_k(l,r)/N \log(N c_k(l,r) / (c_{kl}(l) c_{kr}(r)))$$

$l \setminus r$	c_1	c_2	c_3	c_4
c_1	10	2	0	1
c_2	0	0	5	2
c_3	0	2	0	3
c_4	2	3	0	0

unigram/marginal counts

Trick #1: Recomputing MI the Smart Way

- For test-merging c_2 and c_4 we recompute only rows/columns 2 & 4:

k: $l \setminus r$ | c_1 | c_2 | c_3 | c_4

	c_1	c_2	c_3	c_4
c_1	10	2	0	1
c_2	0	0	5	2
c_3	0	2	0	3
c_4	2	3	0	0

The diagram shows a 5x5 matrix representing a confusion table. The columns are labeled $l \setminus r$, c_1 , c_2 , c_3 , and c_4 . The rows are labeled c_1 , c_2 , c_3 , and c_4 . The matrix contains the following values:
Row c_1 : 10, 2, 0, 1
Row c_2 : 0, 0, 5, 2
Row c_3 : 0, 2, 0, 3
Row c_4 : 2, 3, 0, 0
Arrows from the bottom of each row c_2 , c_3 , and c_4 point to the right edge of the matrix, indicating that only these columns need to be recomputed during the merge process.

Trick #1: Recomputing MI the Smart Way

- For test-merging c_2 and c_4 we recompute only rows/columns 2 & 4:
 - Subtract column/row (2 & 4) from the MI sum:

1 \ r	c_1	c_2	c_3	c_4
c_1	10	2	0	1
c_2	0	0	5	2
c_3	0	2	0	3
c_4	2	3	0	0

The diagram illustrates the step of subtracting row c_2 and column c_4 from the MI sum. Three arrows point from the row and column totals to the corresponding row and column in the matrix. One arrow points from the total row value 10 to the cell containing 2. Another arrow points from the total column value 1 to the cell containing 1. A third arrow points from the total column value 2 to the cell containing 2.

Trick #1: Recomputing MI the Smart Way

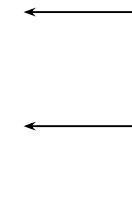
- For test-merging c_2 and c_4 we recompute only rows/columns 2 & 4:

- Subtract column/row (2 & 4) from the MI sum:

(be careful at the intersections)

k:

1 \ r	c_1	c_2	c_3	c_4
c_1	10	2	0	1
c_2	0	0	5	2
c_3	0	2	0	3
c_4	2	3	0	0



Trick #1: Recomputing MI the Smart Way

- For test-merging c_2 and c_4 we recompute only rows/columns 2 & 4:

- Subtract column/row (2 & 4) from the MI sum:

(be careful at the intersections)

k:

1 \ r	c_1	c_2	c_3	c_4
c_1	10	2	0	1
c_2	0	0	5	2
c_3	0	2	0	3
c_4	2	3	0	0




- Add sums of the merged counts (row & column for (c_2' is the merged class):

(watch the intersection again)

k-1:

1 \ r	c_1	c_2'	c_3
c_1	10	3	0
c_2'	2	5	5
c_3	0	5	0



Trick #2: Precompute the Counts-to-be-Subtracted

- Summing loop goes through i, j
 - ... but the single row/column sums do not depend on the (resulting sums after the) merge \Rightarrow can be precomputed
 - only $2k$ logs to compute at each algorithm iteration, instead of k^2
- Then for each “merge-to-be” compute only add-on sums, plus “intersection adjustment”

Formulas for Tricks #1 and #2

- Recap:

$$q_k(l,r) = p_k(l,r) \log(p_k(l,r) / (p_{kl}(l) p_{kr}(r)))$$

the same, but using counts:

$$q_k(l,r) = c_k(l,r)/N \log(N c_k(l,r)/(c_{kl}(l) c_{kr}(r)))$$

- Define further (row+column a sum):

$$s_k(a) = \sum_{l=1..k} q_k(l,a) + \sum_{r=1..k} q_k(a,r) - q_k(a,a)$$

- Then, the subtraction part of Trick #1 amounts to

$$\text{sub}_k(a,b) = s_k(a) + s_k(b) - q_k(a,b) - q_k(b,a)$$

precomputed

Intersection adjustment

	c ₂	c ₃	c ₄
c ₂	0	5	2
c ₃	2	0	3
c ₄	3	0	0

remaining intersection adjustment

Formulas - cont.

- After-merge add-on:

$$\text{add}_k(a,b) = \sum_{l=1..k, l \neq a,b} q_k(l,a+b) + \sum_{r=1..k, r \neq a,b} q_k(a+b,r) + q_k(a+b,a+b)$$

- a+b is the new (merged) class

- Hint: use the definition of q_k as a “macro”, and then:

$$p_k(a+b,r) = p_k(a,r) + p_k(b,r) \quad (\text{same for other sums, equivalent})$$

- The above sums cannot be precomputed
- Mutual Information after merge of class a,b:
 - $I_{k-1}(a,b) = I_k - \text{sub}_k(a,b) + \text{add}_k(a,b)$
 - I_k is the “old” MI, kept from previous iteration of the algorithm

Trick #3: Ignore Zero Counts

- Many bigrams are 0
 - e.g. in the Canadian Hansards corpus, < 0.1 % of bigrams are non-zero)
- Consider non-zero bigrams only:
 - e.g. create linked lists of non-zero counts in columns and rows
 - similar effect: use hashes (store non-zero-count bigrams)
- Update links after merge (after step 3)

Trick #4: Use Updated Loss of MI

- We are now down to $|V|^4$: $|V|$ merges, each merge takes $|V|^2$ “test-merges”, each test-merge involves order-of- $|V|$ operations ($\text{add}_k(i,j)$ term, slide 34)
- Observation:
 - many numbers (s_k, q_k) needed to compute the mutual information loss due to a merge of $i+j$ **do not change**: namely, those which are not in the vicinity of either i nor j .
- Idea:
 - keep the **MI loss matrix** for all pairs of classes, and (after a merge) update only those cells which have been influenced by the merge.

Formulas for Trick #4 (s_{k-1}, L_{k-1})

- Keep a matrix of "losses" $L_k(d,e)$ [symmetry: $L_k(d,e) = L_k(e,d)$]
- Init: $L_k(d,e) = \text{sub}_k(d,e) - \text{add}_k(d,e)$ [then $I_{k-1}(d,e) = I_k - L_k(d,e)$]
- Suppose a,b are now the two classes merged into a
- Update (k-1: index used for the next iteration; $i,j \neq a,b$):

$$s_{k-1}(i) = s_k(i) - q_k(i,a) - q_k(a,i) - q_k(i,b) - q_k(b,i) + q_{k-1}(a,i) + q_{k-1}(i,a)$$

$$\begin{aligned} L_{k-1}(i,j) = & L_k(i,j) - s_k(i) + s_{k-1}(i) - s_k(j) + s_{k-1}(j) + \\ & + q_k(i+j,a) + q_k(a,i+j) + q_k(i+j,b) + q_k(b,i+j) - \\ & - q_{k-1}(i+j,a) - q_{k-1}(a,i+j) \end{aligned}$$

Completing Trick #4

- $s_{k-1}(a)$ must be computed using the “Init” sum (see the prev. slide).
- $L_{k-1}(a,i) = L_{k-1}(i,a)$ must be computed in a similar way, for all $i \neq a,b$.
- $s_{k-1}(b), L_{k-1}(b,i), L_{k-1}(i,b)$ are not needed anymore (keep track of such data, i.e. mark every class already merged into some other class and do not use it anymore).
- Keep track of the minimal loss during the $L_k(i,j)$ update process (so that the next merge to be taken is obvious immediately after finishing the update step).

Efficient Implementation

Data Structures: (N - # of bigrams in data [fixed])

- $\text{Hist}(k)$ - history of merges
 $\text{Hist}(k) = (a,b)$ merged when the remaining number of classes was k
- $c_k(i,j)$ - bigram class counts [updated after merge]
- $c_{kl}(i), c_{kr}(i)$ - unigram (marginal) counts [updated]
- $L_k(a,b)$ - table of losses; upper-right triangle [updated]
- $s_k(a)$ - “subtraction” subterms [optionally updated]
- $q_k(i,j)$ - subterms involving a log [optionally updated]

The optionally updated data structures will give linear improvement only in the subsequent steps, but at least $s_k(i)$ is necessary in the initialization phase (1st iteration)

Implementation: the Initialization Phase

1. Read data in, set $k=|V|$, init counts $c_k(l,r)$; then $\forall l,r,a,b; a < b$:
2. Init unigram counts $c_{kl}(l), c_{kr}(r)$:

$$c_{kl}(l) = \sum_{r=1..k} c_k(l,r), \quad c_{kr}(r) = \sum_{l=1..k} c_k(l,r)$$

[must take care of start & end of data!]

3. Init $q_k(l,r)$: use the 2nd formula (count-based) on slide 27,

$$q_k(l,r) = c_k(l,r)/N \log(N c_k(l,r)/(c_{kl}(l) c_{kr}(r)))$$

4. Init $s_k(a) = \sum_{l=1..k} q_k(l,a) + \sum_{r=1..k} q_k(a,r) - q_k(a,a)$

5. Init $L_k(a,b) = s_k(a) + s_k(b) - q_k(a,b) - q_k(b,a) - q_k(a+b,a+b) +$
 $- \sum_{l=1..k, l \neq a,b} q_k(l,a+b) - \sum_{r=1..k, r \neq a,b} q_k(a+b,r)$

Implementation: Select & Update

6. Select the best pair (a,b) to merge into a
watch the candidates when computing $L_k(a,b)$; save to $\text{Hist}(k)$
7. Optionally, update $q_k(i,j)$ for all $i,j \neq b$, get $q_{k-1}(i,j)$
remember those $q_k(i,j)$ values needed for the updates below
8. Optionally, update $s_k(i)$ for all $i \neq b$, to get $s_{k-1}(i)$
again, remember the $s_k(i)$ values for the “loss table” update
9. Update the loss table, $L_k(i,j)$, to $L_{k-1}(i,j)$, using the tabulated q_k ,
 q_{k-1} , s_k and s_{k-1} values, or compute the needed $q_k(i,j)$ and $q_{k-1}(i,j)$
values dynamically from the counts:


$$c_k(i+j,b) = c_k(i,b) + c_k(j,b); \quad c_{k-1}(a,i) = c_k(a+b,i)$$

Towards the Next Iteration

10. During the $L_k(i,j)$ update, keep track of the minimal loss of MI, and the two classes which caused it.
11. Remember such best merge in $\text{Hist}(k)$.
12. Get rid of all s_k , q_k , L_k values.
13. Set $k = k - 1$; stop if $k = 1$.
14. Start the next iteration
 - either by the optional updates (steps 7 and 8), or
 - directly updating $L_k(i,j)$ again (step 9).

Using the Hierarchy


- Natural Form of Classes
- follows from the sequence of merges:



evaluation assessment analysis understanding opinion

Numbering the Classes (within the Hierarchy)

- Binary branching
- Assign 0/1 to the left/right branch at every node:



Word Classes in Applications

- Even in the era of neural embeddings, Brown classes have modern, practical applications (such as POS tagging)
- Especially useful in low-resource and domain-specific scenarios (tens of millions words)
- Provide compact, interpretable word classes that can replace sparse one-hot or n-gram features.

Moodle Quiz

Moodle Quiz



<https://dl1.cuni.cz/course/view.php?id=18547>