

Statistical Methods in Natural Language Processing

7. Word Representations

Pavel Pecina, Jindřich Helcl

2 December, 2025

Course Segments

1. Introduction, probability, essential information theory
2. Statistical language modelling (n-gram)
3. Statistical properties of words
4. Word representations
5. Hidden Markov models, Tagging

Recap from Last Week

The New Model

Rewrite the n-gram LM using classes:

- Original definition: [$k = 1 \dots n$]

$$p_k(w_i|h_i) = c(h_i, w_i) / c(h_i) \quad [\text{history: } (k-1) \text{ words}]$$

- Introduce classes:

$$p_k(w_i|h_i) = p(w_i|c_i) p_k(c_i|h_i)$$

- history: classes, too: [for trigram: $h_i = c_{i-2}, c_{i-1}$, bigram: $h_i = c_{i-1}$]
- Smoothing as usual
 - over $p_k(w_i|h_i)$, where each is defined as above (except uniform which is $1/|V|$)

Creating the Word-to-Class Map (Brown's Classes)

- Consider bigram model for now.
- Bigram estimate:

$$p_2(c_i|h_i) = p_2(c_i|c_{i-1}) = c(c_{i-1}, c_i) / c(c_{i-1}) = c(r(w_{i-1}), r(w_i)) / c(r(w_{i-1}))$$

- Form of the model:
 - just raw bigram for now:
- Maximize over r (given $r \rightarrow$ fixed p, p_2):
 - define objective

$$L(r) = 1/|T| \sum_{i=1..|T|} \log(p(w_i|r(w_i)) p_2(r(w_i))|r(w_{i-1})))$$

$$r_{\text{best}} = \operatorname{argmax}_r L(r) \quad (L(r) = \text{norm. logprob of training data ... as usual})$$

Simplifying the Objective Function

- Start from $L(r) = 1/|T| \sum_{i=1..|T|} \log(p(w_i|r(w_i))) p_2(r(w_i)|r(w_{i-1}))$:
 $1/|T| \sum_{i=1..|T|} \log(p(w_i|r(w_i)) \underline{p(r(w_i))} p_2(r(w_i)|r(w_{i-1})) / \underline{p(r(w_i))}) =$
 $1/|T| \sum_{i=1..|T|} \log(\underline{p(w_i,r(w_i))} p_2(r(w_i)|r(w_{i-1})) / p(r(w_i))) =$
 $1/|T| \sum_{i=1..|T|} \log(\underline{p(w_i)}) + 1/|T| \sum_{i=1..|T|} \log(p_2(r(w_i)|r(w_{i-1})) / p(r(w_i))) =$
 $-H(W) + 1/|T| \sum_{i=1..|T|} \log(p_2(r(w_i)|r(w_{i-1})) \underline{p(r(w_{i-1}))} / (\underline{p(r(w_{i-1}))} p(r(w_i))))$
=
- $-H(W) + 1/|T| \sum_{i=1..|T|} \log(\underline{p(r(w_i),r(w_{i-1}))} / (\underline{p(r(w_{i-1}))} p(r(w_i)))) =$
 $-H(W) + \sum_{d,e \in C} p(d,e) \log(p(d,e) / (p(d) p(e))) =$
 $-H(W) + I(D,E)$ (event E picks class adjacent (to the right) to the one picked by D)
- Since W does not depend on r, we ended up with maximizing I(D,E)

The Greedy Algorithm

- Define merging operation on the mapping $r: V \rightarrow C$:
 - merge: $R \times C \times C \rightarrow R' \times C-1: (r,k,l) \rightarrow r',C'$ such that
 - $C^{-1} = \{C - \{k,l\} \cup \{m\}\}$ (throw out k and l, add new m $\notin C$)
 - $r'(w) = m \quad \text{for } w \in r_{INV}(\{k,l\}),$
 $r(w) \quad \text{otherwise.}$
- 1. Start with each word in its own class ($C = V$), $r = \text{id}$.
- 2. Merge two classes k,l into one, m , such that
$$(k,l) = \operatorname{argmax}_{k,l} I_{\text{merge}(r,k,l)}(D,E).$$
- 3. Set new $(r,C) = \text{merge}(r,k,l)$.
- 4. Repeat 2 and 3 until $|C|$ reaches predetermined size.

Complexity Issues

Still too complex:

- $|V|$ iterations of the steps 2 and 3.
- $|V|^2$ steps to maximize $\text{argmax}_{k,l}$ (selecting k,l freely from $|C|$, which is in the order of $|V|^2$)
- $|V|^2$ steps to compute $I(D,E)$ (sum within sum, all classes, also: includes log)
⇒ total: $|V|^5$
- i.e., for $|V| = 100$, about 10^{10} steps (several hours!)
- but $|V| \sim 50,000$ or more

Formula breakdown

- Mutual Information at k^{th} iteration (= k classes):
 - $I_k = \sum_{l,r \in C} p_k(l,r) \log(p_k(l,r) / (p_{kl}(l) p_{kr}(r)))$
- For each pair of classes at iteration k , we define:
 - $q_k(l,r) = p_k(l,r) \log(p_k(l,r) / (p_{kl}(l) p_{kr}(r)))$
- So:
 - $I_k = \sum_{l,r \in C} q_k(l,r)$
- $q_k(l,r)$ using bigram counts $c_k(l,r)$:
$$q_k(l,r) = c_k(l,r)/N \log(N c_k(l,r) / (c_{kl}(l) c_{kr}(r)))$$

$l \setminus r$	c_1	c_2	c_3	c_4
c_1	10	2	0	1
c_2	0	0	5	2
c_3	0	2	0	3
c_4	2	3	0	0

unigram/marginal counts

Trick #1: Recomputing MI the Smart Way

- For test-merging c_2 and c_4 we recompute only rows/columns 2 & 4:

- Subtract column/row (2 & 4) from the MI sum:

(be careful at the intersections)

- add sums of the merged counts (row & column for (c_2' is the merged class):

(watch the intersection again)

$l \setminus r$	c_1	c_2	c_3	c_4
c_1	10	2	0	1
c_2	0	0	5	2
c_3	0	2	0	3
c_4	2	3	0	0

↑ ↑

$l \setminus r$	c_1	c_2'	c_3
c_1	10	3	0
c_2'	2	5	5
c_3	0	5	0

↑

Trick #2: Precompute the Counts-to-be-Subtracted

- Summing loop goes through i, j
 - ... but the single row/column sums do not depend on the (resulting sums after the) merge \Rightarrow can be precomputed
 - only $2k$ logs to compute at each algorithm iteration, instead of k^2
- Then for each “merge-to-be” compute only add-on sums, plus “intersection adjustment”

Formulas for Tricks #1 and #2

- Recap:

$$q_k(l,r) = p_k(l,r) \log(p_k(l,r) / (p_{kl}(l) p_{kr}(r)))$$

the same, but using counts:

$$q_k(l,r) = c_k(l,r)/N \log(N c_k(l,r)/(c_{kl}(l) c_{kr}(r)))$$

- Define further (row+column a sum):

$$s_k(a) = \sum_{l=1..k} q_k(l,a) + \sum_{r=1..k} q_k(a,r) - q_k(a,a)$$

- Then, the subtraction part of Trick #1 amounts to

$$\text{sub}_k(a,b) = s_k(a) + s_k(b) - q_k(a,b) - q_k(b,a)$$

precomputed

Intersection adjustment

	c ₂	c ₃	c ₄
c ₂	0	5	2
c ₃	2	0	3
c ₄	3	0	0

remaining intersection adjustment

Formulas - cont.

- After-merge add-on:

$$\text{add}_k(a,b) = \sum_{l=1..k, l \neq a,b} q_k(l,a+b) + \sum_{r=1..k, r \neq a,b} q_k(a+b,r) + q_k(a+b,a+b)$$

- a+b is the new (merged) class

- Hint: use the definition of q_k as a “macro”, and then:

$$p_k(a+b,r) = p_k(a,r) + p_k(b,r) \quad (\text{same for other sums, equivalent})$$

- The above sums cannot be precomputed
- Mutual Information after merge of class a,b:
 - $I_{k-1}(a,b) = I_k - \text{sub}_k(a,b) + \text{add}_k(a,b)$
 - I_k is the “old” MI, kept from previous iteration of the algorithm

Trick #3: Ignore Zero Counts

- Many bigrams are 0
 - e.g. in the Canadian Hansards corpus, < .1 % of bigrams are non-zero)
- Consider non-zero bigrams only:
 - e.g. create linked lists of non-zero counts in columns and rows
 - similar effect: use hashes (store non-zero-count bigrams)
- Update links after merge (after step 3)

Trick #4: Use Updated Loss of MI

- We are now down to $|V|^4$: $|V|$ merges, each merge takes $|V|^2$ “test-merges”, each test-merge involves order-of- $|V|$ operations ($\text{add}_k(i,j)$ term, slide 34)
- Observation:
 - many numbers (s_k, q_k) needed to compute the mutual information loss due to a merge of $i+j$ **do not change**: namely, those which are not in the vicinity of either i nor j .
- Idea:
 - keep the MI loss matrix for all pairs of classes, and (after a merge) update only those cells which have been influenced by the merge.

Formulas for Trick #4 (s_{k-1}, L_{k-1})

- Keep a matrix of "losses" $L_k(d,e)$ [symmetry: $L_k(d,e) = L_k(e,d)$]
- Init: $L_k(d,e) = \text{sub}_k(d,e) - \text{add}_k(d,e)$ [then $I_{k-1}(d,e) = I_k - L_k(d,e)$]
- Suppose a,b are now the two classes merged into a
- Update (k-1: index used for the next iteration; $i,j \neq a,b$):

$$s_{k-1}(i) = s_k(i) - q_k(i,a) - q_k(a,i) - q_k(i,b) - q_k(b,i) + q_{k-1}(a,i) + q_{k-1}(i,a)$$

$$\begin{aligned} L_{k-1}(i,j) = & L_k(i,j) - s_k(i) + s_{k-1}(i) - s_k(j) + s_{k-1}(j) + \\ & + q_k(i+j,a) + q_k(a,i+j) + q_k(i+j,b) + q_k(b,i+j) - \\ & - q_{k-1}(i+j,a) - q_{k-1}(a,i+j) \end{aligned}$$

Word Representations

Word Representations in n-gram LMs


Word representation in n-gram language models

- by themselves (strings) as dictionary keys (or integer indices to a vocabulary list)
- by class numbers
- context-independent, needs to be modeled with longer n-grams
- also independent on each other (does not capture semantic similarity between words/classes)

Word Representations in NLP Tasks

Many NLP tasks depend on good word representations.


- Example: part-of-speech tagging
 - for given word x , produce a part of speech tag y
 - using task-specific training data T_{pos} and test data S_{pos}



Word Representations in NLP Tasks

Many NLP tasks depend on good word representations.


- Example: part-of-speech tagging
 - for given word x , produce a part of speech tag y
 - using task-specific training data T_{pos} and test data S_{pos}
- OOV problem
 - S_{pos} contains words (or word/tag combinations) not present in T_{pos}
 - If we use T_{pos} with discrete word representations we cannot decide about OOVs



Word Representations in NLP Tasks

Many NLP tasks depend on good word representations.

- Example: part-of-speech tagging
 - for given word x , produce a part of speech tag y
 - using task-specific training data T_{pos} and test data S_{pos}
- OOV problem
 - S_{pos} contains words (or word/tag combinations) not present in T_{pos}
 - If we use T_{pos} with discrete word representations we cannot decide about OOVs
- Idea: Re-use word representations from a language model
 - trained on much larger data (larger than T_{pos}) with much larger vocabulary
 - how can this be done with our n-gram LMs?



Representations as the Product of LMs


Using word representations from n-gram LMs to solve POS tagging

- originally words are represented with dictionary keys
 - no direct comparability, no similarity modeled

Representations as the Product of LMs

Using word representations from n-gram LMs to solve POS tagging


- originally words are represented with dictionary keys
 - no direct comparability, no similarity modeled
- representing a word by its class
 - we can represent words that were not seen in T_{pos} !
 - to train a POS classifier, we can pre-process the task-specific data T_{pos} and S_{pos} by replacing words by their class as given by a class n-gram model
 - works, but it is too coarse
 - no similarity between classes
 - words either belong or do not belong to one class



Beyond Word Classes

Capturing word similarities

- word classes divide the vocabulary into equivalence classes
 - words within a class considered similar
- real groups of similar words are more general
 - non-disjoint, not transitive, ambiguous ...
 - e.g. *book+pay* vs. *book+text*




Beyond Word Classes

Capturing word similarities

- word classes divide the vocabulary into equivalence classes
 - words within a class considered similar
- real groups of similar words are more general
 - non-disjoint, not transitive, ambiguous ...
 - e.g. $\{book, pay\}$ vs. $\{book, text\}$


Each word should be able to belong to more classes!



Word Embeddings

Capturing word similarities

- Previously: each word can belong to more than one class
 - stored as a boolean vector
 - a word either belongs to a class or not
 - enables looking at similarity of classes (in terms of words they contain)
 - but we still cannot distinguish words within the same (set of) class(es)




Word Embeddings

Capturing word similarities

- Previously: each word can belong to more than one class
 - stored as a boolean vector
 - a word either belongs to a class or not
 - enables looking at similarity of classes (in terms of words they contain)
 - but we still cannot distinguish words within the same (set of) class(es)

- Use real numbers instead of booleans?
 - stored as vector of floats
 - each word belong to all classes *to different extents*
 - fine-grained categorization of vocabulary
 - these vectors are called **word embeddings**



Embedding Matrix

With pre-defined vocabulary V , let $E \in \mathbb{R}^{|V| \times d}$ be an **embedding matrix**

- Each row of E corresponds to a word from the vocabulary
- d is the **embedding size** (or dimension)
 - a hyper-parameter and needs to be decided before obtaining the embeddings

hello	-1.2	0.3	5.0	...
the	0.1	3.6	2.2	...
of	3.3	-1.0	3.1	...
dog	-1.0	7.5	7.1	...

- OOV words?
 - Usually collapsed into a special `<oov>` token that will get its embedding

Obtaining Embeddings

Getting word embeddings

- implicitly in end-to-end scenario
 - in neural-network-based models with sufficient data volumes (such as tasks like machine translation or language modeling), embeddings are trained along with the rest of the network
- using off-the-shelf pre-trained embeddings
 - embeddings for many languages readily available (FastText, Hugging Face, ...)
 - download the embedding matrix and train your classifier on the task-specific data
- **learning** from data

Learning Word Embeddings - Word2vec

Ideas of the Word2vec model (Mikolov et al., 2013)

- Use language modeling data for obtaining word embeddings, then re-use word embeddings for a specific task (such as POS tagging)
- Keep the word embedding model simple
 - smaller model means we can use more data for the same cost
- Embedding of a word should be based on the word's context
 - “you shall know a word by the company it keeps”
 - i.e. the Distributional Hypothesis (see Lecture 5, slide 14)

Word2vec Model Architecture

Two-layer linear neural network

$$w = (0,0,0,1,0,0)$$

1. word embedding (here w is a one-hot representation of the input word)

$$h = w \times E$$


2. output projection with softmax

$$y = \text{softmax}(h \times W + b)$$

$E \in \mathbb{R}^{|V| \times d}$, $W \in \mathbb{R}^{d \times |V|}$, and $b \in \mathbb{R}^{|V|}$ are *trainable* parameters

- The output vector y can be represented as probability distribution over vocabulary: the probability of each word appearing in the context of the input word w .
- Note the shape of W (similar to E , only transposed): columns of W are called “output embeddings” or “context embeddings”

Word2vec Model Architecture



Word2vec Training

Training objective: minimizing cross-entropy loss

- cross-entropy: distance between *two* distributions
 - in our case, the predicted distribution (output of softmax), vs the “true” distribution, i.e. the one-hot vector representing the target word
- goal: find parameters E , W , b that minimize the cross-entropy
 - using gradient descent,
 - and a large training dataset to sample from

Word2vec Training

Training is done by sampling input and target word(s) from data.

- **CBOW** (continuous bag-of-words): For a given target word w , sum embeddings of the context words and predict w (one training example).
- **Skip-gram**: For a given input word w , predict words in its context (one training example per context word)




Image credit: Libovický et al., Slides for NPFL129 lecture 6, 2025

Word2vec Data Sampling (skip-gram case)

1. All human beings are born free and equal in dignity ... → (All, humans)
(All, beings)

2. All human beings are born free and equal in dignity ... → (human, All)
(human, beings)
(human, are)
(beings, All)

3. All human beings are born free and equal in dignity ... → (beings, human)
(beings, are)
(beings, born)
(are, human)

4. All human beings are born free and equal in dignity ... → (are, beings)
(are, born)
(are, free)

Word2vec Implementation

The *softmax* operation normalizes probabilities across the vocabulary.

- For vector $z = (z_1, z_2, \dots, z_n)$ the softmax component is $\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^n e^{z_j}}$
- For large $n=|V|$ this slows down the training (computing the norm. factor)
- Improvement: Use sigmoid instead of softmax
 - does not model the probability for all context words given an input
 - works as binary classification (0/1) for a given pair of input and context words
- The cross-entropy loss given (input, context) word pair becomes:

$$-\log \sigma(e^\top \cdot c)$$

where e and c are the input embedding and context output embedding respectively and

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Word2vec Implementation, cont.

Using sigmoid instead of softmax gives us:

$$-\log \sigma(e^\top \cdot c)$$

- This is only for word pairs that do occur together (positive samples)
- For better training we need to also train on negative ones
 - a process called *negative sampling*
- Taking into account pairs of words that *do not* occur together, we can complete the loss function:

$$L = -\log \sigma(e^\top \cdot c_p) - \sum_{c_n} \log(1 - \sigma(e^\top \cdot c_n))$$

Word2vec Implementation, cont.

Using sigmoid instead of softmax gives us:

$$-\log \sigma(e^\top \cdot c)$$

- This is only for word pairs that do occur together (positive samples)
- For better training we need to also train on negative ones
 - a process called *negative sampling*
- Taking into account pairs of words that *do not* occur together, we can complete the loss function:

$$L = -\log \sigma(e^\top \cdot c_p) - \sum_{c_n} \log(1 - \sigma(e^\top \cdot c_n))$$

loss from the positive example

Word2vec Implementation, cont.

Using sigmoid instead of softmax gives us:

$$-\log \sigma(e^\top \cdot c)$$

- This is only for word pairs that do occur together (positive samples)
- For better training we need to also train on negative ones
 - a process called *negative sampling*
- Taking into account pairs of words that *do not* occur together, we can complete the loss function:
loss from the negative example

$$L = -\log \sigma(e^\top \cdot c_p) - \sum_{c_n} \log(1 - \sigma(e^\top \cdot c_n))$$

loss from the positive example

Word2vec Implementation, cont.

Using sigmoid instead of softmax gives us:

$$-\log \sigma(e^\top \cdot c)$$

- This is only for word pairs that do occur together (positive samples)
 - For better training we need to also train on negative ones
 - a process
 - Taking into account that we sample more negative examples than positive, usually around 2-5
- do not occur together, we can loss from the negative example*

$$L = -\log \sigma(e^\top \cdot c_p) - \sum_{c_n} \log(1 - \sigma(e^\top \cdot c_n))$$

loss from the positive example

Word2vec Implementation, cont.

Using sigmoid instead of softmax gives us:

$$-\log \sigma(e^\top \cdot c)$$

- This is only for word pairs that do occur together (positive samples)
 - For better training we need to also train on negative ones
 - a process
 - Taking into account that we sample more negative examples than positive, usually around 2-5
- do not occur together, we can calculate loss from the negative example*


$$L = -\log \sigma(e^\top \cdot c_p) - \sum c_n \log(1 - \sigma(e^\top \cdot c_n))$$

loss from the positive example


note that the input word e stays the same

Word2vec Vector Properties I


Word2vec vectors seem to enable meaningful arithmetic operations



Male-Female



Verb tense



Country-Capital

Word2vec Vector Properties II

Word2vec vectors encode distributional semantics: similarity in meaning → proximity in vector space.




Image from <https://www.ruder.io/word-embeddings-1/> 43

Suitable Tasks

Examples of tasks where using pre-trained word embeddings helps:

- Text classification (spam detection, news category)
- Sequence labeling (part-of-speech tagging, named-entity recognition)
- Sentiment analysis
- Document clustering
- Information retrieval / search ranking

Notable Implementations

- word2vec itself probably only used in tutorials such as
<https://www.tensorflow.org/text/tutorials/word2vec>
- GloVe (<https://nlp.stanford.edu/projects/glove/>)
 - count global cooccurrences, i.e. not just within a small window
- fastText (<https://fasttext.cc/>)
 - takes into account subwords and characters - good for OOVs
 - provide both trained models for many languages and code to train your own
- contextual embeddings, such as BERT
 - produce different word embeddings in different contexts, using a (large) language model

Moodle Quiz

Moodle Quiz



<https://dl1.cuni.cz/course/view.php?id=18547>