

Statistical Methods in Natural Language Processing

8. Hidden Markov Models

Pavel Pecina, Jindřich Helcl

9 December, 2025

Course Segments

1. Introduction, probability, essential information theory
2. Statistical language modelling (n-gram)
3. Statistical properties of words
4. Word representations
5. Hidden Markov models, Tagging

Recap from Last Week

Embedding Matrix

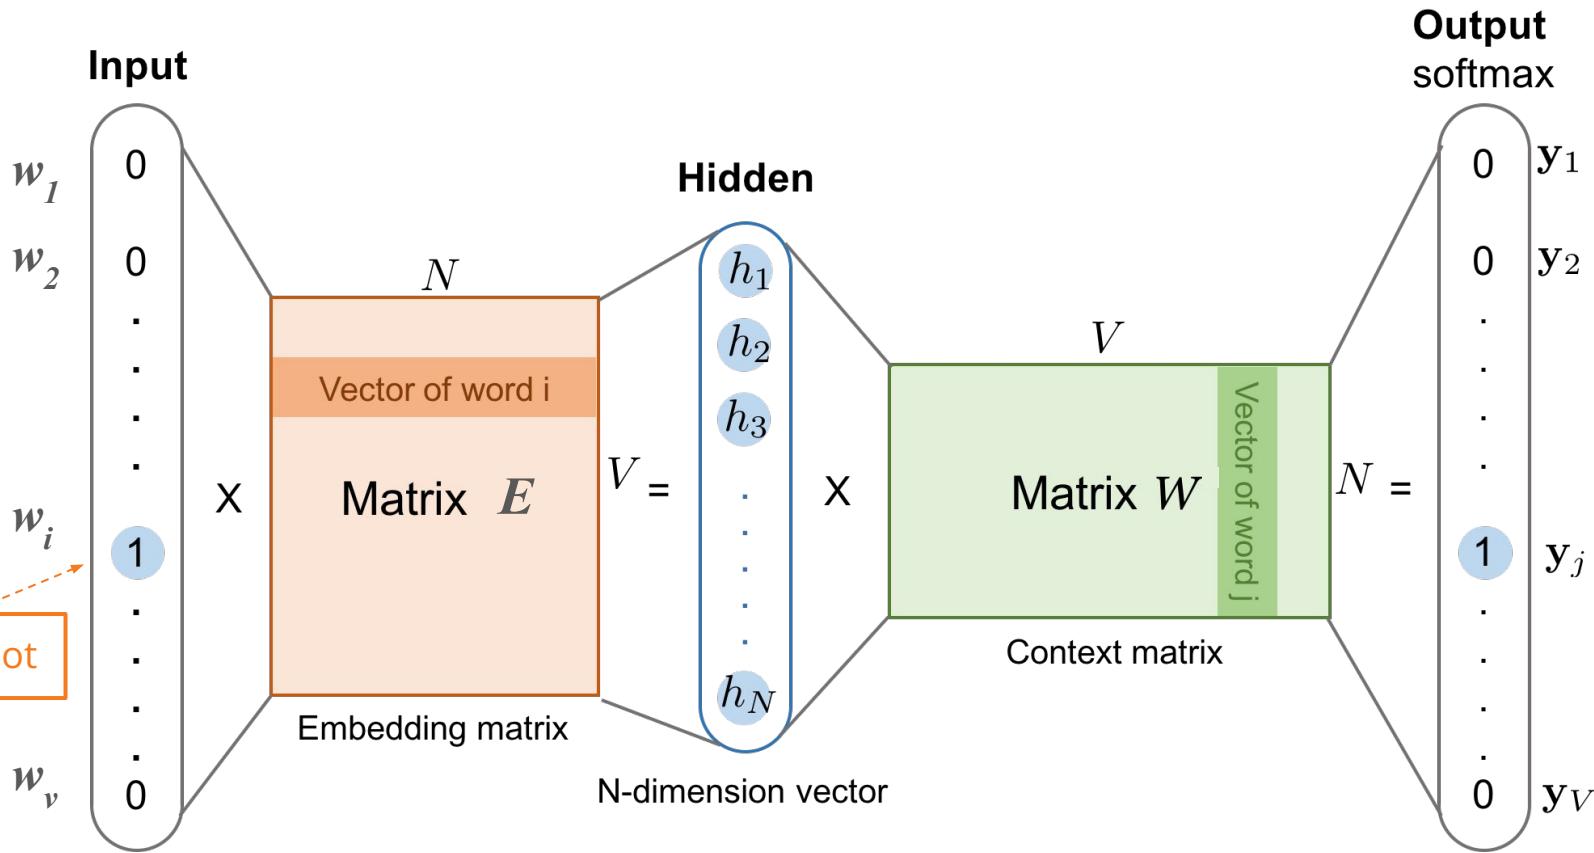
With pre-defined vocabulary V , let $E \in \mathbb{R}^{|V| \times d}$ be an **embedding matrix**

- Each row of E corresponds to a word from the vocabulary
- d is the **embedding size** (or dimension)
 - a hyper-parameter and needs to be decided before obtaining the embeddings

hello	-1.2	0.3	5.0	...
the	0.1	3.6	2.2	...
of	3.3	-1.0	3.1	...
dog	-1.0	7.5	7.1	...

- OOV words?
 - Usually collapsed into a special `<oov>` token that will get its embedding

Word2vec Model Architecture



Word2vec Training

Training is done by sampling input and target word(s) from data.

- **CBOW** (continuous bag-of-words): For a given target word w , sum embeddings of the context words and predict w (one training example).
- **Skip-gram**: For a given input word w , predict words in its context (one training example per context word)

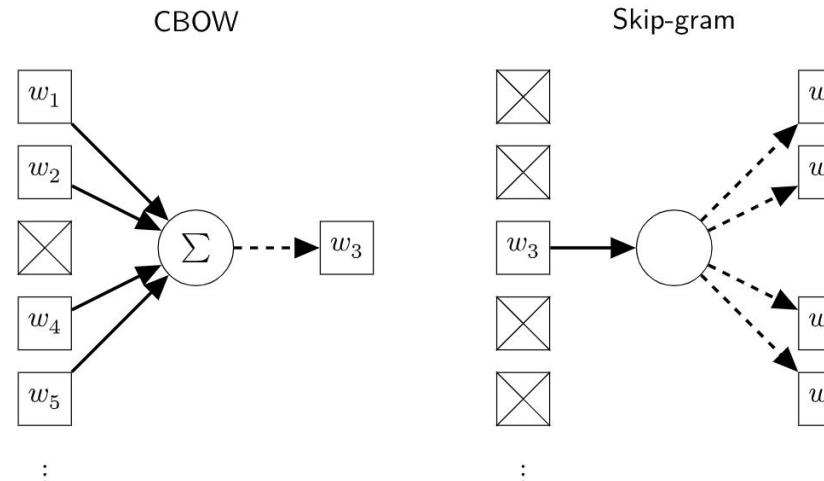


Image credit: Libovický et al., Slides for NPFL129 lecture 6, 2025

Word2vec Implementation, cont.

Using sigmoid instead of softmax gives us:

$$-\log \sigma(e^\top \cdot c)$$

- This is only for word pairs that do occur together (positive samples)
 - For better training we need to also train on negative ones
 - a process
 - Taking into account that we sample more negative examples than positive, usually around 2-5
- do not occur together, we can loss from the negative example*

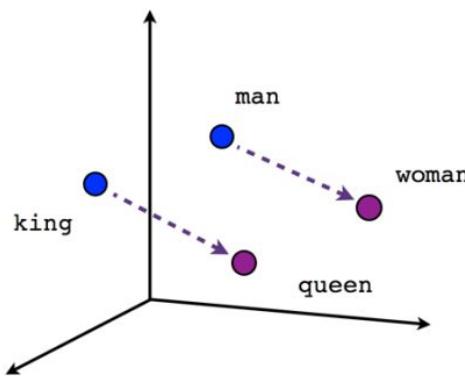
$$L = -\log \sigma(e^\top \cdot c_p) - \sum c_n \log(1 - \sigma(e^\top \cdot c_n))$$

loss from the positive example

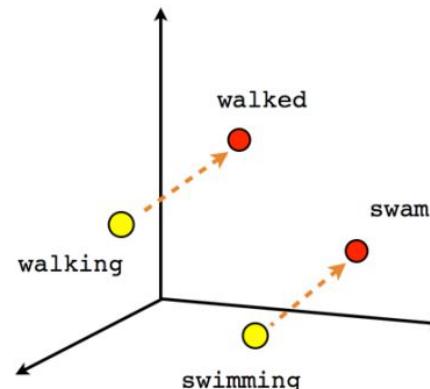
note that the input word e stays the same

Word2vec Vector Properties I

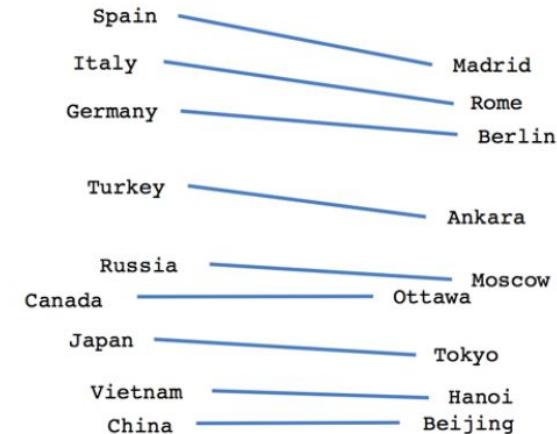
Word2vec vectors seem to enable meaningful arithmetic operations



Male-Female



Verb tense



Country-Capital

Word2vec Vector Properties II

Word2vec vectors encode distributional semantics: similarity in meaning → proximity in vector space.

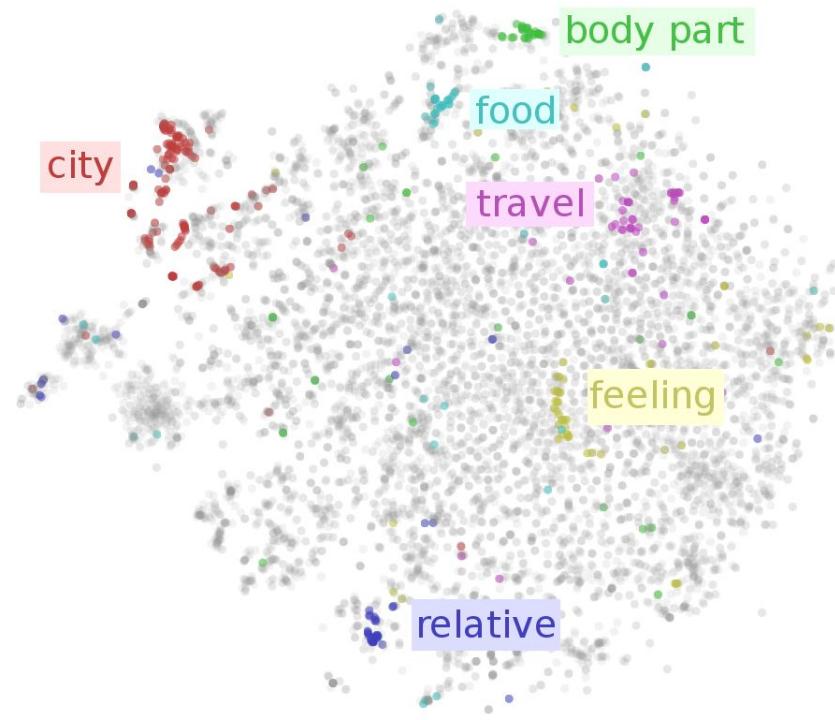


Image from <https://www.ruder.io/word-embeddings-1/>

Hidden Markov Models

Markov Process (Review)

- Bayes formula (chain rule):

$$P(W) = P(w_1, w_2, \dots, w_T) = \prod_{i=1..T} p(w_i | w_1, w_2, \dots, w_{i-n+1}, \dots, w_{i-1})$$

Markov Process (Review)

- Bayes formula (chain rule):

$$P(W) = P(w_1, w_2, \dots, w_T) = \prod_{i=1..T} p(w_i | \cancel{w_1, w_2, \dots, w_{i-1}}) w_{i-n+1}, \dots, w_{i-1})$$

- n-gram language models:

- Markov process (chain) of the order $n-1$:

$$P(W) = P(w_1, w_2, \dots, w_T) = \prod_{i=1..T} p(w_i | w_{i-n+1}, w_{i-n+2}, \dots, w_{i-1})$$



Markov Process (Review)

- Bayes formula (chain rule):

$$P(W) = P(w_1, w_2, \dots, w_T) = \prod_{i=1..T} p(w_i | \cancel{w_1}, \cancel{w_2}, \dots, \cancel{w_{i-n+1}}, \dots, \cancel{w_{i-1}})$$

- n-gram language models:

- Markov process (chain) of the order n-1:

$$P(W) = P(w_1, w_2, \dots, w_T) = \prod_{i=1..T} p(w_i | w_{i-n+1}, w_{i-n+2}, \dots, w_{i-1})$$

Using just one distribution (Ex.: trigram model: $p(w_i | w_{i-2}, w_{i-1})$):

Positions: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16

Words: My car **broke down**, and within hours Bob 's car **broke down**, too .

$$p(\text{,}|\text{broke down}) = p(w_5|w_3, w_4) = p(w_{14}|w_{12}, w_{13})$$

Markov Properties

- Markov Chain can generalize to any process (not just words):
 - Sequence of random variables: $X = (X_1, X_2, \dots, X_T)$
 - Sample space S (*states*), size N : $S = \{s_0, s_1, s_2, \dots, s_N\}$

Markov Properties

- Markov Chain can generalize to any process (not just words):
 - Sequence of random variables: $X = (X_1, X_2, \dots, X_T)$
 - Sample space S (*states*), size N : $S = \{s_0, s_1, s_2, \dots, s_N\}$
- Two properties
 1. Limited history (context, horizon):

$$\forall i \in 1..T; P(X_i | X_1, \dots, X_{i-1}) = P(X_i | X_{i-1})$$



Markov Properties

- Markov Chain can generalize to any process (not just words):
 - Sequence of random variables: $X = (X_1, X_2, \dots, X_T)$
 - Sample space S (*states*), size N : $S = \{s_0, s_1, s_2, \dots, s_N\}$
- Two properties

1. Limited history (context, horizon):

$$\forall i \in 1..T; P(X_i | X_1, \dots, X_{i-1}) = P(X_i | X_{i-1})$$



2. Time invariance (Markov Chain is stationary, homogeneous)

$$\forall i \in 1..T, \forall y, x \in S; P(X_i=y | X_{i-1}=x) = p(y|x)$$



ok ... same distribution

Long History Possible

- What if we want trigrams:

1 7 3 7 9 0 6 7 3 [4 5] [...]


Long History Possible

- What if we want trigrams:

1 7 3 7 9 0 6 7 3 

- Formally, use transformation:

- Define new variables Q_i , such that $X_i = \{Q_{i-1}, Q_i\}$
 - And then $P(X_i|X_{i-1}) = P(Q_{i-1}, Q_i|Q_{i-2}, Q_{i-1}) = P(Q_i|Q_{i-2}, Q_{i-1})$

Long History Possible

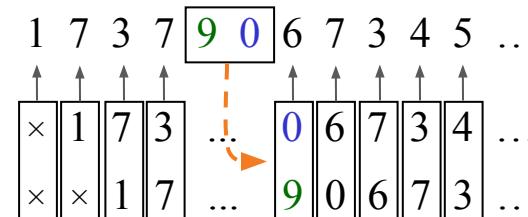
- What if we want trigrams:



- Formally, use transformation:

- Define new variables Q_i , such that $X_i = \{Q_{i-1}, Q_i\}$
- And then $P(X_i|X_{i-1}) = P(Q_{i-1}, Q_i|Q_{i-2}, Q_{i-1}) = P(Q_i|Q_{i-2}, Q_{i-1})$

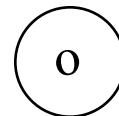
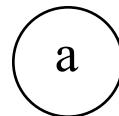
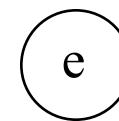
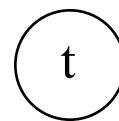
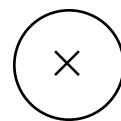
Predicting (X_i):



History ($X_{i-1} = \{Q_{i-2}, Q_{i-1}\}$):

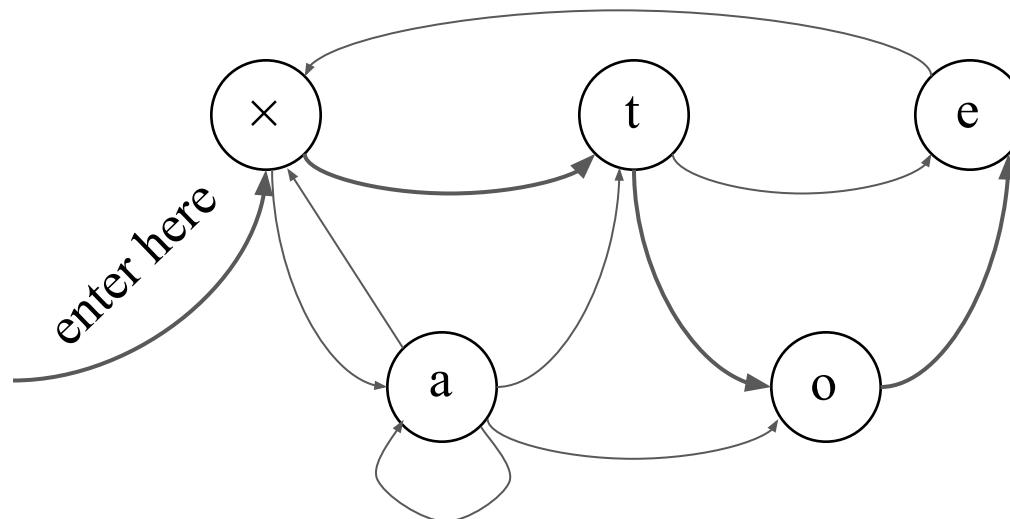
Graph Representation: State Diagram (Bigram case)

- Nodes: States, $S = \{s_0, s_1, s_2, \dots, s_N\}$



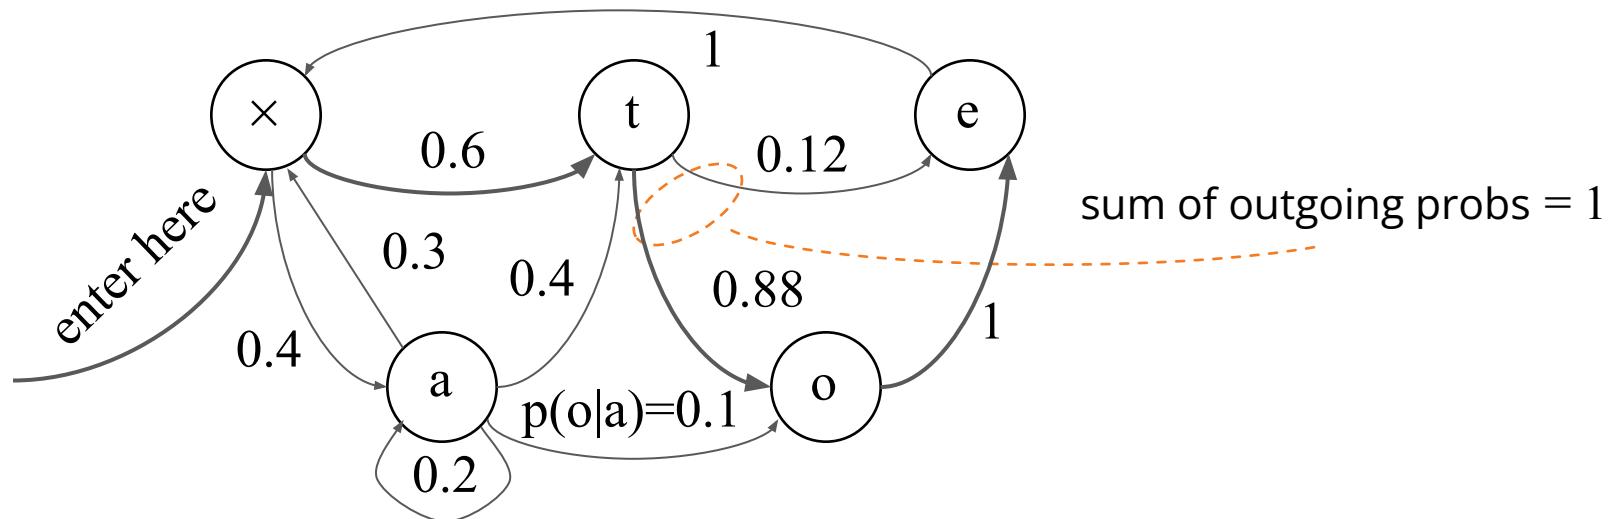
Graph Representation: State Diagram (Bigram case)

- Nodes: States, $S = \{s_0, s_1, s_2, \dots, s_N\}$
- Arcs: Transitions with probabilities, $P(X_i|X_{i-1})$, X_i generates s_i



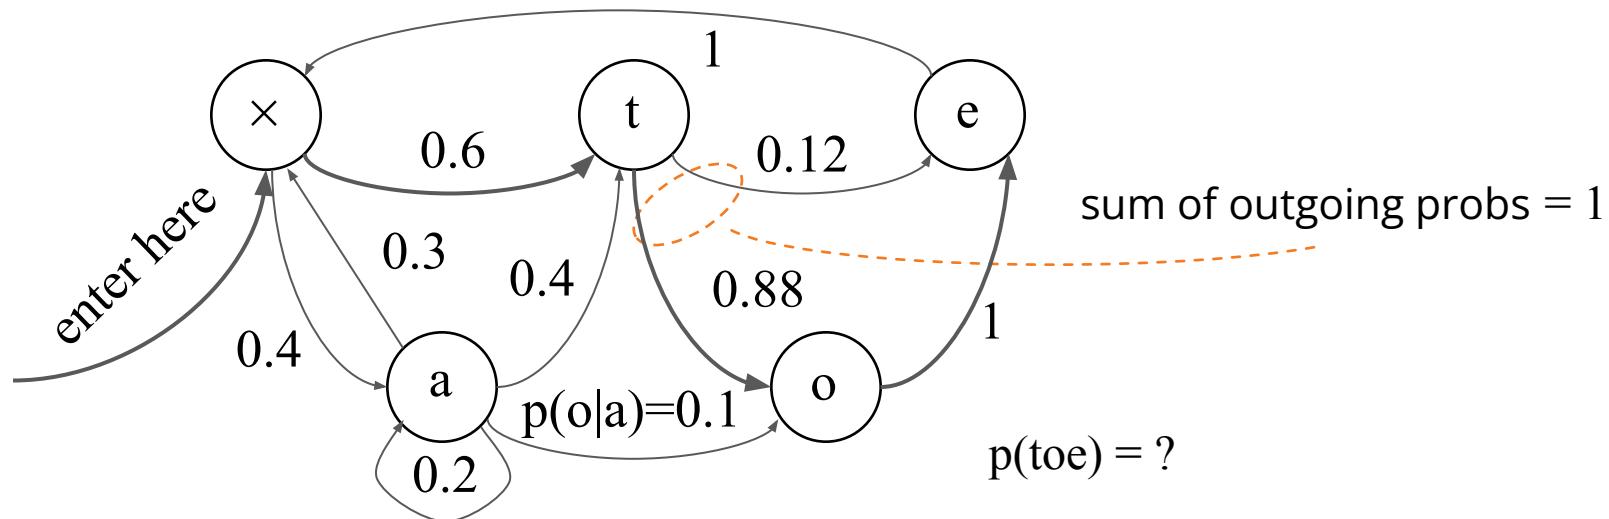
Graph Representation: State Diagram (Bigram case)

- Nodes: States, $S = \{s_0, s_1, s_2, \dots, s_N\}$
- Arcs: Transitions with probabilities, $P(X_i|X_{i-1})$, X_i generates s_i



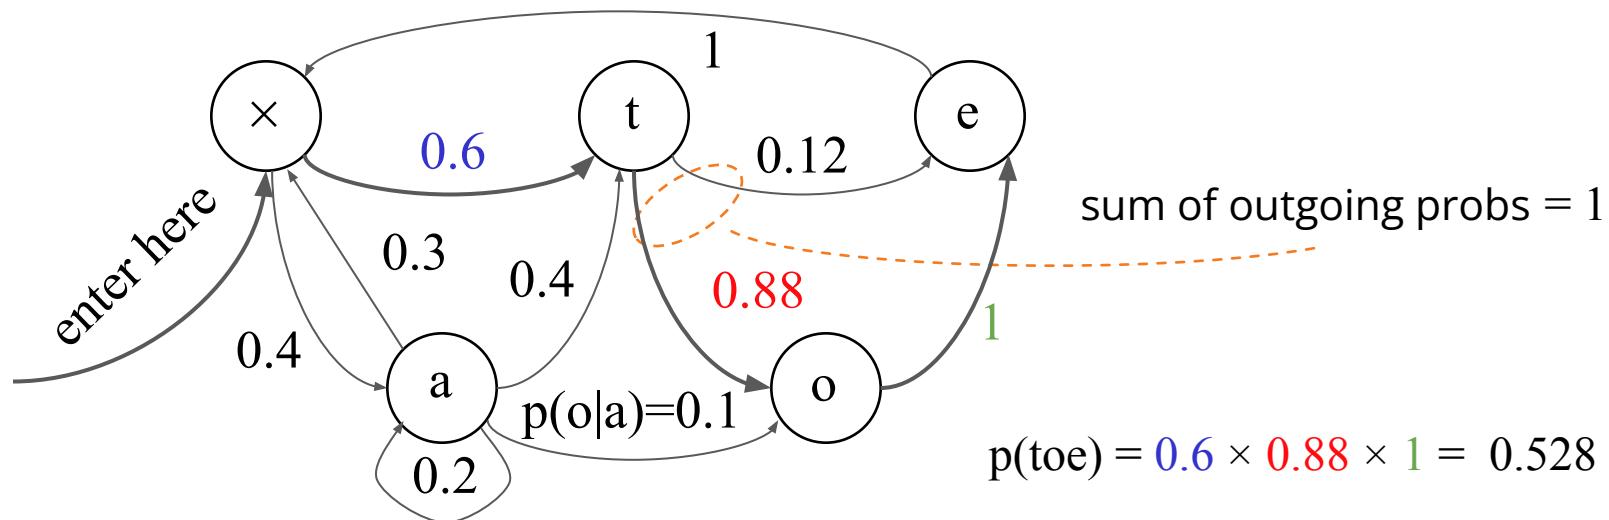
Graph Representation: State Diagram (Bigram case)

- Nodes: States, $S = \{s_0, s_1, s_2, \dots, s_N\}$
- Arcs: Transitions with probabilities, $P(X_i|X_{i-1})$, X_i generates s_i



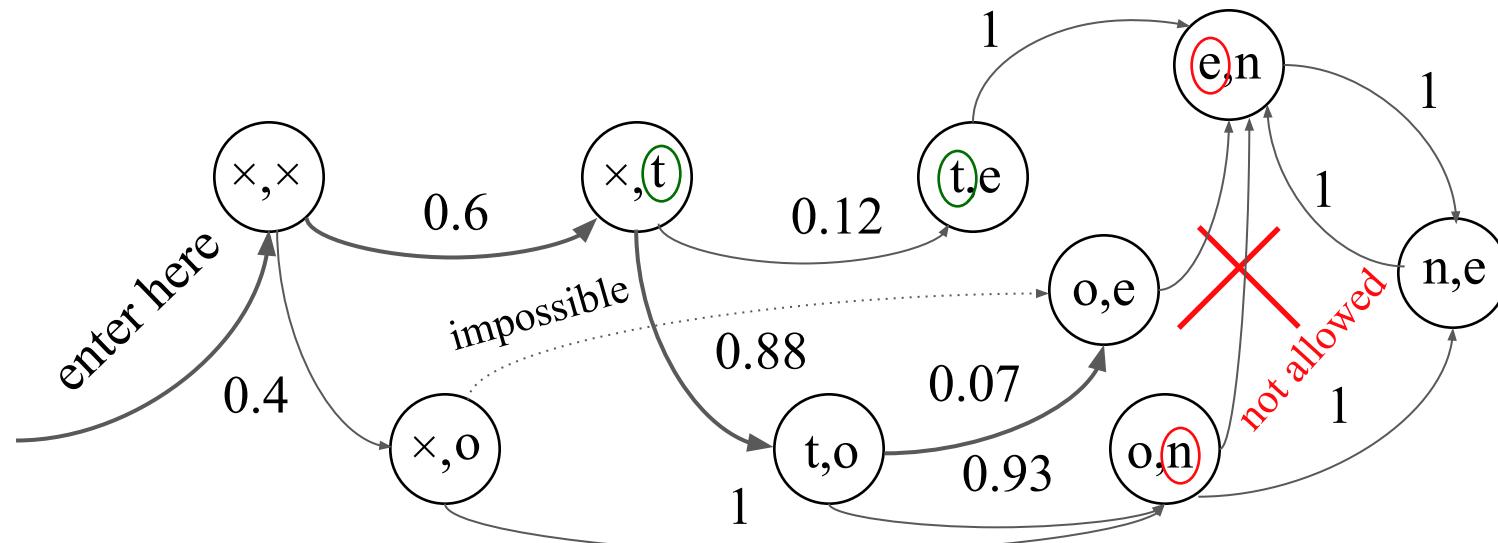
Graph Representation: State Diagram (Bigram case)

- Nodes: States, $S = \{s_0, s_1, s_2, \dots, s_N\}$
- Arcs: Transitions with probabilities, $P(X_i|X_{i-1})$, X_i generates s_i



The Trigram Case

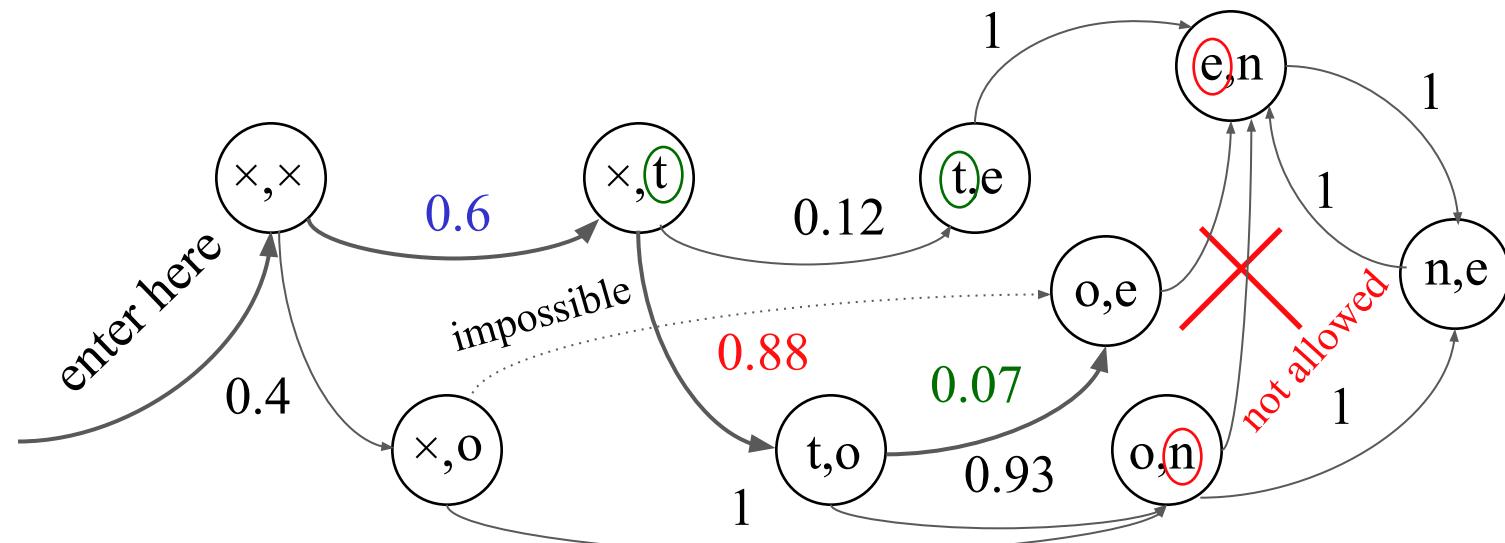
- Nodes: Pairs of states (s_k, s_l) , $S = \{s_0, s_1, s_2, \dots, s_N\}$
- Arcs: Transitions with probabilities, $P(X_i|X_{i-1})$, X_i generates (s_k, s_l)



$$p(\text{toe}) = ?$$

The Trigram Case

- Nodes: Pairs of states (s_k, s_l) , $S = \{s_0, s_1, s_2, \dots, s_N\}$
- Arcs: Transitions with probabilities, $P(X_i|X_{i-1})$, X_i generates (s_k, s_l)



$$p(\text{toe}) = 0.6 \times 0.88 \times 0.07 = 0.037$$

$$p(\text{one}) = ?$$

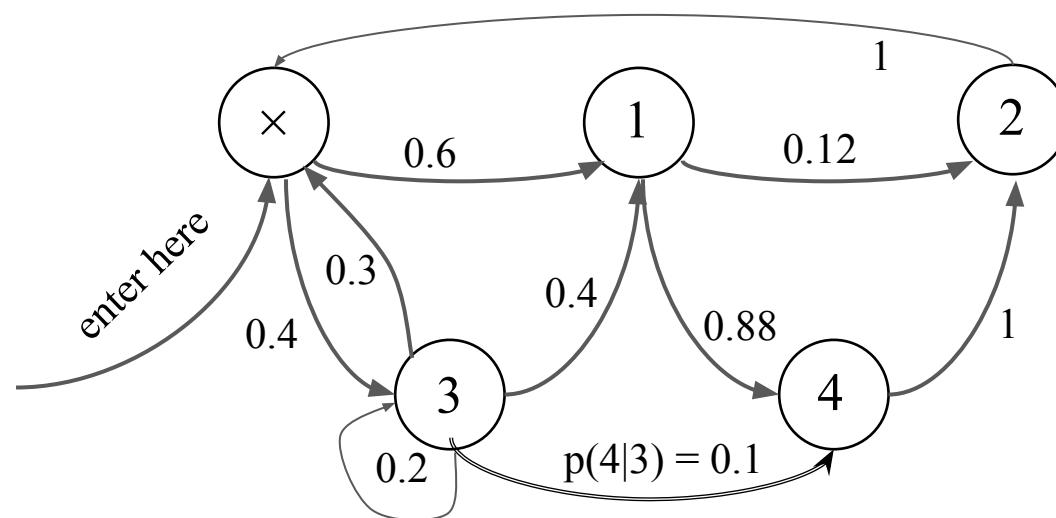
Finite State Automaton

- States ~ symbols of the [input / output] alphabet
 - pairs (or more): last element of the n-tuple
- Arcs ~ transitions (sequence of states)
 - Classical FSA: alphabet symbols on arcs:
 - possible transformation: arcs \leftrightarrow nodes
- Possible thanks to the “limited history” Markov Property
- So far: **Visible** Markov Models (VMM)



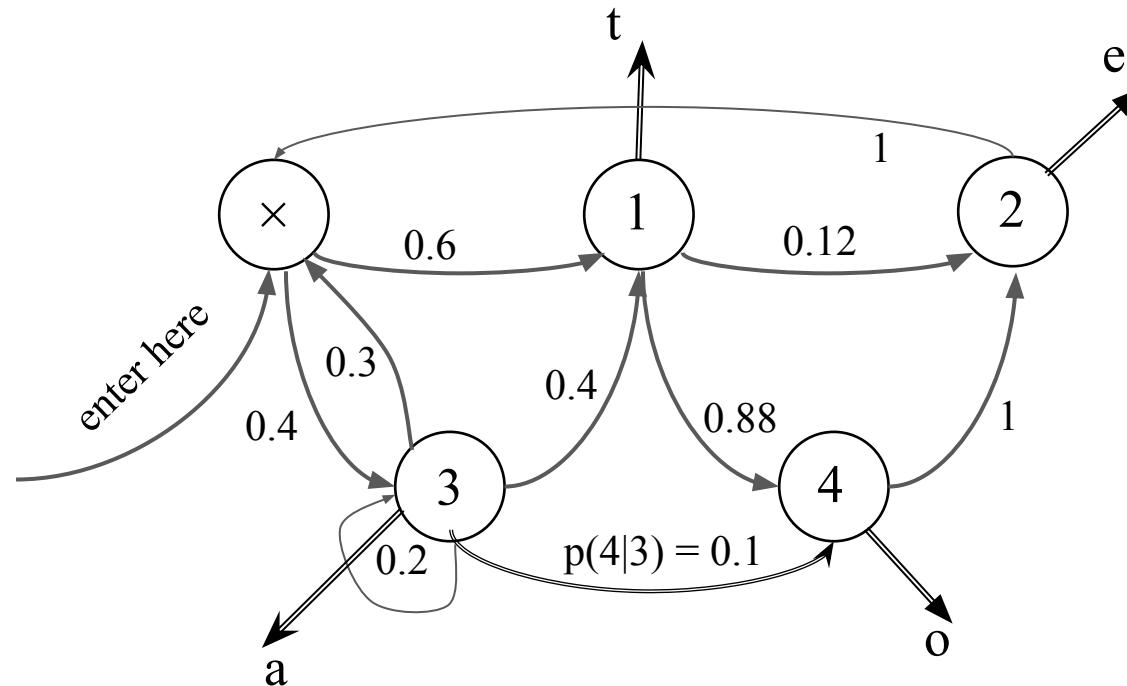
Hidden Markov Models

- The simplest HMM: states generate [*observable*] output (using the “data” alphabet) but remain “*invisible*”:



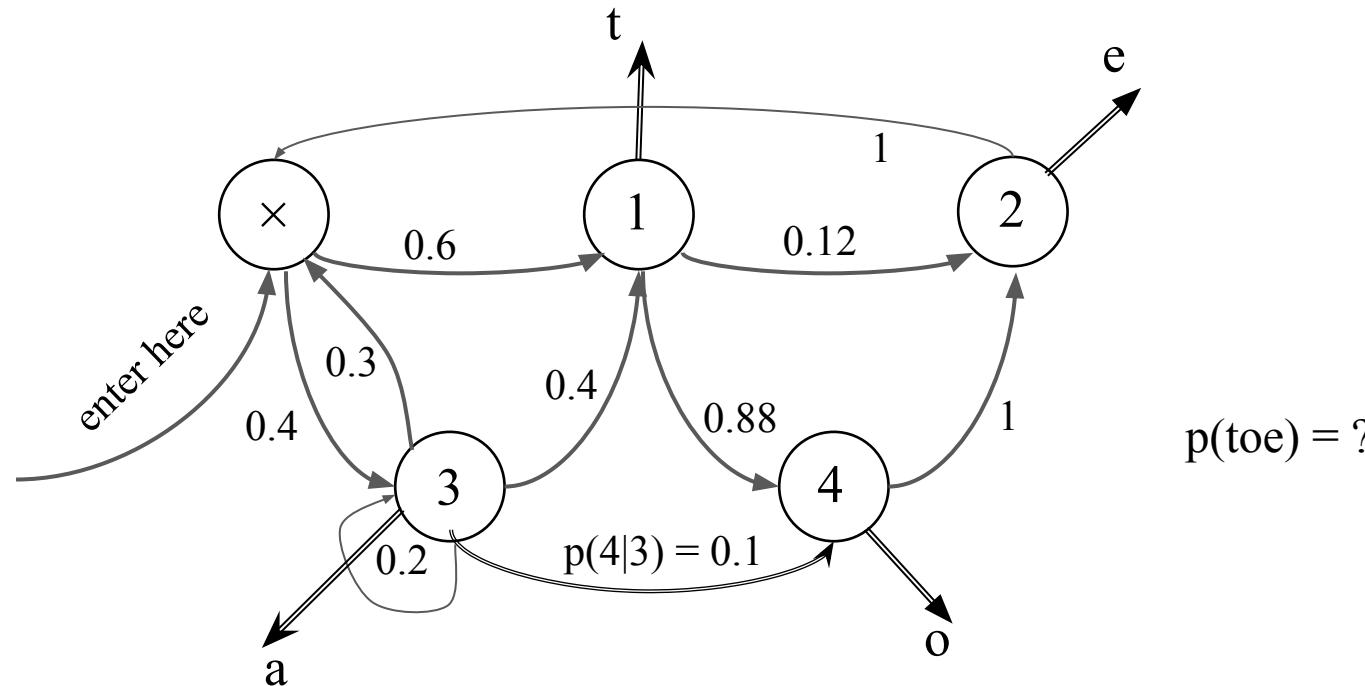
Hidden Markov Models

- The simplest HMM: states generate [*observable*] output (using the “data” alphabet) but remain “*invisible*”:



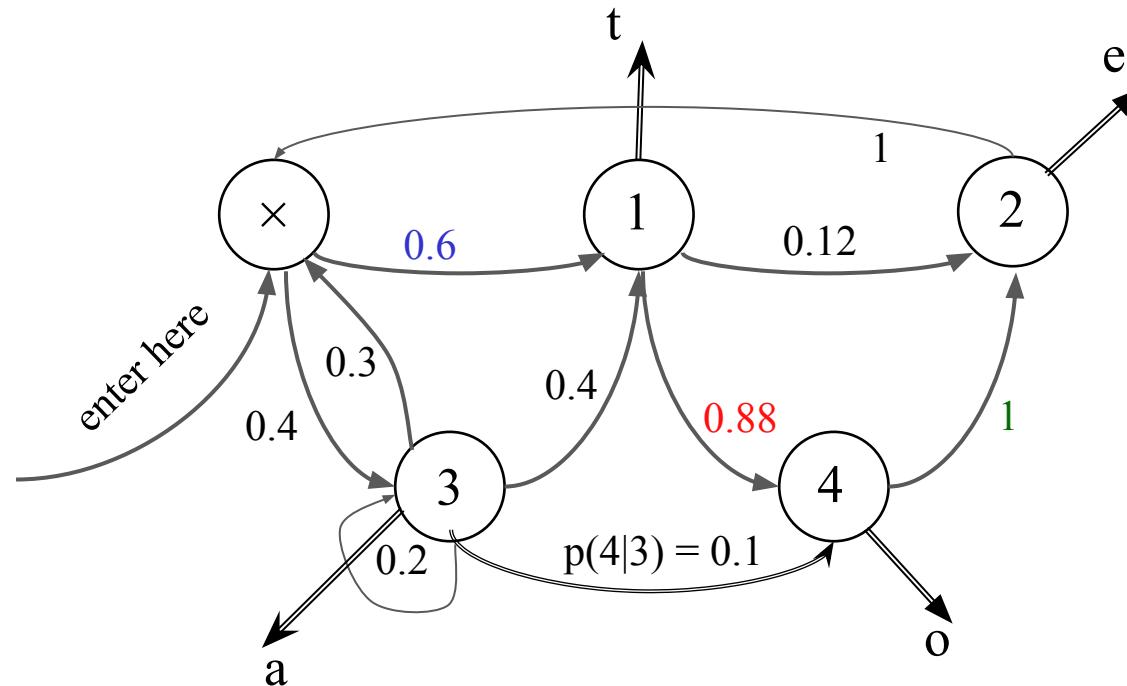
Hidden Markov Models

- The simplest HMM: states generate [*observable*] output (using the “data” alphabet) but remain “*invisible*”:



Hidden Markov Models

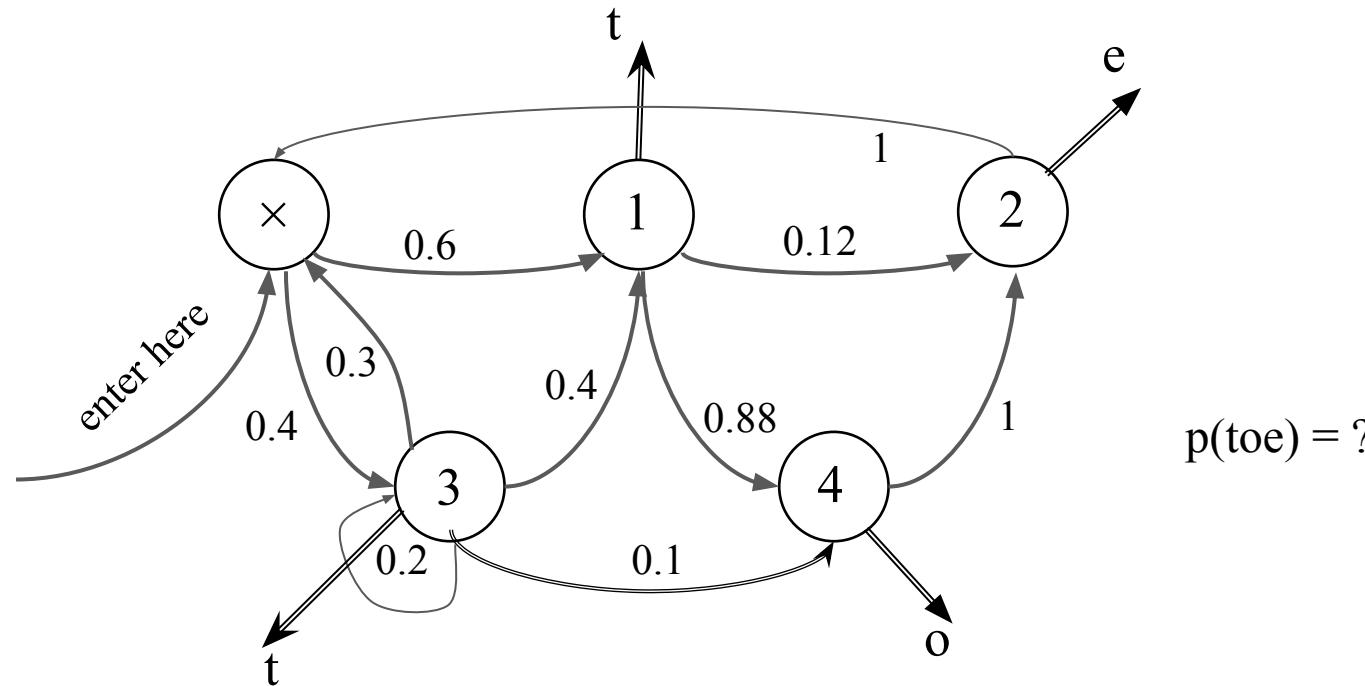
- The simplest HMM: states generate [observable] output (using the “data” alphabet) but remain “*invisible*”:



$$p(\text{toe}) = 0.6 \times 0.88 \times 1 = 0.528$$

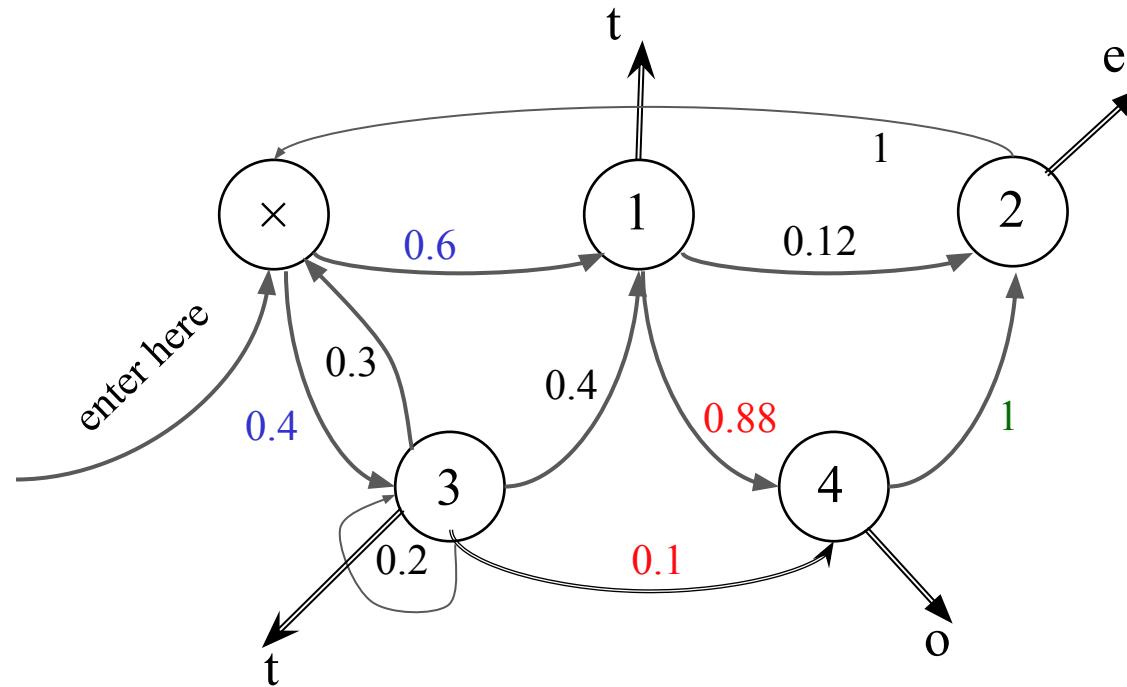
Added Flexibility

- So far, no change; but different states may generate the same output (why not?):



Added Flexibility

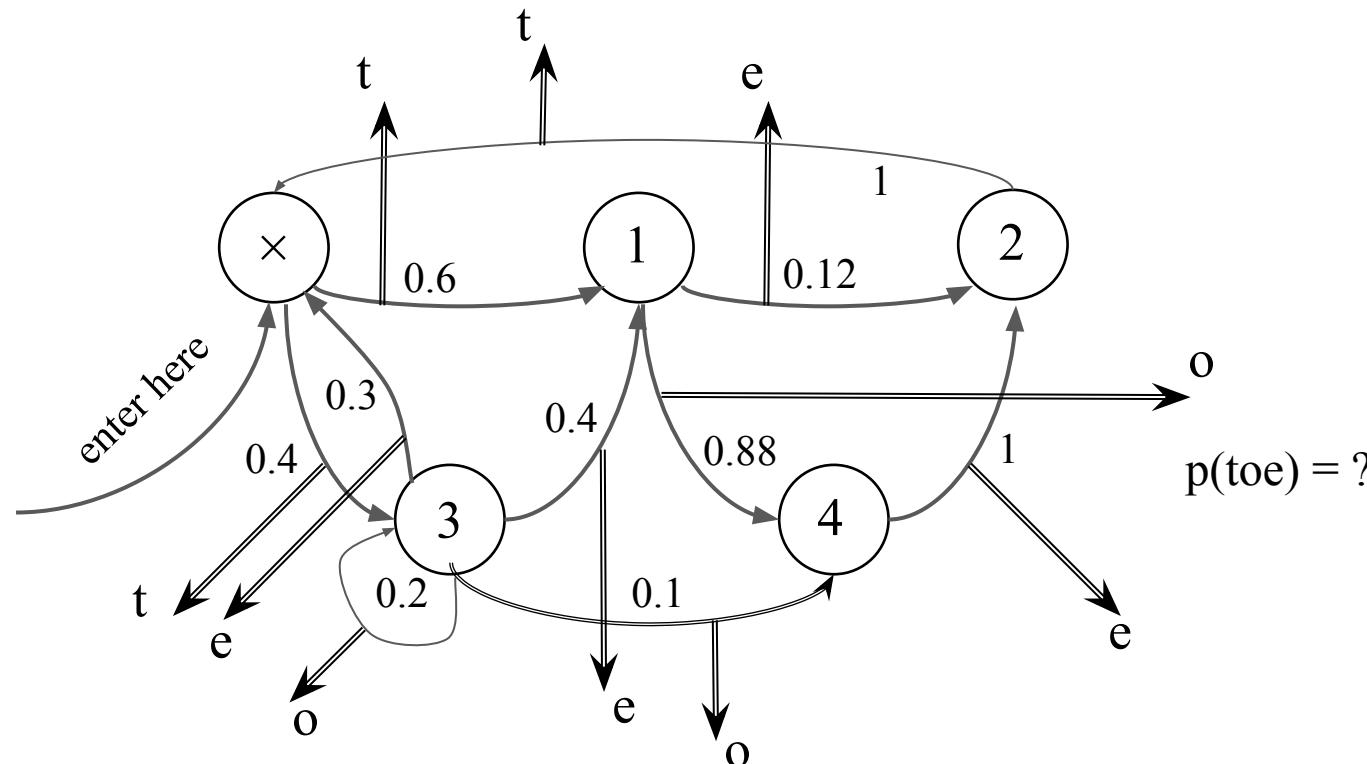
- So far, no change; but different states may generate the same output (why not?):



$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 \approx 0.568$$

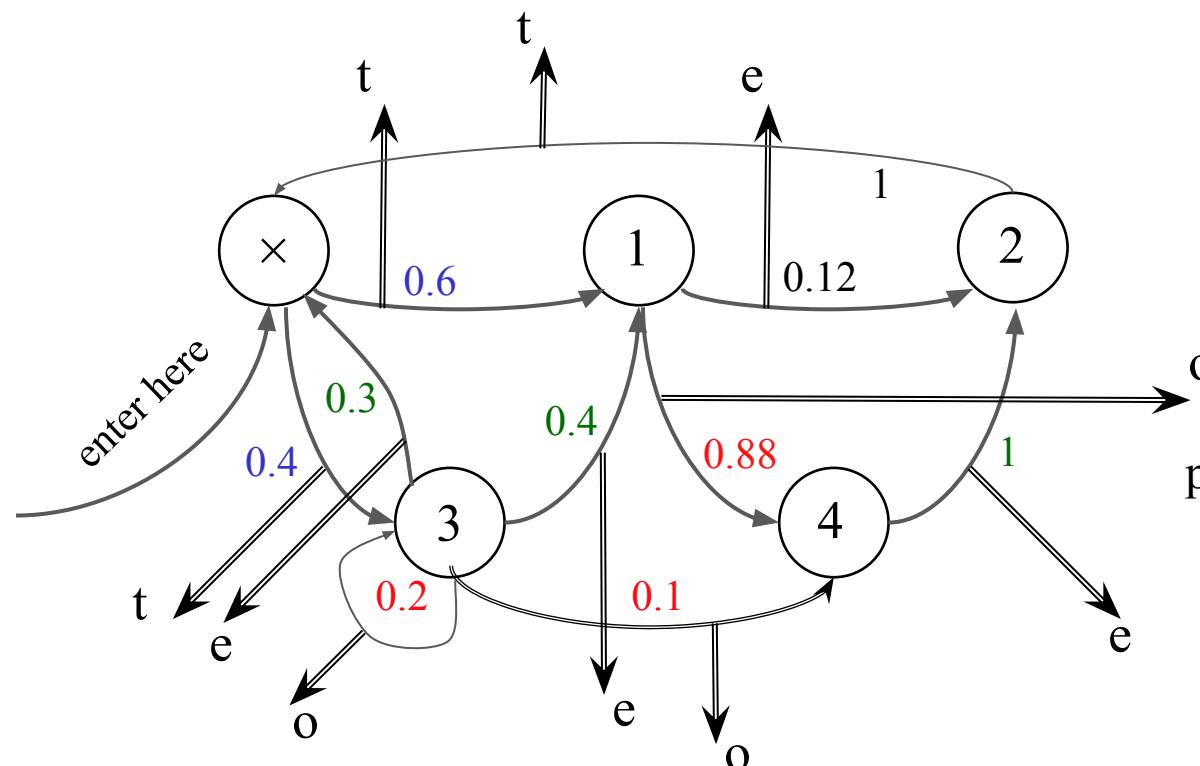
Output from Arcs...

- Added flexibility: Generate output from arcs, not states:



Output from Arcs...

- Added flexibility: Generate output from arcs, not states:

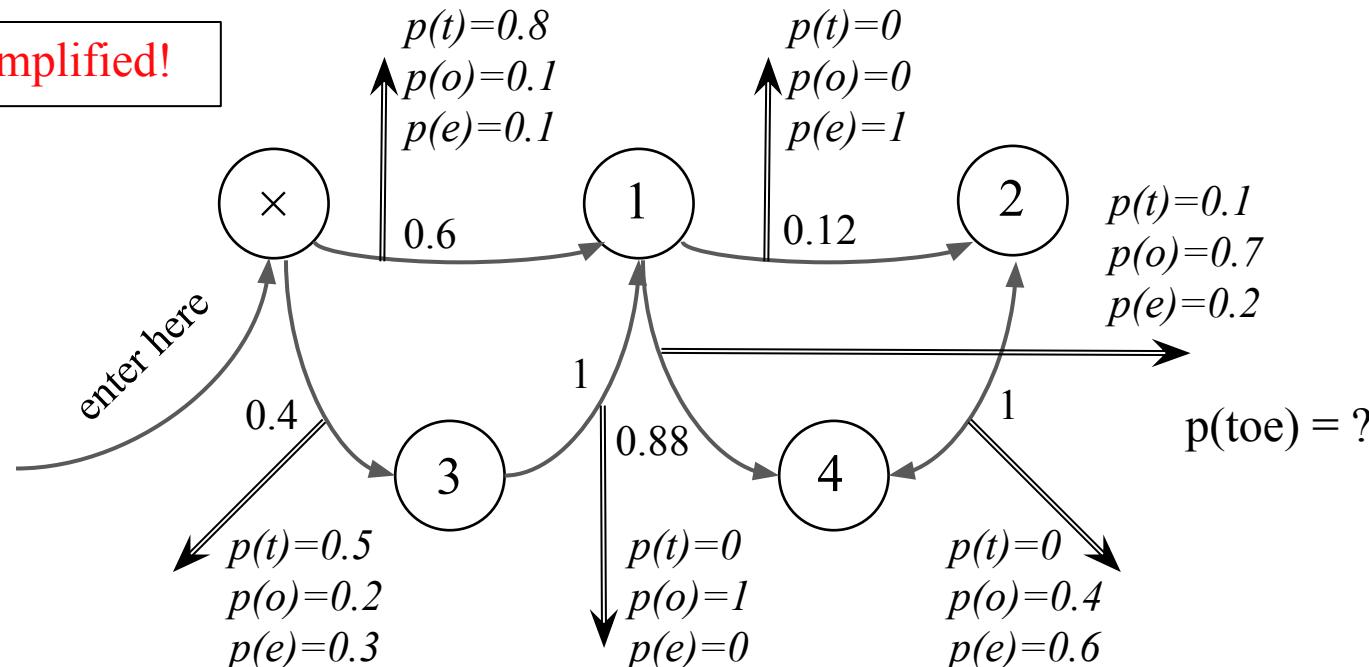


$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 + 0.4 \times 0.2 \times 0.3 + 0.4 \times 0.2 \times 0.4 \approx 0.624$$

... and Finally, Add Output Probabilities

- Maximum flexibility: [Unigram] distribution (sample space: output alphabet) at each output arc:

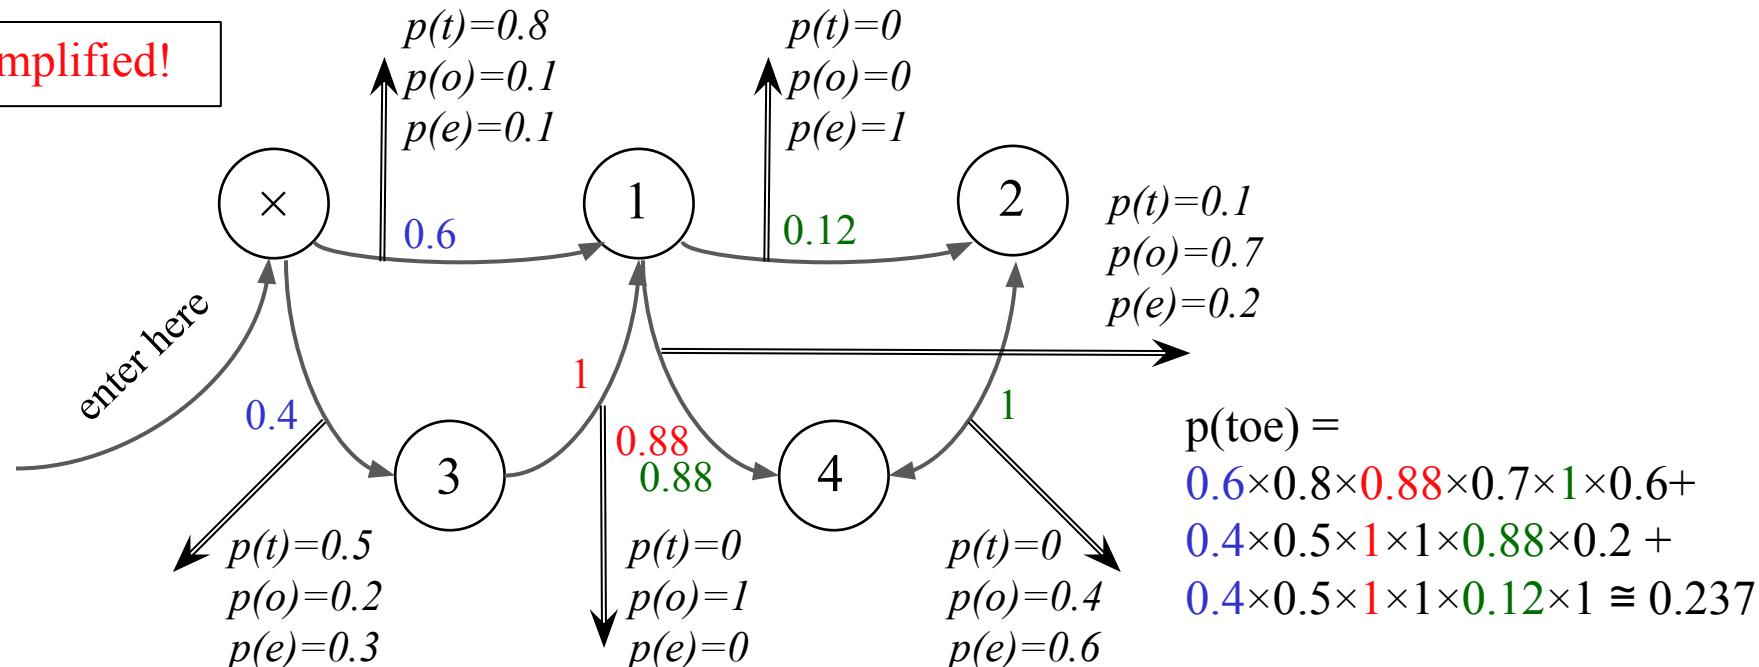
!simplified!



... and Finally, Add Output Probabilities

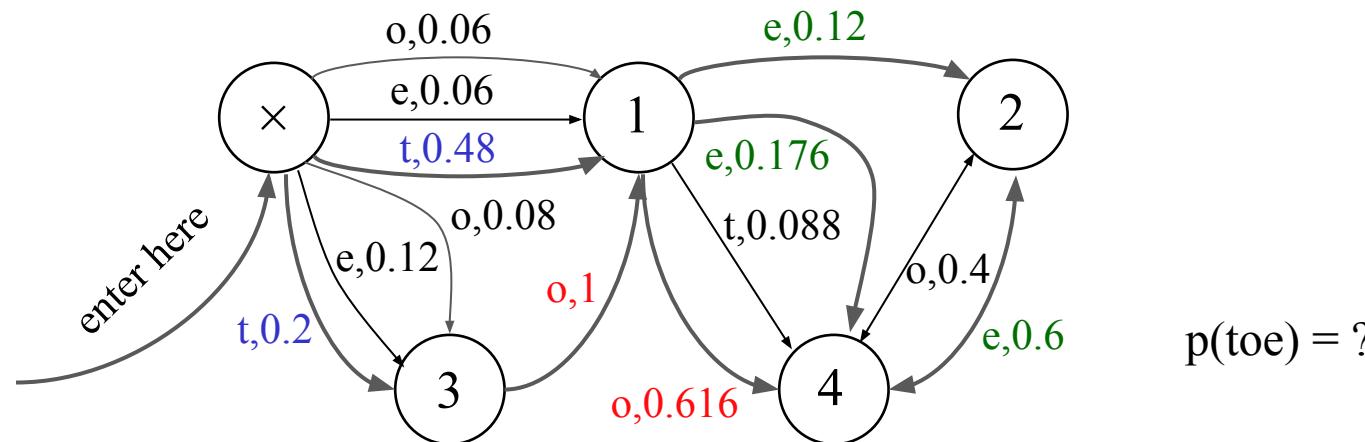
- Maximum flexibility: [Unigram] distribution (sample space: output alphabet) at each output arc:

!simplified!



Slightly Different View

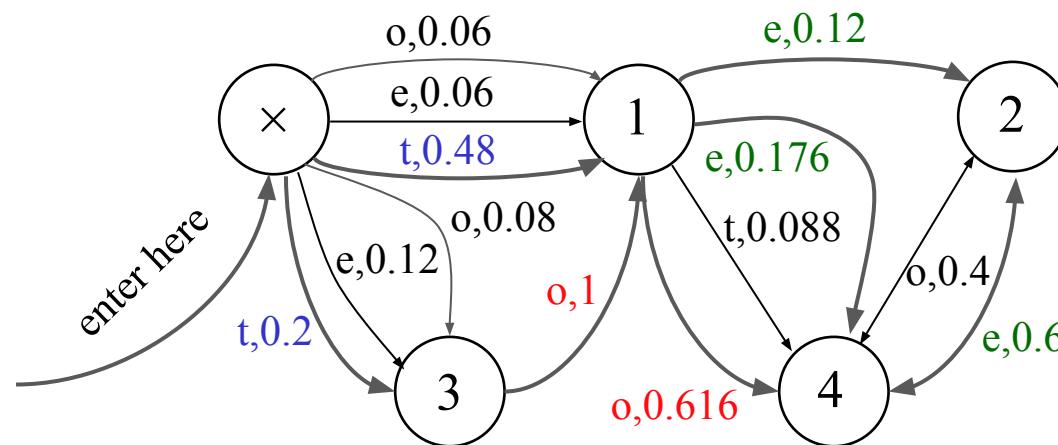
- Allow for multiple arcs from $s_i \rightarrow s_j$, mark them by output symbols, get rid of output distributions:



- In the future, we will use the view more convenient for the problem at hand.

Slightly Different View

- Allow for multiple arcs from $s_i \rightarrow s_j$, mark them by output symbols, get rid of output distributions:



$$p(\text{toe}) = 0.48 \times 0.616 \times 0.6 + \\ 0.2 \times 1 \times 0.176 + \\ 0.2 \times 1 \times 0.12 \approx 0.237$$

- In the future, we will use the view more convenient for the problem at hand.

Formalization

- HMM (the general case): five-tuple (S, s_0, Y, P_S, P_Y) , where:
 - $S = \{s_1, s_2, \dots, s_T\}$ is the set of states, s_0 is the initial state,
 - $Y = \{y_1, y_2, \dots, y_V\}$ is the output alphabet,
 - $P_S(s_j|s_i)$ is the set of prob. distributions of transitions,
 - size of P_S : $|S|^2$
 - $P_Y(y_k|s_i, s_j)$ is the set of output (emission) probability distributions
 - size of P_Y : $|S|^2 \times |Y|$
- Example:
 - $S = \{x, 1, 2, 3, 4\}, s_0 = x$
 - $Y = \{t, o, e\}$

Formalization - Example

- Example (for graph on slide 36):
 - $S = \{x, 1, 2, 3, 4\}$, $s_0 = x$
 - $Y = \{t, o, e\}$
 - P_S :

	\times	1	2	3	4
\times	0	0.6	0	0.4	0
1	0	0	0.12	0	0.88
2	0	0	0	0	1
3	0	1	0	0	0
4	0	0	1	0	0

→ $\Sigma = 1$

P_Y :

e	\times	1	2	3	4
o	\times	1	2	3	4
t	\times	1	2	3	4
\times		0.8		0.5	0.7
1					0.1
2					0
3		0			
4			0		

$\Sigma = 1$

Using the HMM

- The generation algorithm (of limited value :-)):
 1. Start in $s = s_0$
 2. Move from s to s' with probability $P_s(s'|s)$
 3. Output (emit) symbol y_k with probability $P_s(y_k|s,s')$
 4. Repeat from step 2 (until somebody says enough)
- More interesting usage:
 1. Given an output sequence $Y = \{y_1, y_2, \dots, y_k\}$, compute its probability.
 2. Given an output sequence $Y = \{y_1, y_2, \dots, y_k\}$, compute the most likely sequence of states which has generated it.
...plus variations: e.g., n best state sequences

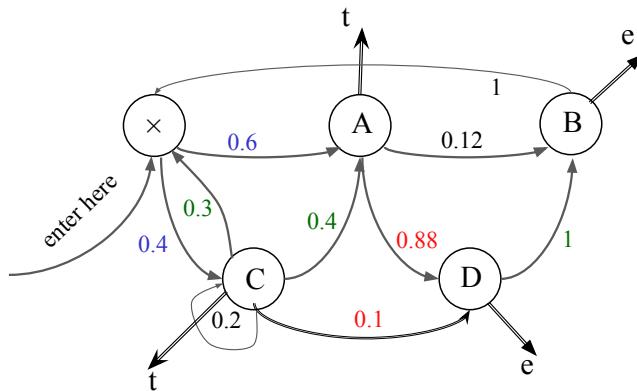
Trellis

HMM: The Two Tasks

- HMM (the general case): five-tuple (S, S_0, Y, P_S, P_Y) , where:
 - $S = \{s_1, s_2, \dots, s_T\}$ is the set of states, S_0 is the initial state,
 - $Y = \{y_1, y_2, \dots, y_V\}$ is the output alphabet,
 - $P_S(s_j|s_i)$ is the set of prob. distributions of transitions,
 - $P_Y(y_k|s_i, s_j)$ is the set of output (emission) probability distributions
- Given an HMM & an output sequence $Y = \{y_1, y_2, \dots, y_k\}$:
 - (Task 1) compute the probability of Y ;
 - (Task 2) compute the most likely sequence of states which has generated Y .

Trellis: HMM “roll-out”

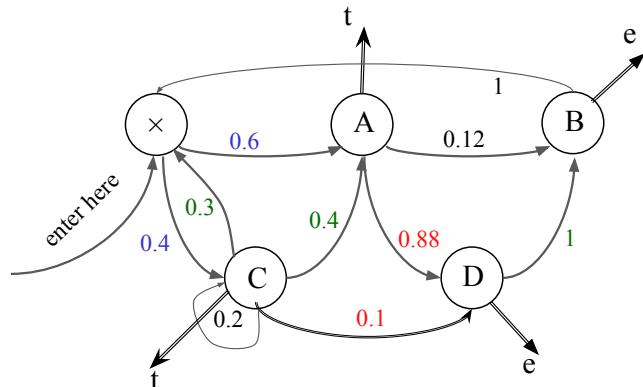
HMM:



$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 \approx 0.568$$

Trellis: HMM “roll-out”

HMM:



$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 \approx 0.568$$

Trellis:

time/position:

0



“roll-out”

- Trellis state: (HMM state, position)

Y:

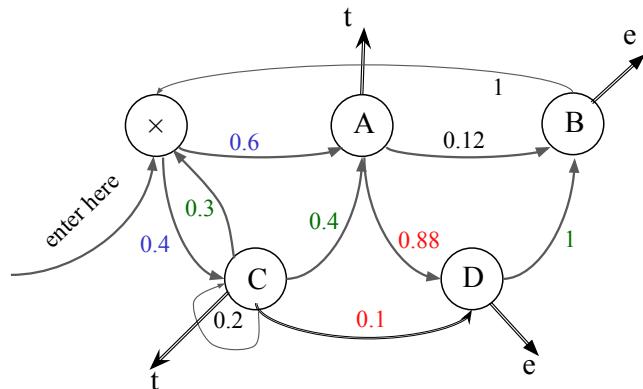
t

o

e

Trellis: HMM “roll-out”

HMM:

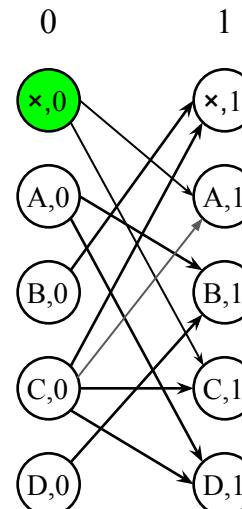


$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 \approx 0.568$$

Trellis:

time/position:

“roll-out”

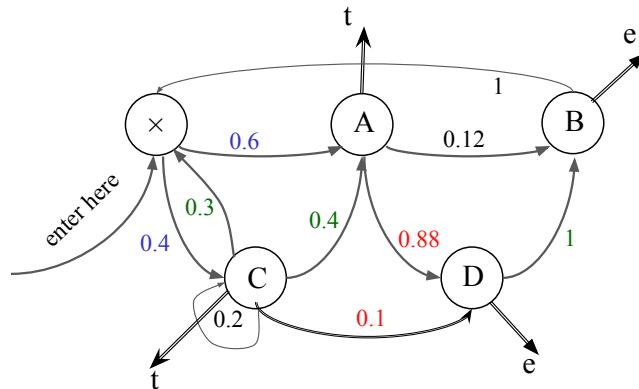


- Trellis state: (HMM state, position)

Y: t o e

Trellis: HMM “roll-out”

HMM:

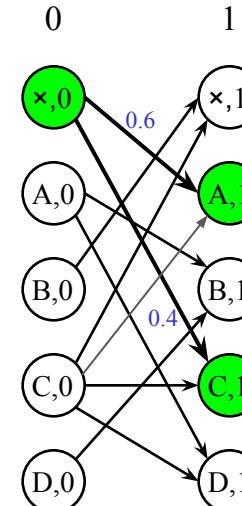


$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 \approx 0.568$$

Trellis:

time/position:

“roll-out”



- Trellis state: (HMM state, position)

Y:

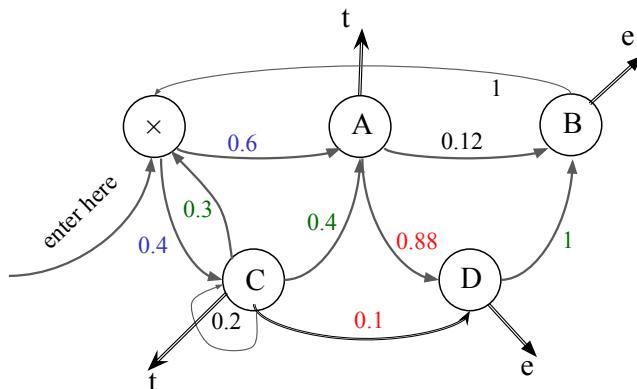
t

o

e

Trellis: HMM “roll-out”

HMM:

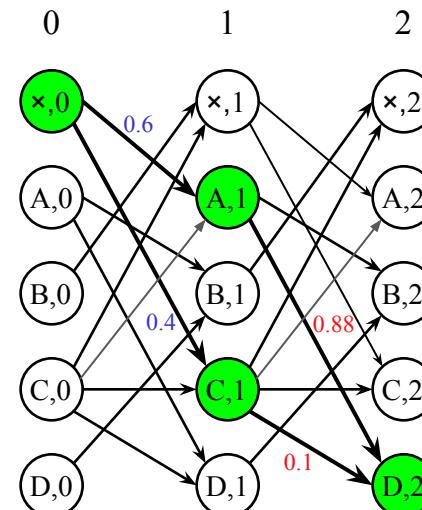


$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 \approx 0.568$$

Trellis:

time/position:

“roll-out”

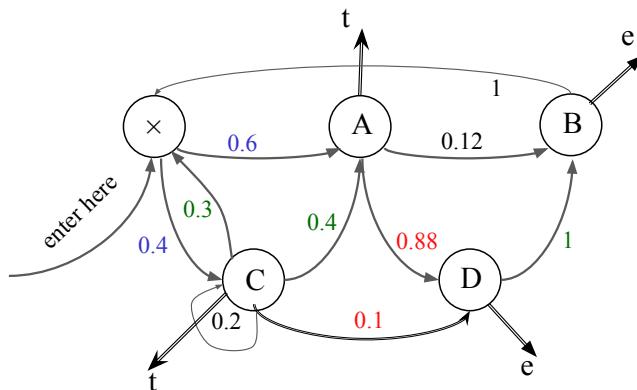


- Trellis state: (HMM state, position)

Y: t o e

Trellis: HMM “roll-out”

HMM:

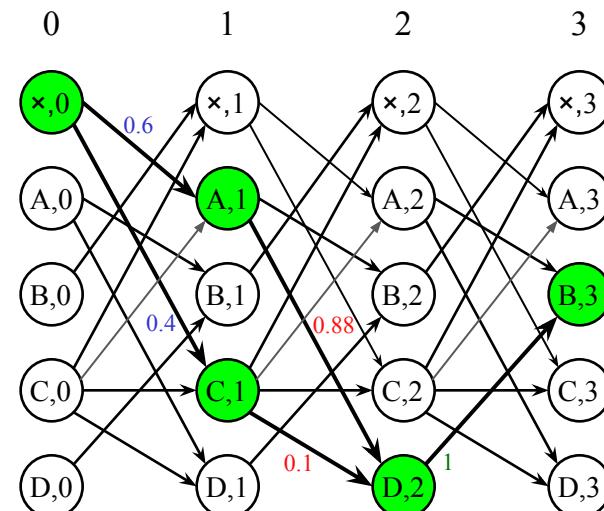


$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 \approx 0.568$$

Trellis:

time/position:

“roll-out”



- Trellis state: (HMM state, position)

Y:

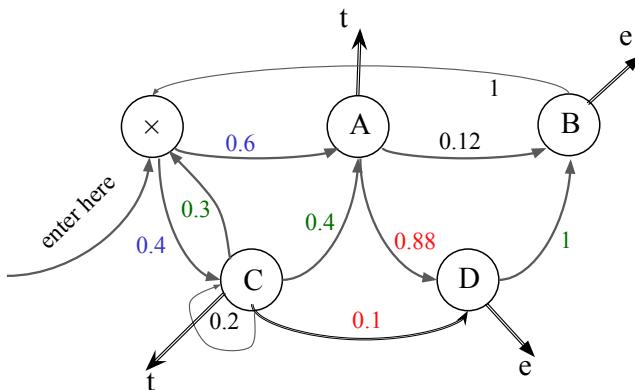
t

o

e

Trellis: HMM “roll-out”

HMM:

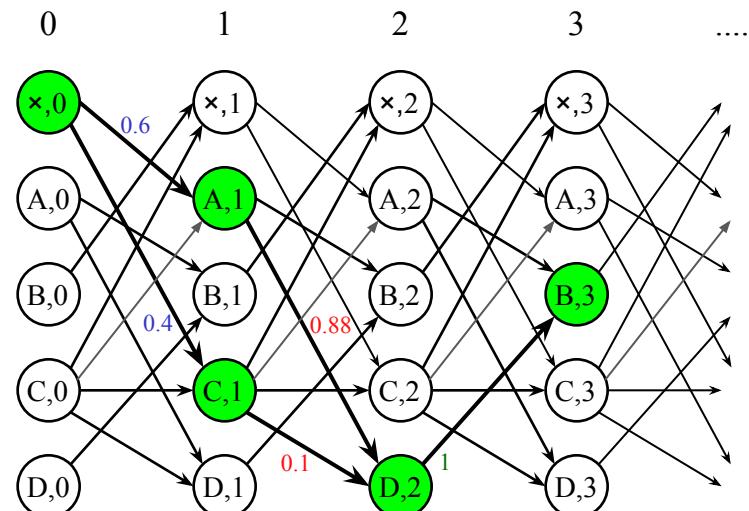


$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 \approx 0.568$$

Trellis:

time/position:

“roll-out”



- Trellis state: (HMM state, position)

Y:

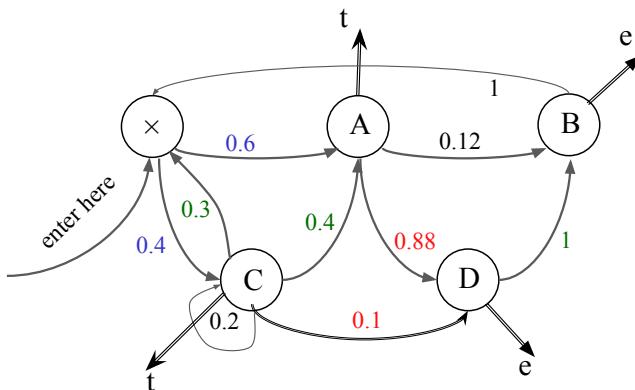
t

o

e

Trellis: HMM “roll-out”

HMM:

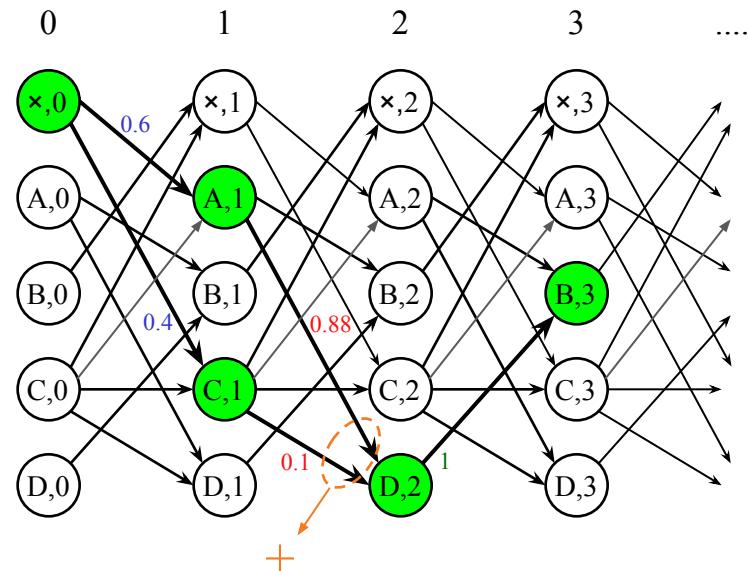


$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 \approx 0.568$$

Trellis:

time/position:

“roll-out”



- Trellis state: (HMM state, position)
- each state: holds one number (prob): α
- probability of Y: $\sum \alpha$ in the last state

Y:

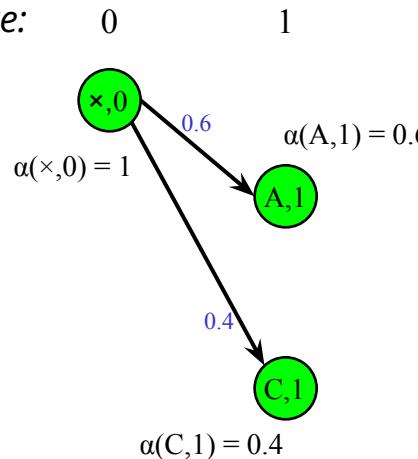
$$\begin{aligned} \alpha(x,0) &= 1 & \alpha(A,1) &= 0.6 & \alpha(D,2) &= 0.568 & \alpha(B,3) &= 0.568 \\ \alpha(C,1) &= 0.4 & & & & & \end{aligned}$$

Creating the Trellis: The Start

- Start in the start state (\times),

- set its $\alpha(\times, 0)$ to 1

position/stage:



- Create the first stage:

- get the first “output” symbol y_1
 - create the first stage (column)
 - but only those Trellis states which generate y_1

- set their $\alpha(state, 1)$ to the $P_s(state | \times) \alpha(\times, 0)$

- and forget about the 0-th stage

Y: t

Trellis: The Next Step

- Suppose we are in stage i
- Creating the next stage: *position/stage:* 1 2
 - create all trellis states in the next stage which generate y_{i+1} , but only those reachable from any of the stage-i states
 - set their $\alpha(state, i+1)$ to:
 $P_s(state|prev.state) \times \alpha(prev.state, i)$
(add up all such numbers on arcs going to a common Trellis state)
...and forget about stage i

A,1

C,1

$$\alpha(A,1) = 0.6$$

$$\alpha(C,1) = 0.4$$

Y:

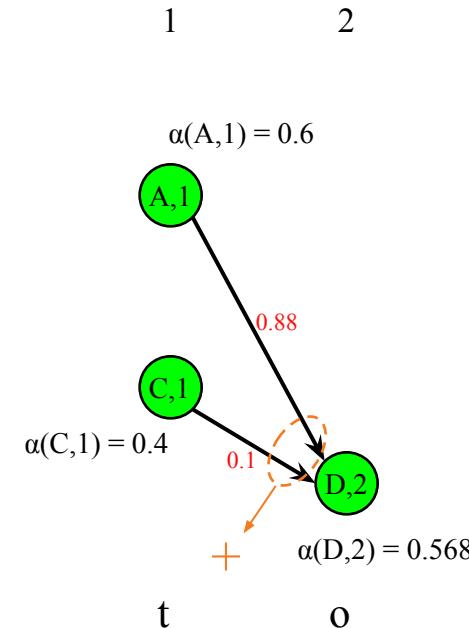
t

o

Trellis: The Next Step

- Suppose we are in stage i
- Creating the next stage: *position/stage:* 1 2
 - create all trellis states in the next stage which generate y_{i+1} , but only those reachable from any of the stage-i states
 - set their $\alpha(state, i+1)$ to:
 $P_s(state|prev.state) \times \alpha(prev.state, i)$
(add up all such numbers on arcs going to a common Trellis state)
...and forget about stage i

Y:

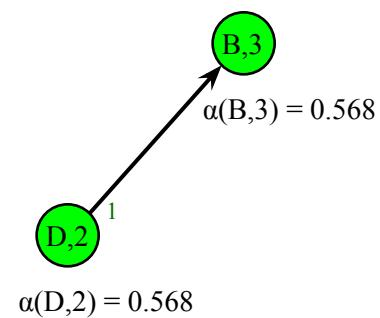


Trellis: The Last Step

- Continue until output exhausted
 - $|Y| = 3$: until stage 3 *position/stage:*
- Add together all the $\alpha(state, |Y|)$
- That's the P(Y)
- Observation (pleasant):
 - memory usage max: $2|S|$
 - multiplications max: $|S|^2|Y|$

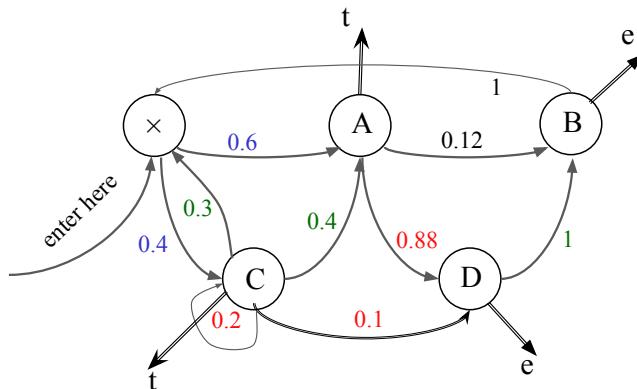
Y:

0 e



Trellis (again)

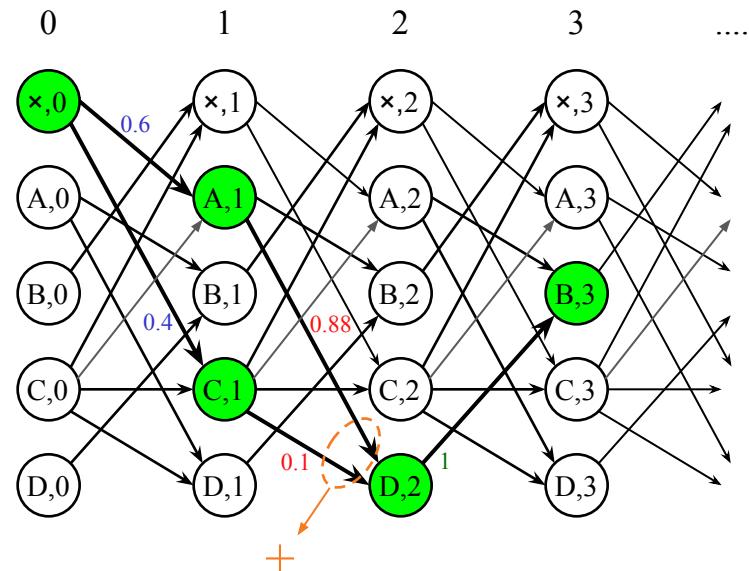
HMM:



Trellis:

time/position:

"roll-out"



- Trellis state: (HMM state, position)
- each state: holds one number (prob): α
- probability of Y: $\sum \alpha$ in the last state

Y:

t o e

$$\begin{aligned} \alpha(x,0) &= 1 & \alpha(A,1) &= 0.6 & \alpha(D,2) &= 0.568 & \alpha(B,3) &= 0.568 \\ \alpha(C,1) &= 0.4 & & & & & \end{aligned}$$

HMM: The Two Tasks

- HMM (the general case): five-tuple (S, S_0, Y, P_S, P_Y) , where:
 - $S = \{s_1, s_2, \dots, s_T\}$ is the set of states, S_0 is the initial state,
 - $Y = \{y_1, y_2, \dots, y_V\}$ is the output alphabet,
 - $P_S(s_j|s_i)$ is the set of prob. distributions of transitions,
 - $P_Y(y_k|s_i, s_j)$ is the set of output (emission) probability distributions
- Given an HMM & an output sequence $Y = \{y_1, y_2, \dots, y_k\}$:
 - (Task 1) compute the probability of Y ;
 - (Task 2) compute the most likely sequence of states which has generated Y .

Moodle Quiz

Moodle Quiz



<https://dl1.cuni.cz/course/view.php?id=18547>

Trellis: The General Case (still, bigrams)

- B

General Trellis: The Next Step

- B

Trellis: The Complete Example

- B

The Case of Trigrams

- B

Trigrams with Classes

- B

Class Trigrams: the Trellis

- B

Overlapping Classes

- B

Overlapping Classes: Trellis Example

- B

Trellis: Remarks

- So far, we went left to right (computing α)
- Same result: going right to left (computing β)
 - supposed we know where to start (finite data)
- In fact, we might start in the middle going left and right
- Important for parameter estimation
 - (Forward-Backward Algorithm alias Baum-Welch)
- Implementation issues:
 - scaling/normalizing probabilities, to avoid too small numbers & addition problems with many transitions