

# Statistical Methods in Natural Language Processing

## 9. Hidden Markov Models and Baum Welch.

Pavel Pecina, Jindřich Helcl

16 December, 2025

# Course Segments

1. Introduction, probability, essential information theory
2. Statistical language modelling (n-gram)
3. Statistical properties of words
4. Word representations
5. Hidden Markov models, Tagging

## Recap from Last Week

# Markov Properties

- Markov Chain can generalize to any process (not just words):
  - Sequence of random variables:  $X = (X_1, X_2, \dots, X_T)$
  - Sample space  $S$  (*states*), size  $N$ :  $S = \{s_0, s_1, s_2, \dots, s_N\}$
- Two properties

1. Limited history (context, horizon):

$$\forall i \in 1..T; P(X_i | X_1, \dots, X_{i-1}) = P(X_i | X_{i-1})$$



2. Time invariance (Markov Chain is stationary, homogeneous)

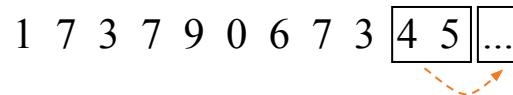
$$\forall i \in 1..T, \forall y, x \in S; P(X_i=y | X_{i-1}=x) = p(y|x)$$



ok ... same distribution

# Long History Possible

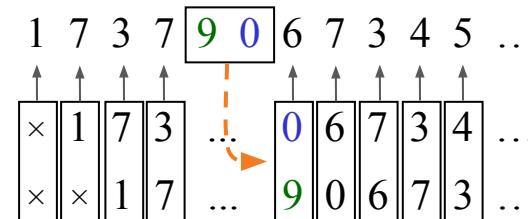
- What if we want trigrams:



- Formally, use transformation:

- Define new variables  $Q_i$ , such that  $X_i = \{Q_{i-1}, Q_i\}$
- And then  $P(X_i|X_{i-1}) = P(Q_{i-1}, Q_i|Q_{i-2}, Q_{i-1}) = P(Q_i|Q_{i-2}, Q_{i-1})$

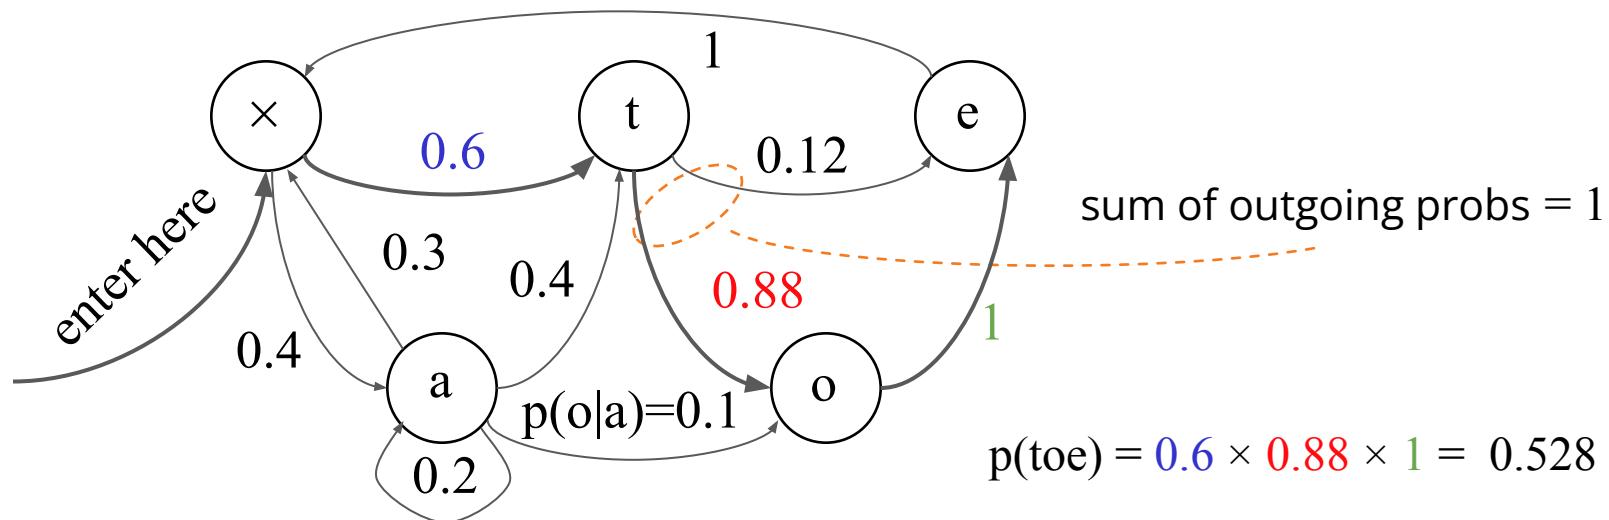
Predicting ( $X_i$ ):



History ( $X_{i-1} = \{Q_{i-2}, Q_{i-1}\}$ ):

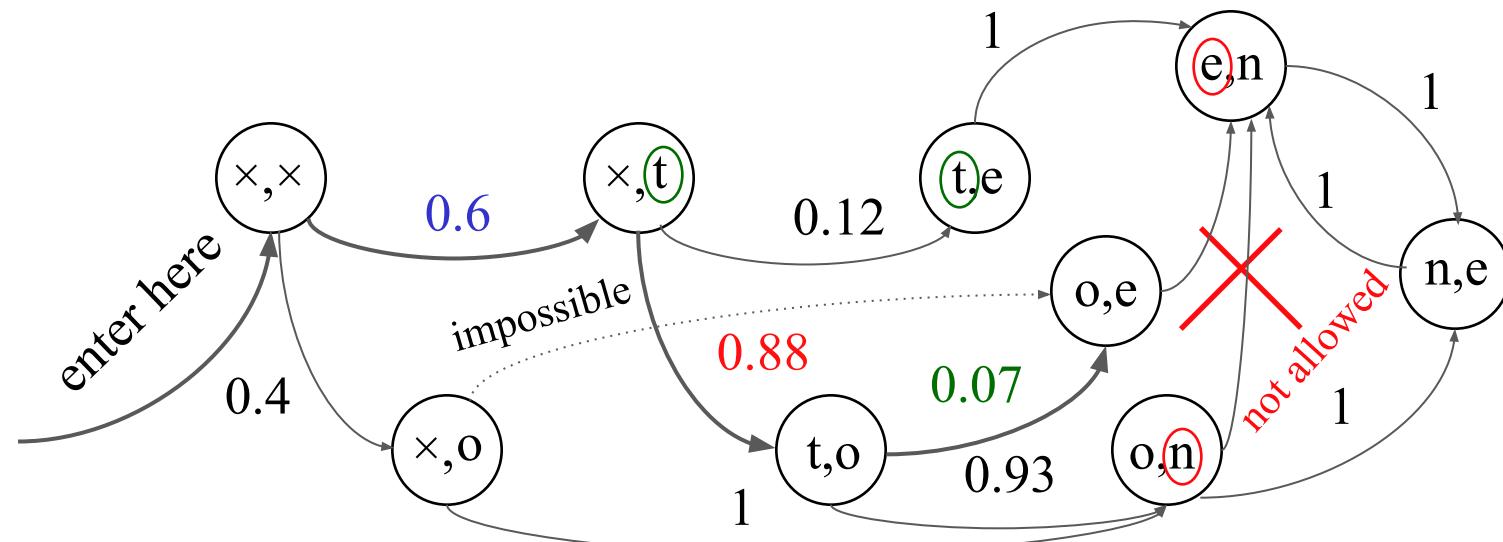
# Markov Models: Bigram Case

- Nodes: States,  $S = \{s_0, s_1, s_2, \dots, s_N\}$
- Arcs: Transitions with probabilities,  $P(X_i|X_{i-1})$ ,  $X_i$  generates  $s_i$



# Markov Models: Trigram Case

- Nodes: Pairs of states  $(s_k, s_l)$ ,  $S = \{s_0, s_1, s_2, \dots, s_N\}$
- Arcs: Transitions with probabilities,  $P(X_i|X_{i-1})$ ,  $X_i$  generates  $(s_k, s_l)$



$$p(\text{toe}) = 0.6 \times 0.88 \times 0.07 = 0.037$$

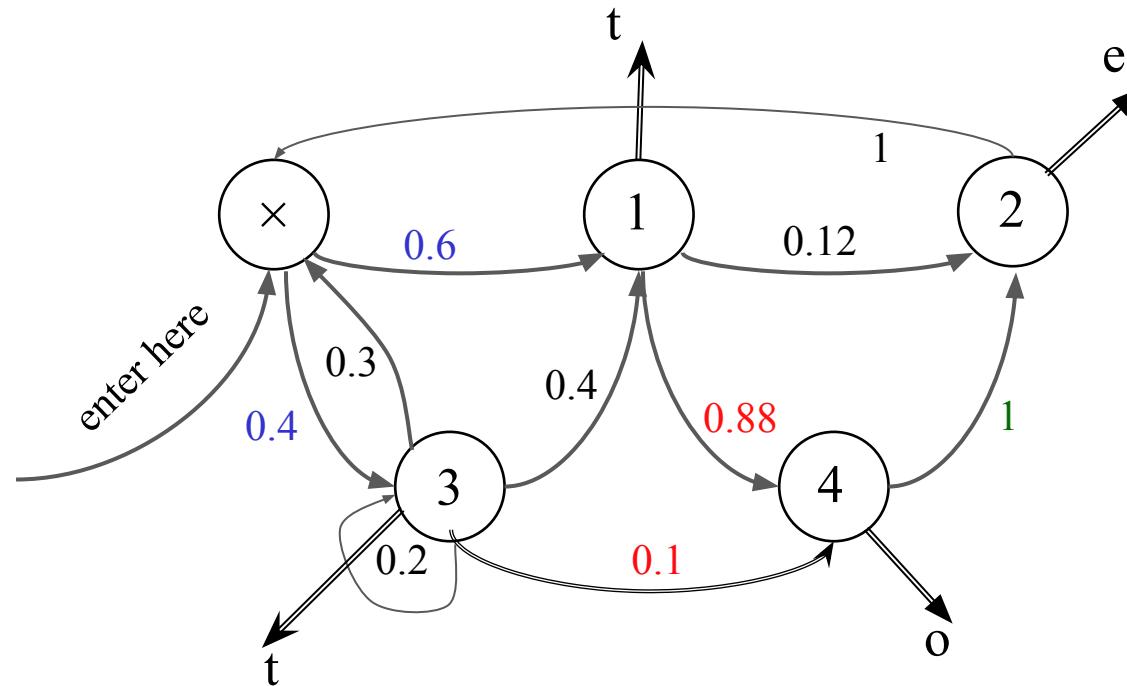
$$p(\text{one}) = ?$$

# Finite State Automaton

- States ~ symbols of the [input / output] alphabet
  - pairs (or more): last element of the n-tuple
- Arcs ~ transitions (sequence of states)
  - Classical FSA: alphabet symbols on arcs:
  - possible transformation: arcs  $\leftrightarrow$  nodes
- Possible thanks to the “limited history” Markov Property
- So far: **Visible** Markov Models (VMM)

# Hidden Markov Models

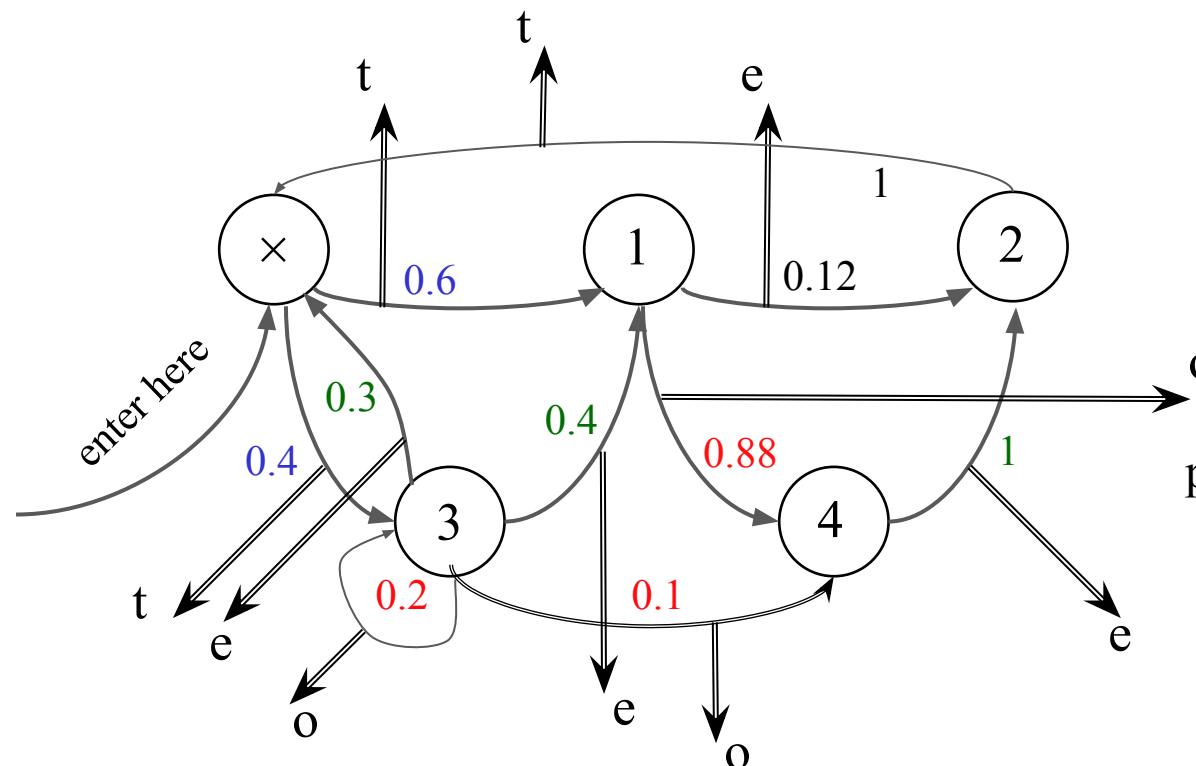
- The simplest HMM: states generate [*observable*] output (using the “data” alphabet) but remain “*invisible*”



$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 \approx 0.568$$

# Output from Arcs...

- Added flexibility: Generate output from arcs, not states:

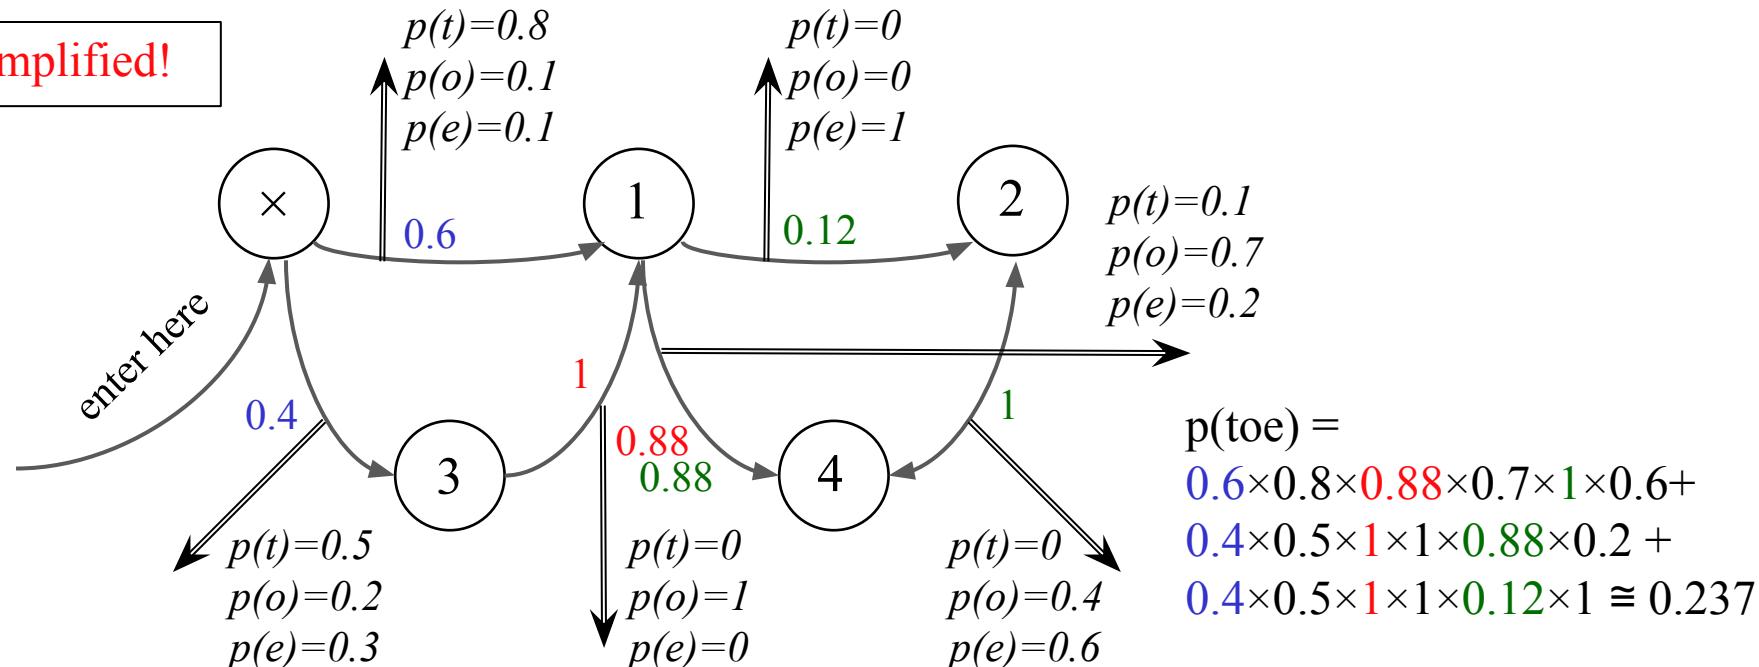


$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 0.1 \times 1 + 0.4 \times 0.2 \times 0.3 + 0.4 \times 0.2 \times 0.4 \approx 0.624$$

# ... and Finally, Add Output Probabilities

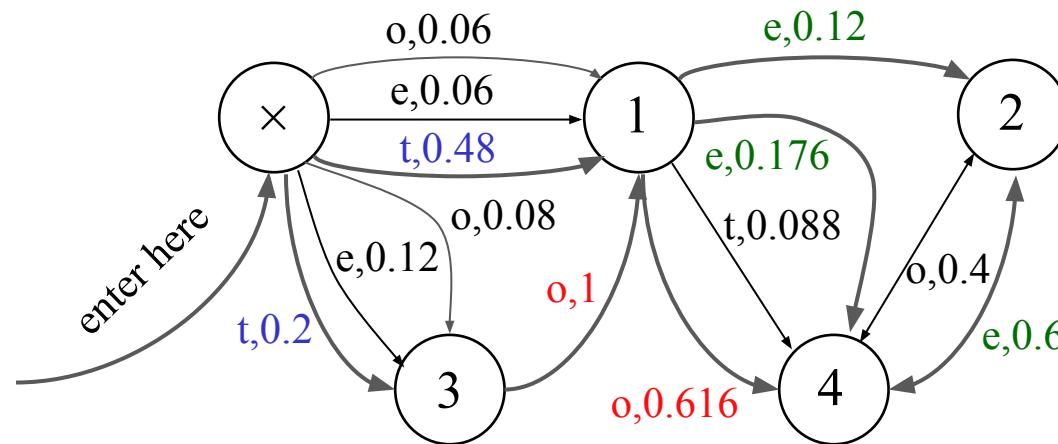
- Maximum flexibility: [Unigram] distribution (sample space: output alphabet) at each output arc:

!simplified!



# Slightly Different View

- Allow for multiple arcs from  $s_i \rightarrow s_j$ , mark them by output symbols, get rid of output distributions:



$$p(\text{toe}) = 0.48 \times 0.616 \times 0.6 + \\ 0.2 \times 1 \times 0.176 + \\ 0.2 \times 1 \times 0.12 \approx 0.237$$

# Formalization

- HMM (the general case): five-tuple  $(S, s_0, Y, P_S, P_Y)$ , where:
  - $S = \{s_1, s_2, \dots, s_T\}$  is the set of states,  $s_0$  is the initial state,
  - $Y = \{y_1, y_2, \dots, y_V\}$  is the output alphabet,
  - $P_S(s_j|s_i)$  is the set of prob. distributions of transitions,
    - size of  $P_S$ :  $|S|^2$
  - $P_Y(y_k|s_i, s_j)$  is the set of output (emission) probability distributions
    - size of  $P_Y$ :  $|S|^2 \times |Y|$
- Example:
  - $S = \{x, 1, 2, 3, 4\}, s_0 = x$
  - $Y = \{t, o, e\}$

# Formalization - Example

- Example (for graph on slide 36):
  - $S = \{x, 1, 2, 3, 4\}$ ,  $s_0 = x$
  - $Y = \{t, o, e\}$
  - $P_S$ :

	$\times$	1	2	3	4
$\times$	0	0.6	0	0.4	0
1	0	0	0.12	0	0.88
2	0	0	0	0	1
3	0	1	0	0	0
4	0	0	1	0	0

→  $\Sigma = 1$

$P_Y$ :

e	$\times$	1	2	3	4
o	$\times$	1	2	3	4
t	$\times$	1	2	3	4
$\times$		0.8		0.5	0.7
1					0.1
2					0
3		0			
4			0		

$\Sigma = 1$

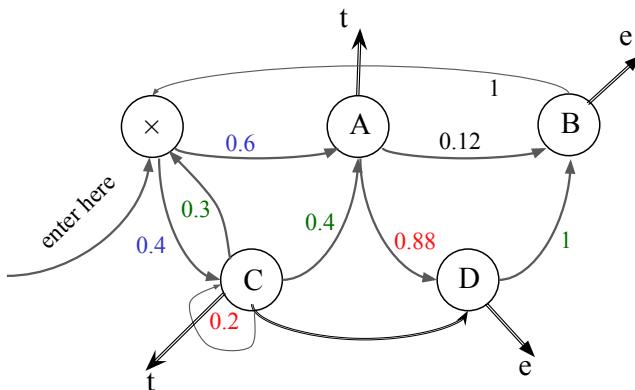
# HMM: The Two Tasks

- HMM (the general case): five-tuple  $(S, S_0, Y, P_S, P_Y)$ , where:
  - $S = \{s_1, s_2, \dots, s_T\}$  is the set of states,  $S_0$  is the initial state,
  - $Y = \{y_1, y_2, \dots, y_V\}$  is the output alphabet,
  - $P_S(s_j|s_i)$  is the set of prob. distributions of transitions,
  - $P_Y(y_k|s_i, s_j)$  is the set of output (emission) probability distributions
- Given an HMM & an output sequence  $Y = \{y_1, y_2, \dots, y_k\}$ :
  - (Task 1) compute the probability of  $Y$ ;
  - (Task 2) compute the most likely sequence of states which has generated  $Y$ .

# Trellis

# Trellis: HMM “roll-out”

HMM:

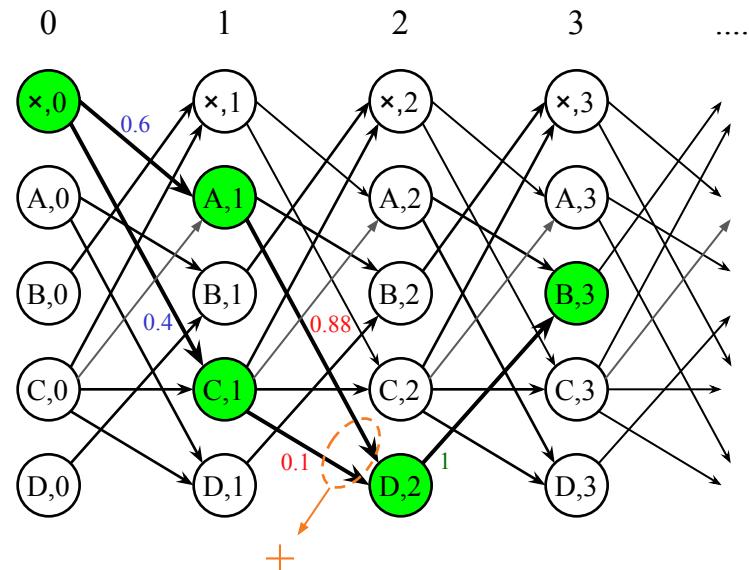


$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 1 \times 1 \approx 0.568$$

Trellis:

*time/position:*

“roll-out”



- Trellis state: (HMM state, position)
- each state: holds one number (prob):  $\alpha$
- probability of Y:  $\sum \alpha$  in the last state

Y:

t

o

e

$$\begin{aligned} \alpha(x,0) &= 1 & \alpha(A,1) &= 0.6 & \alpha(D,2) &= 0.568 & \alpha(B,3) &= 0.568 \\ \alpha(C,1) &= 0.4 & & & & & \end{aligned}$$

# Creating the Trellis: The Start

- Start in the start state ( $\times$ ),

- set its  $\alpha(\times, 0)$  to 1

position/stage:

0                    1

- Create the first stage:

- get the first “output” symbol  $y_1$

- create the first stage (column)

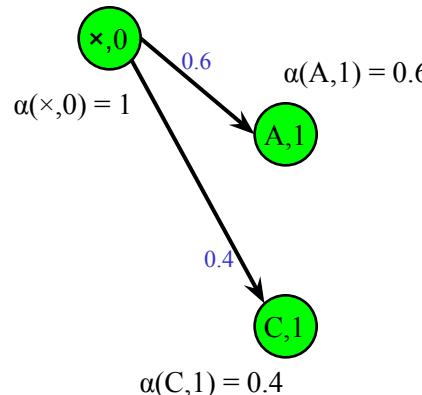
- but only those Trellis states

which generate  $y_1$

- set their  $\alpha(state, 1)$  to the

$P_s(state|\times) \alpha(\times, 0)$

- and forget about the 0-th stage



Y:                    t

# Trellis: The Next Step

- Suppose we are in stage i
- Creating the next stage: *position/stage:* 1 2
  - create all trellis states in the next stage which generate  $y_{i+1}$ , but only those reachable from any of the stage-i states
  - set their  $\alpha(state, i+1)$  to:  
 $P_s(state|prev.state) \times \alpha(prev.state, i)$   
(add up all such numbers on arcs going to a common Trellis state)  
...and forget about stage i

A,1

C,1

$$\alpha(A,1) = 0.6$$

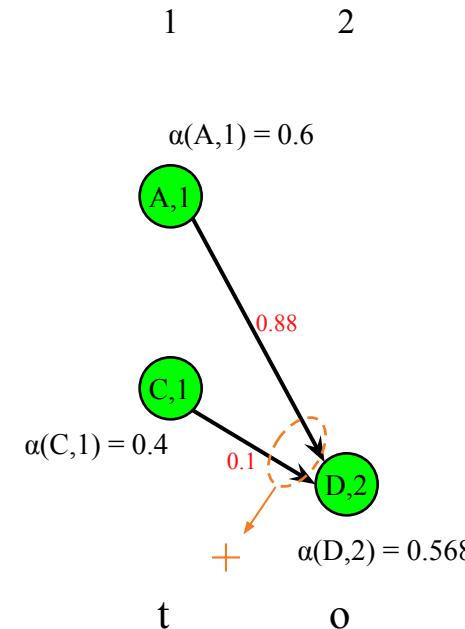
$$\alpha(C,1) = 0.4$$

Y: t o

# Trellis: The Next Step

- Suppose we are in stage  $i$
- Creating the next stage:  $position/stage:$ 
  - create all trellis states in the next stage which generate  $y_{i+1}$ , but only those reachable from any of the stage- $i$  states
  - set their  $\alpha(state, i+1)$  to:  
 $P_s(state|prev.state) \times \alpha(prev.state, i)$   
(add up all such numbers on arcs going to a common Trellis state)  
...and forget about stage  $i$

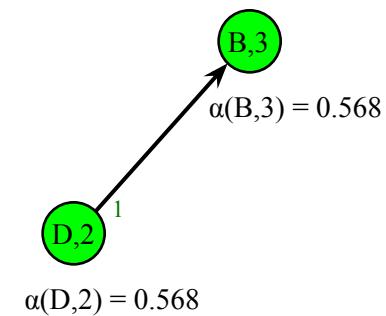
Y:



# Trellis: The Last Step

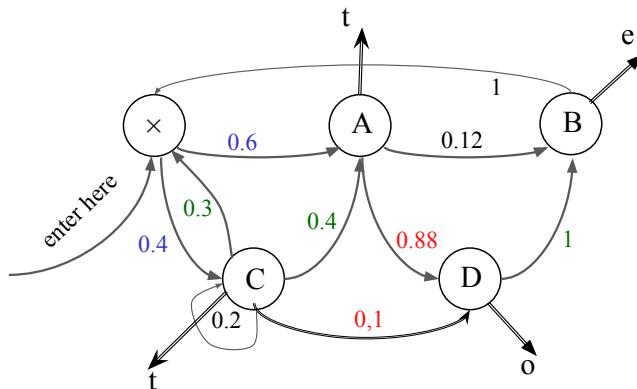
- Continue until output exhausted
  - $|Y| = 3$ : until stage 3 *position/stage:* 2      3
- Add together all the  $\alpha(state, |Y|)$
- That's the P(Y)
- Observation (pleasant):
  - memory usage max:  $2|S|$
  - multiplications max:  $|S|^2|Y|$

Y:



# Trellis (again)

HMM:

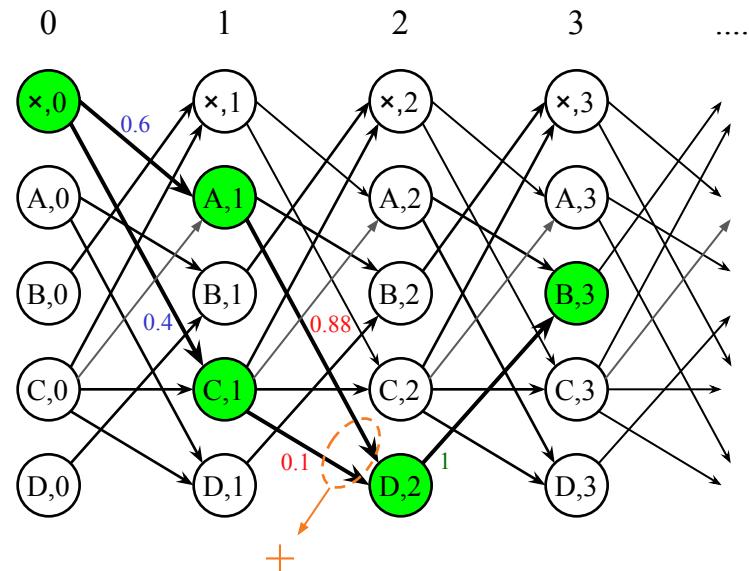


$$p(\text{toe}) = 0.6 \times 0.88 \times 1 + 0.4 \times 1 \times 1 \approx 0.568$$

Trellis:

*time/position:*

"roll-out"



- Trellis state: (HMM state, position)
- each state: holds one number (prob):  $\alpha$
- probability of Y:  $\sum \alpha$  in the last state

Y:

t

o

e

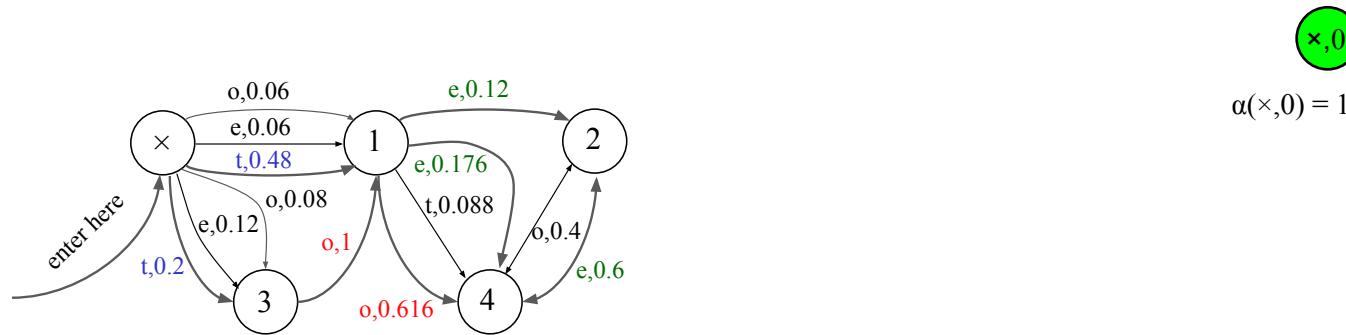
$$\begin{aligned} \alpha(\times, 0) &= 1 & \alpha(A, 1) &= 0.6 & \alpha(D, 2) &= 0.568 & \alpha(B, 3) &= 0.568 \\ & & & & & & \\ \alpha(C, 1) &= 0.4 & & & & & \end{aligned}$$

# HMM: The Two Tasks

- HMM (the general case): five-tuple  $(S, S_0, Y, P_S, P_Y)$ , where:
  - $S = \{s_1, s_2, \dots, s_T\}$  is the set of states,  $S_0$  is the initial state,
  - $Y = \{y_1, y_2, \dots, y_V\}$  is the output alphabet,
  - $P_S(s_j|s_i)$  is the set of prob. distributions of transitions,
  - $P_Y(y_k|s_i, s_j)$  is the set of output (emission) probability distributions
- Given an HMM & an output sequence  $Y = \{y_1, y_2, \dots, y_k\}$ :
  - (Task 1) compute the probability of  $Y$ ;
  - (Task 2) compute the most likely sequence of states which has generated  $Y$ .

# Trellis: The General Case (still, bigrams)

- Start as usual:
  - start state ( $\times$ ), set its  $\alpha(\times, 0)$  to 1. *position/stage:* 0



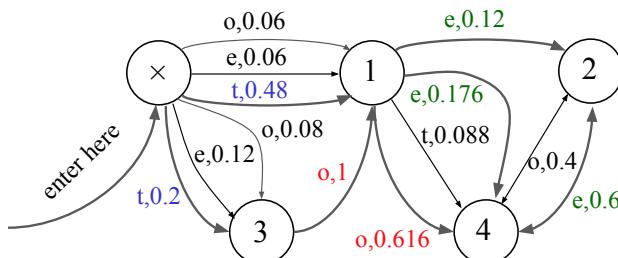
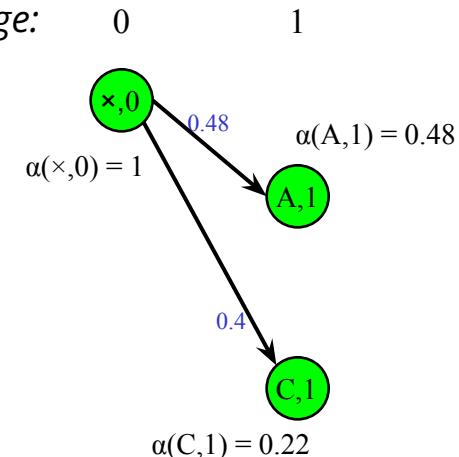
$$\begin{aligned} p(\text{toe}) = & 0.48 \times 0.616 \times 0.6 + \\ & 0.2 \times 1 \times 0.176 + \\ & 0.2 \times 1 \times 0.12 \approx 0.237 \end{aligned}$$

Y:

# General Trellis: The Next Step

We are in stage  $i$ :

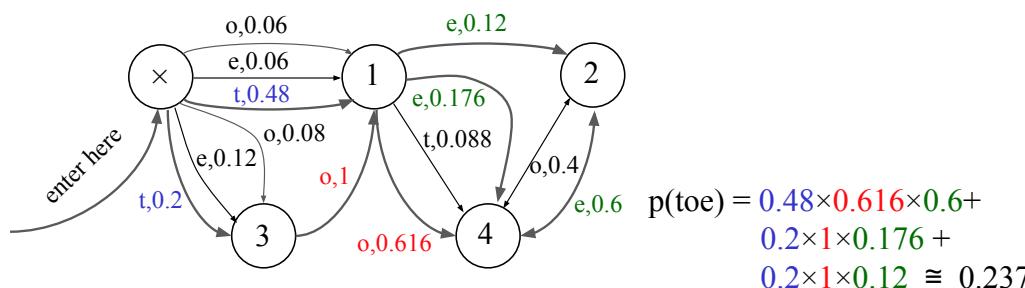
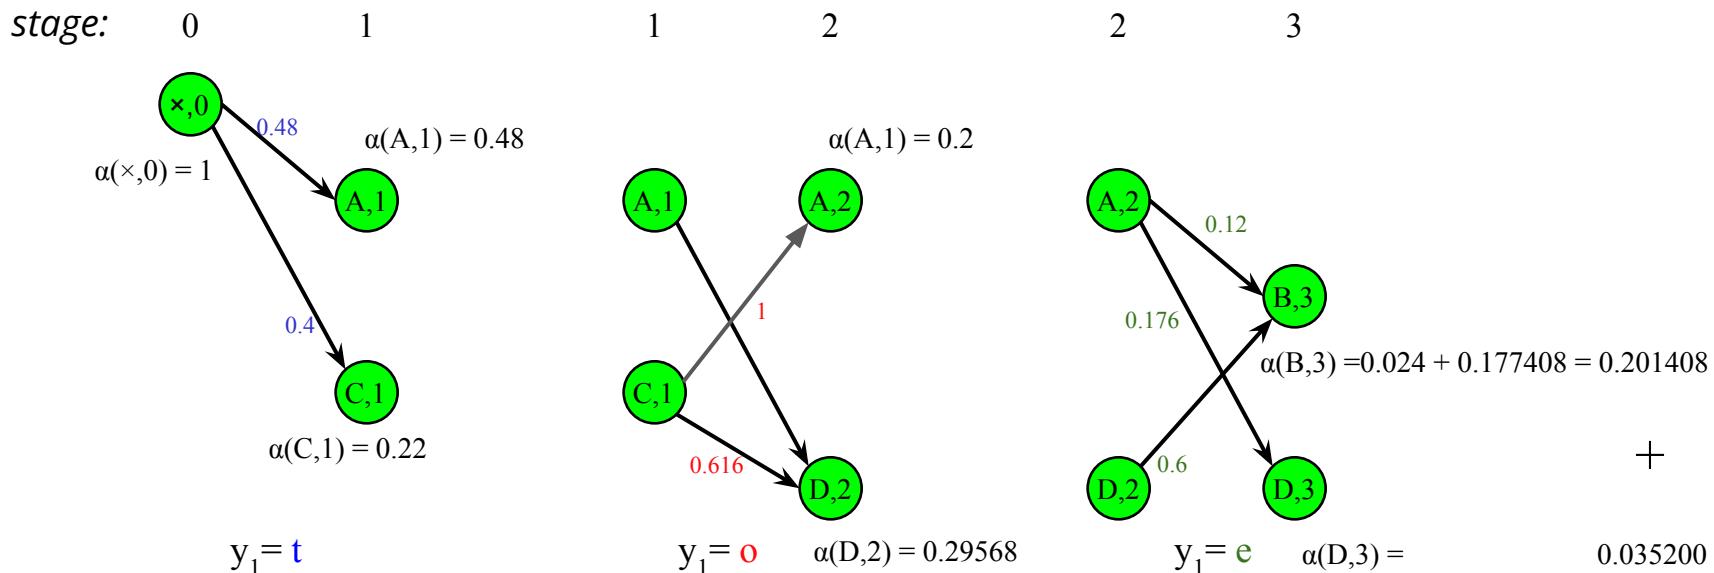
- Generate the next stage  $i+1$  as before (except now arcs generate output, thus use only those arcs marked by the output symbol  $y_{i+1}$ )
- For each generated *state*, compute  $\alpha(state, i+1) = \sum_{\text{inc. arcs}} P_Y(y_{i+1} | state, prev.state) \times \alpha(prev.state, i)$



Y:  $t$

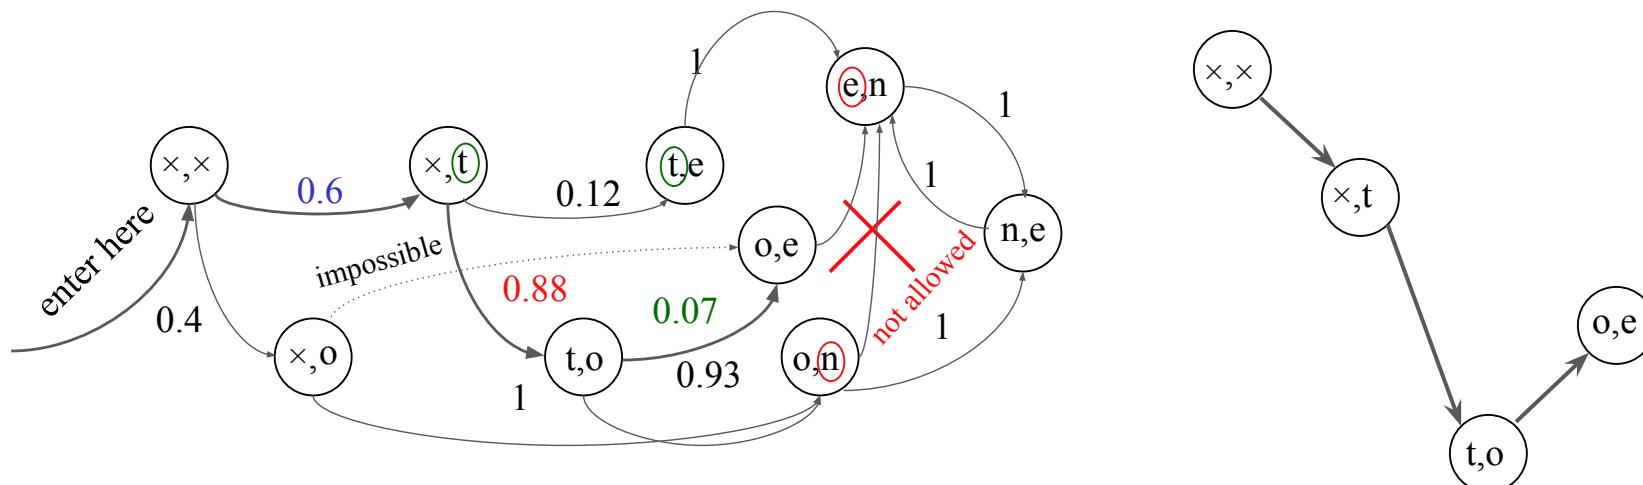
And forget about the previous state ...

# Trellis: The Complete Example



# The Case of Trigrams

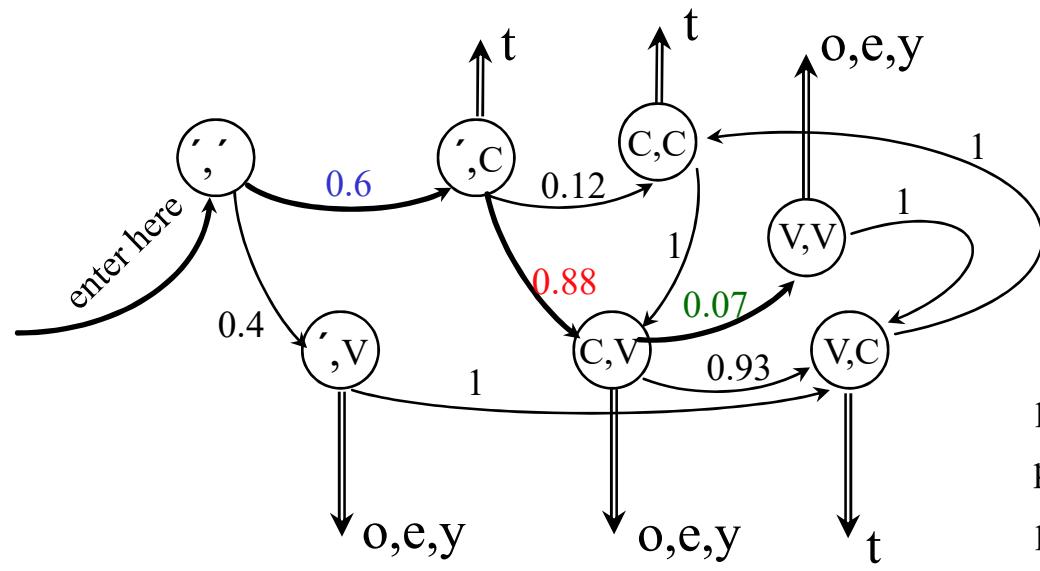
- Like before, but:
  - states correspond to bigrams
  - output function always emits the second output symbol of the pair (state) to which the arc goes:



- Multiple paths not possible → trellis not really needed

# Trigrams with Classes

- More interesting:
  - n-gram class LM:  $p(w_i|w_{i-2}, w_{i-1}) = p(w_i|c_i) p(c_i|c_{i-2}, c_{i-1})$
  - states are pairs of classes  $(c_{i-1}, c_i)$ , and emit “words”:



(letters in our example)

$$\begin{aligned} p(t|C) &= 1 && \text{usual,} \\ p(o|V) &= .3 && \text{non-} \\ p(e|V) &= .6 && \text{overlapping} \\ p(y|V) &= .1 && \text{classes} \end{aligned}$$

$$p(\text{toe}) = .6 \cdot 1 \cdot .88 \cdot .3 \cdot .07 \cdot .6 \cong .00665$$

$$p(\text{teo}) = .6 \cdot 1 \cdot .88 \cdot .6 \cdot .07 \cdot .3 \cong .00665$$

$$p(\text{toy}) = .6 \cdot 1 \cdot .88 \cdot .3 \cdot .07 \cdot .1 \cong .00111$$

$$p(\text{tty}) = .6 \cdot 1 \cdot .12 \cdot 1 \cdot 1 \cdot .1 \cong .0072$$

# Class Trigrams: the Trellis

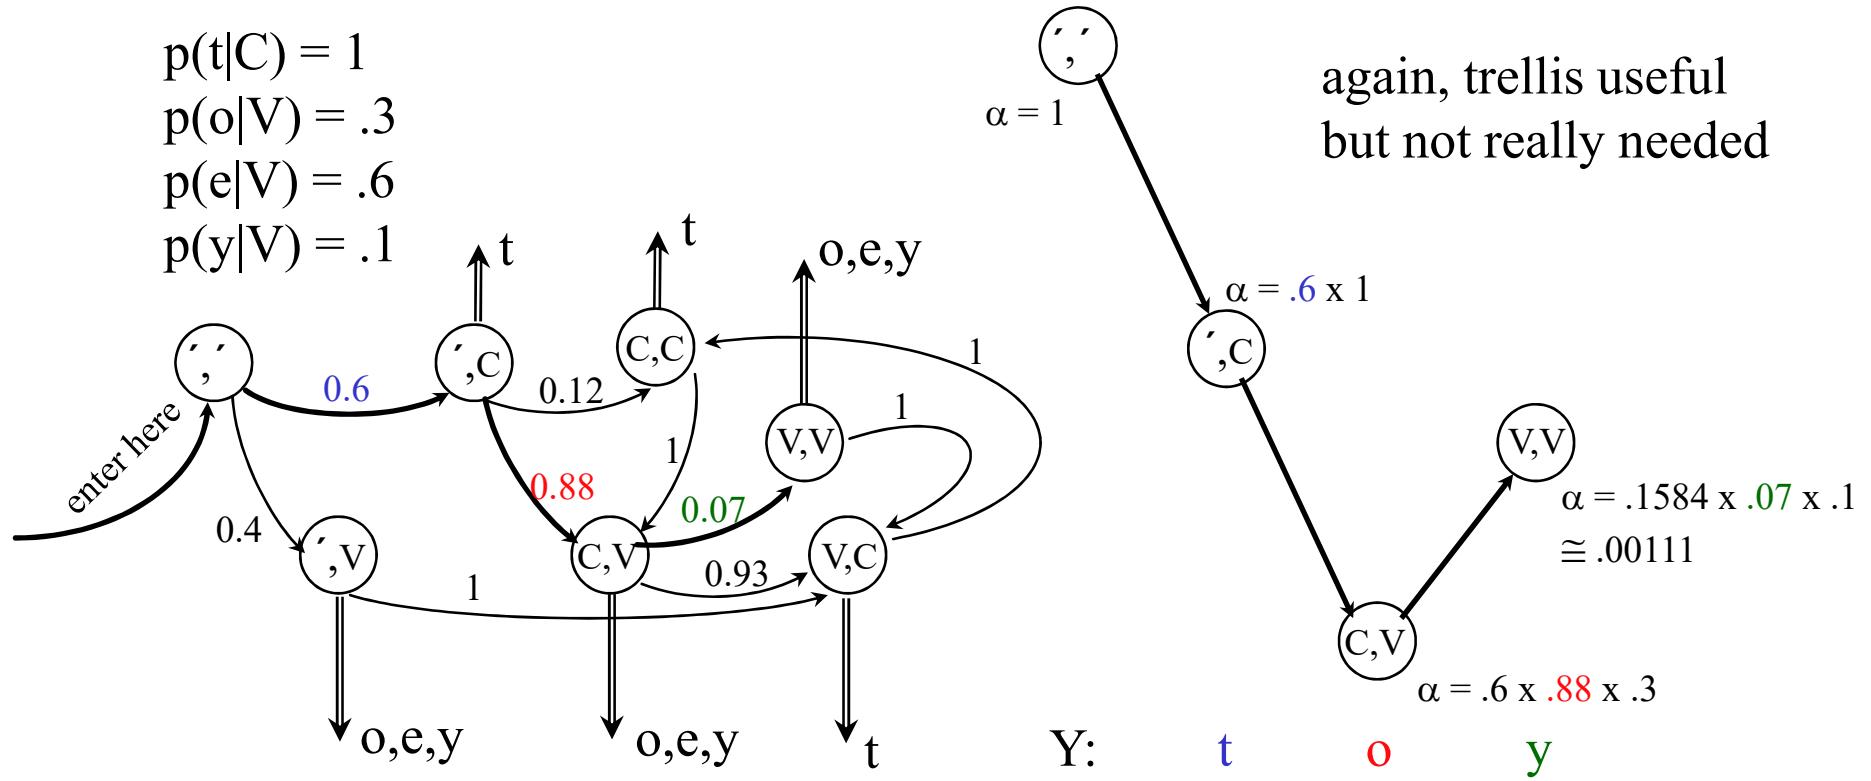
- Trellis generation ( $Y = \text{"toy"}$ ):

$$p(t|C) = 1$$

$$p(o|V) = .3$$

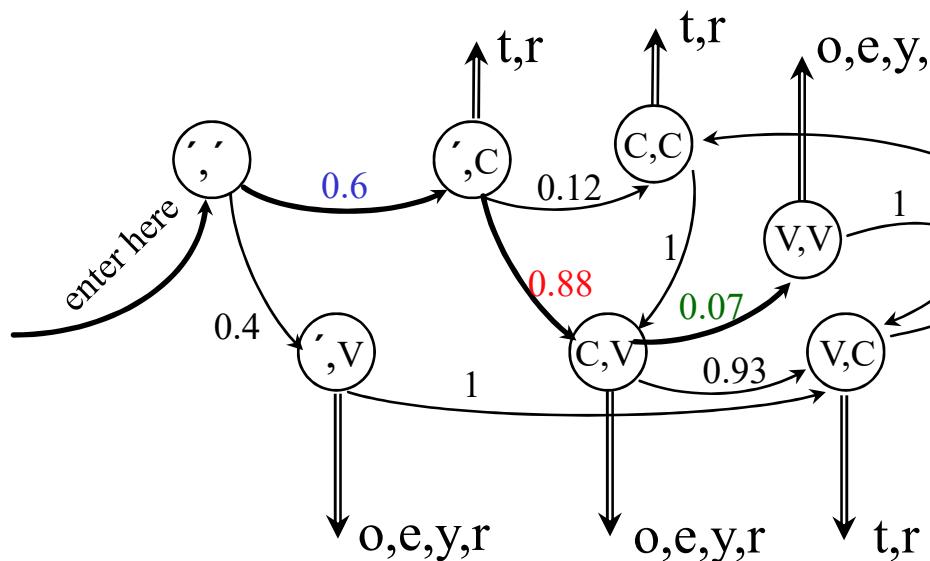
$$p(e|V) = .6$$

$$p(y|V) = .1$$



# Overlapping Classes

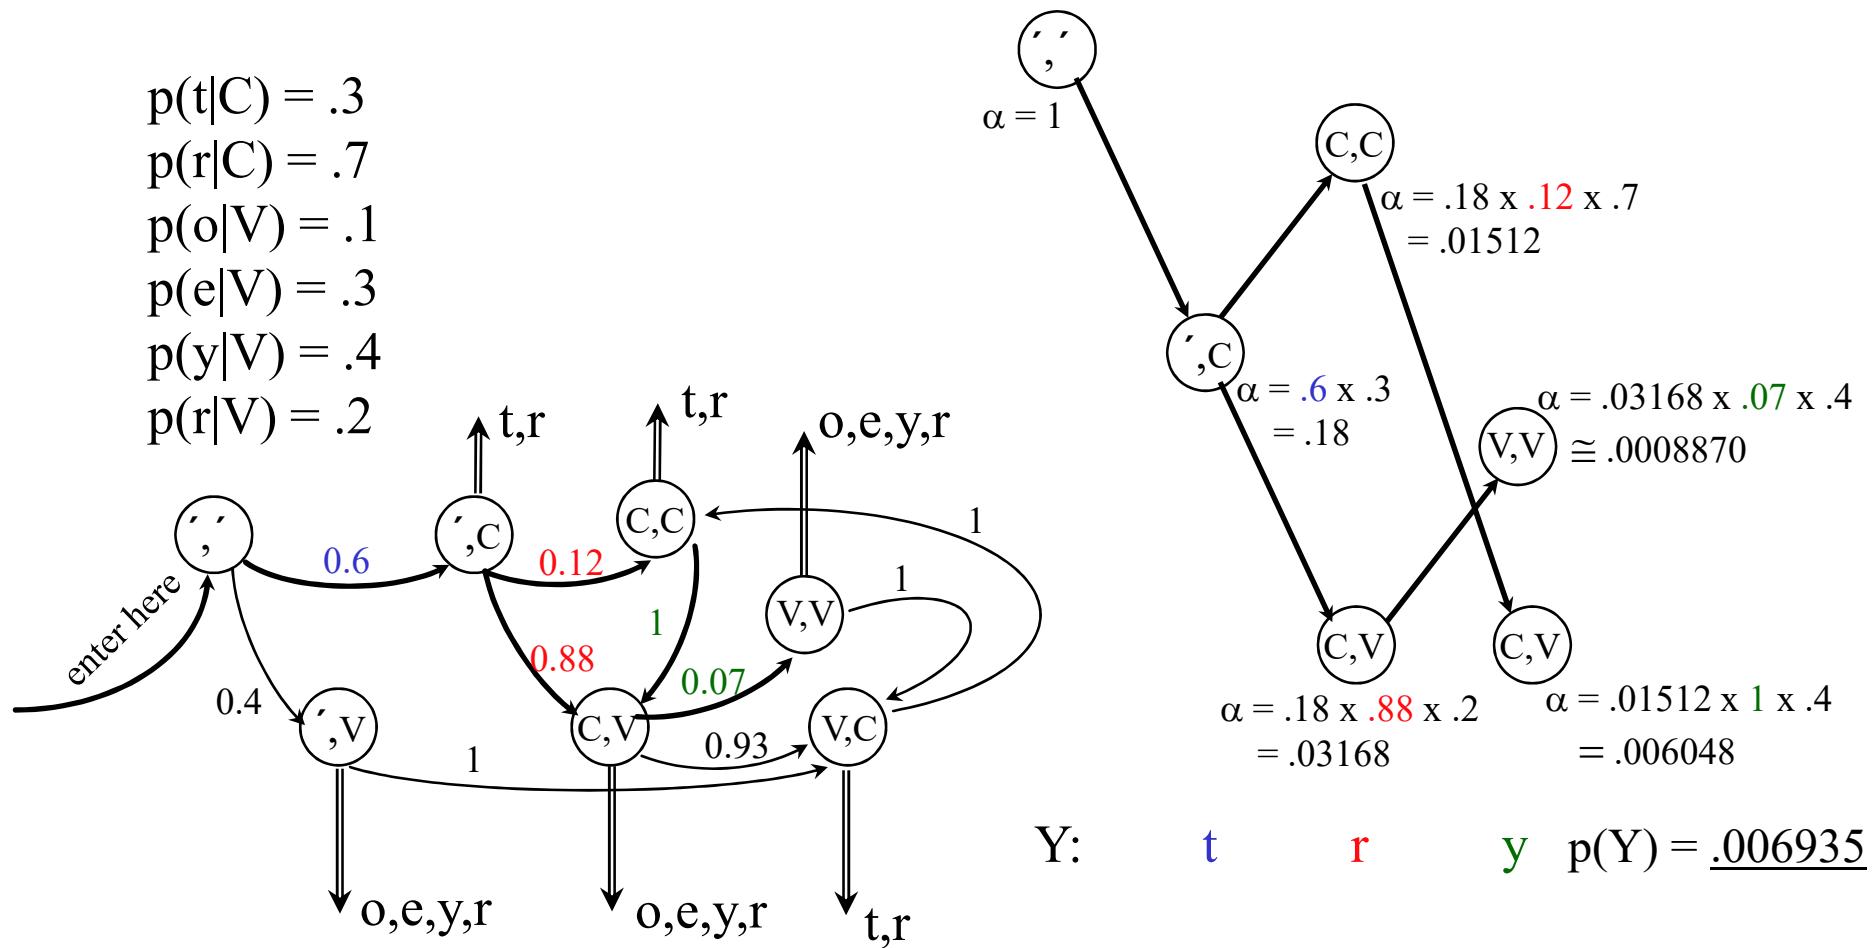
- Imagine that classes may overlap
  - e.g. ‘r’ is sometimes vowel sometimes consonant, belongs to V as well as C:



$$\begin{aligned}
 p(t|C) &= .3 \\
 p(r|C) &= .7 \\
 p(o|V) &= .1 \\
 p(e|V) &= .3 \\
 p(y|V) &= .4 \\
 p(r|V) &= .2
 \end{aligned}$$

$$p(\text{try}) = ?$$

# Overlapping Classes: Trellis Example



# Trellis: Remarks

- So far, we went left to right (computing  $\alpha$ )
- Same result: going right to left (computing  $\beta$ )
  - supposed we know where to start (finite data)
- In fact, we might start in the middle going left and right
- Important for parameter estimation
  - (Forward-Backward Algorithm alias Baum-Welch)
- Implementation issues:
  - scaling/normalizing probabilities, to avoid too small numbers & addition problems with many transitions

# The Viterbi Algorithm

- Solving the task of finding the most likely sequence of states which generated the observed data
- i.e., finding

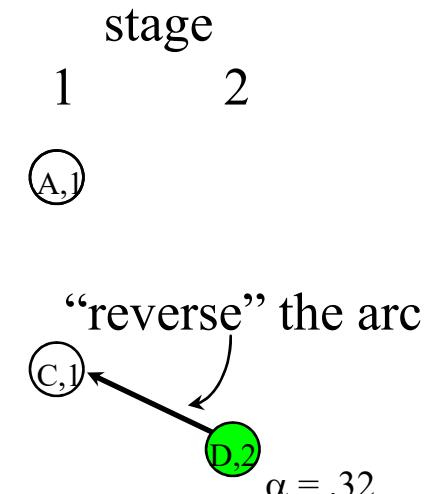
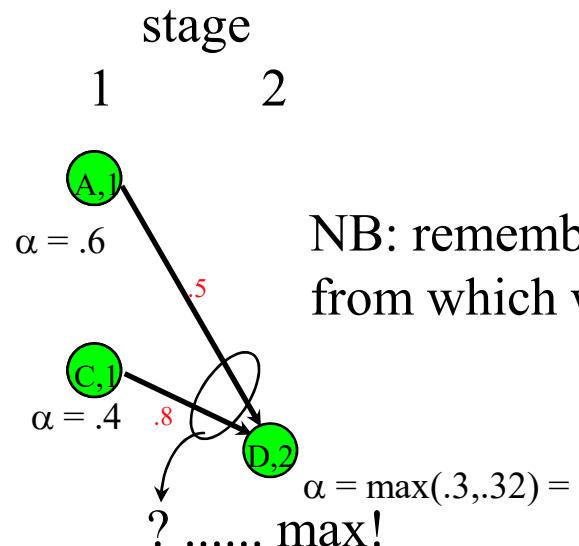
$$S_{\text{best}} = \operatorname{argmax}_S P(S|Y)$$

which is equal to ( $Y$  is constant and thus  $P(Y)$  is fixed):

$$\begin{aligned} S_{\text{best}} &= \operatorname{argmax}_S P(S, Y) = \\ &= \operatorname{argmax}_S P(s_0, s_1, s_2, \dots, s_k, y_1, y_2, \dots, y_k) = \\ &= \operatorname{argmax}_S \prod_{i=1..k} p(y_i | s_i, s_{i-1}) p(s_i | s_{i-1}) \end{aligned}$$

# The Crucial Observation

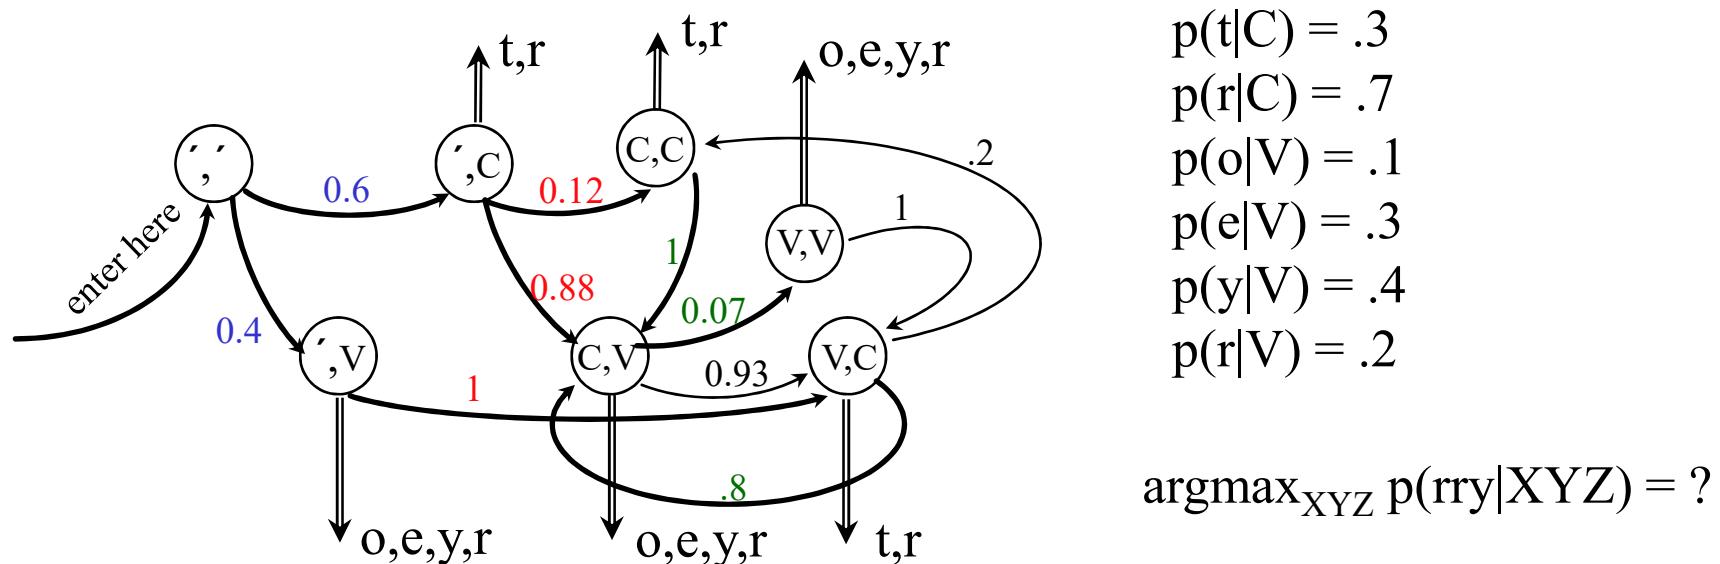
- Imagine the trellis build as before (but do not compute the  $\alpha$ s yet; assume they are o.k.); stage  $i$ :



this is certainly the “backwards” maximum to (D,2)... but  
it cannot change even whenever we go forward (M. Property: Limited History)

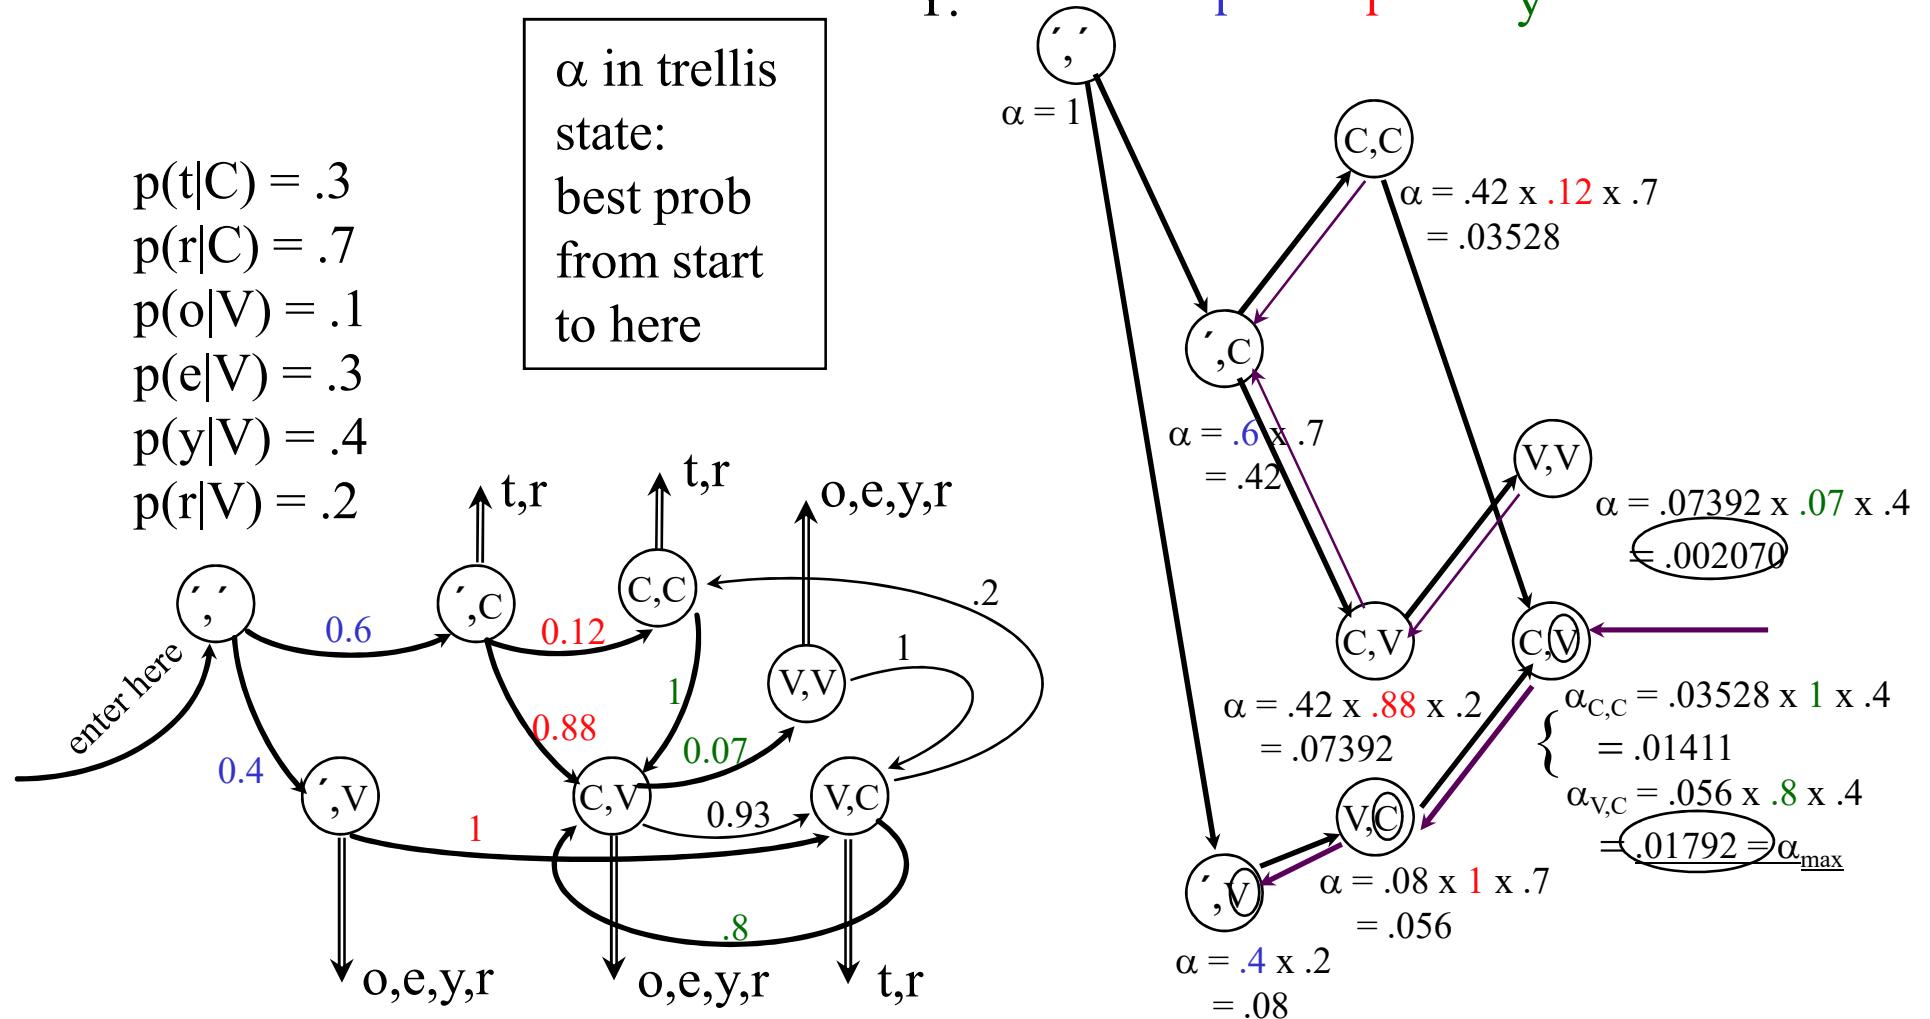
# Viterbi Example

- ‘r’ classification (C or V?, sequence?):



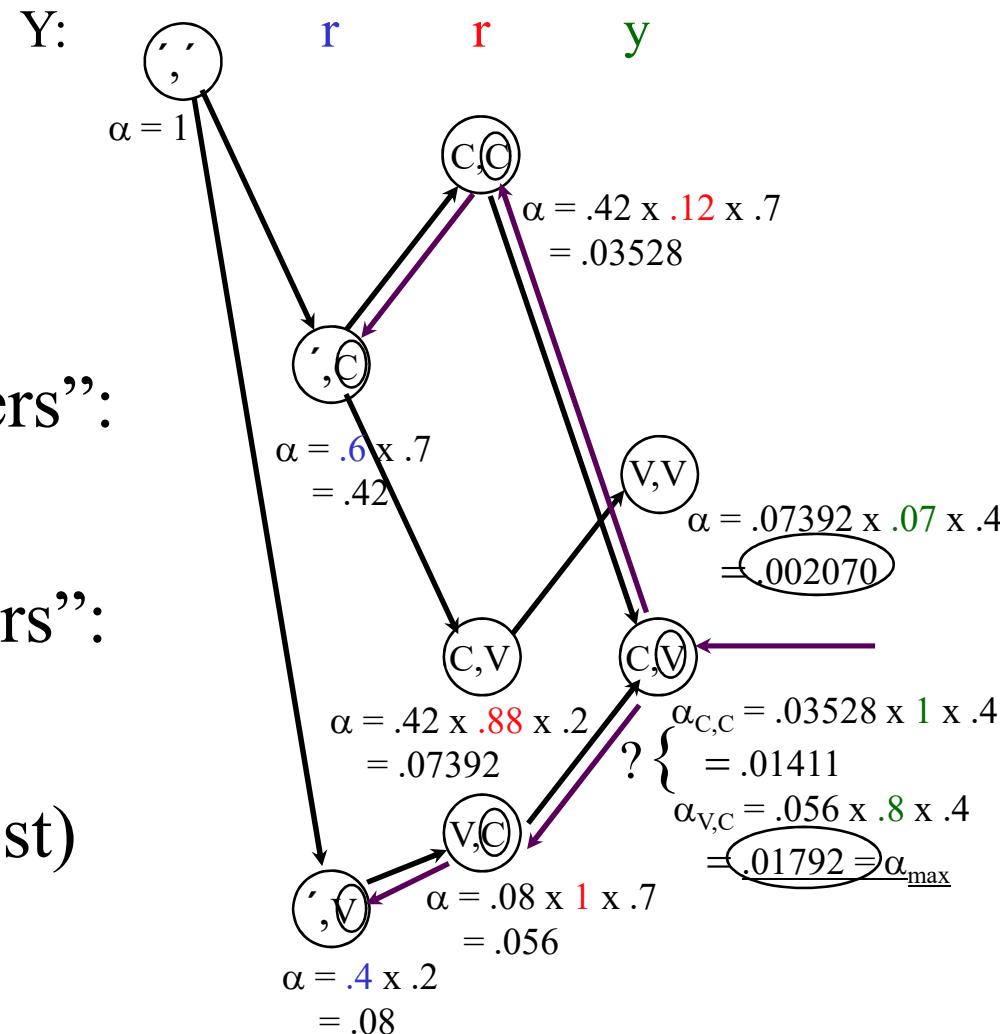
Possible state seq.:  $(\cdot, v)(v, c)(c, v)[VCV]$ ,  $(\cdot, c)(c, c)(c, v)[CCV]$ ,  $(\cdot, c)(c, v)(v, v) [CVV]$

# Viterbi Computation



# n-best State Sequences

- Keep track of n best “back pointers”:
- Ex.: n= 2:  
Two “winners”: VCV (best)  
CCV (2<sup>nd</sup> best)

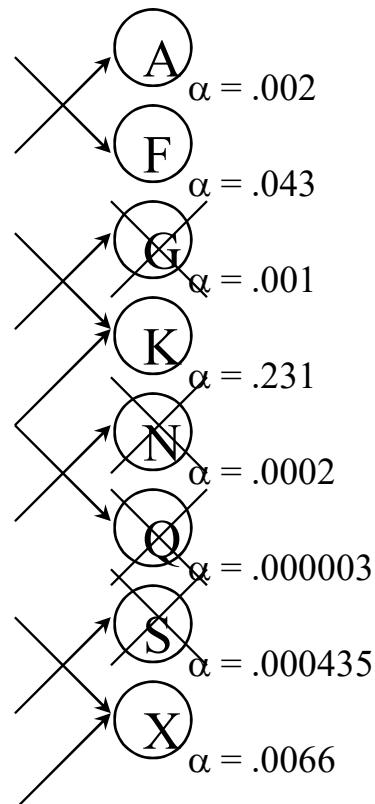


# Tracking Back the n-best paths

- Backtracking-style algorithm:
  - Start at the end, in the best of the n states ( $s_{best}$ )
  - Put the other  $n-1$  best nodes/back pointer pairs on stack, except those leading from  $s_{best}$  to the same best-back state.
- Follow the back “beam” towards the start of the data, spitting out nodes on the way (backwards of course) using always only the best back pointer.
- At every beam split, push the diverging node/back pointer pairs onto the stack (node/beam width is sufficient!).
- When you reach the start of data, close the path, and pop the top-most node/back pointer(width) pair from the stack.
- Repeat until the stack is empty; expand the result tree if necessary.

# Pruning

- Sometimes, too many trellis states in a stage:



criteria: (a)  $\alpha < \text{threshold}$   
(b)  $\sum \pi < \text{threshold}$   
(c) # of states  $> \text{threshold}$   
(get rid of smallest  $\alpha$ )

# HMM Parameter Estimation: the Baum-Welch Algorithm

# HMM: The Tasks

- HMM (the general case):
  - five-tuple  $(S, S_0, Y, P_S, P_Y)$ , where:
    - **$S = \{s_1, s_2, \dots, s_T\}$  is the set of states,  $S_0$  is the initial state,**
    - **$Y = \{y_1, y_2, \dots, y_V\}$  is the output alphabet,**
    - **$P_S(s_j|s_i)$  is the set of prob. distributions of transitions,**
    - **$P_Y(y_k|s_i, s_j)$  is the set of output (emission) probability distributions.**
- Given an HMM & an output sequence  $Y = \{y_1, y_2, \dots, y_k\}$ :
  - ✓(Task 1) compute the probability of  $Y$ ;
  - ✓(Task 2) compute the most likely sequence of states which has generated  $Y$ .
- (Task 3) Estimating the parameters (transition/output distributions)

# A Variant of EM

- Idea (~ EM, for another variant see LM smoothing):
  - Start with (possibly random) estimates of  $P_S$  and  $P_Y$ .
  - Compute (fractional) “counts” of state transitions/emissions taken, from  $P_S$  and  $P_Y$ , given data  $Y$ .
  - Adjust the estimates of  $P_S$  and  $P_Y$  from these “counts” (using the MLE, i.e. relative frequency as the estimate).
- Remarks:
  - many more parameters than the simple four-way smoothing
  - no proofs here; see Jelinek, Chapter 9

# Setting

- HMM (without  $P_S, P_Y$ ) ( $S, S_0, Y$ ), and data  $T = \{y^i \in Y\}_{i=1..|T|}$ 
  - will use  $T \sim |T|$
  - HMM structure is given:  $(S, S_0)$
  - $P_S$ : Typically, one wants to allow “fully connected” graph
    - (i.e. no transitions forbidden  $\sim$  no transitions set to hard 0)
    - why?  $\rightarrow$  we better leave it on the learning phase, based on the data!
    - sometimes possible to remove some transitions ahead of time
  - $P_Y$ : should be restricted (if not, we will not get anywhere!)
    - restricted  $\sim$  hard 0 probabilities of  $p(y|s,s')$
    - “Dictionary”: states  $\leftrightarrow$  words, “m:n” mapping on  $S \times Y$  (in general)

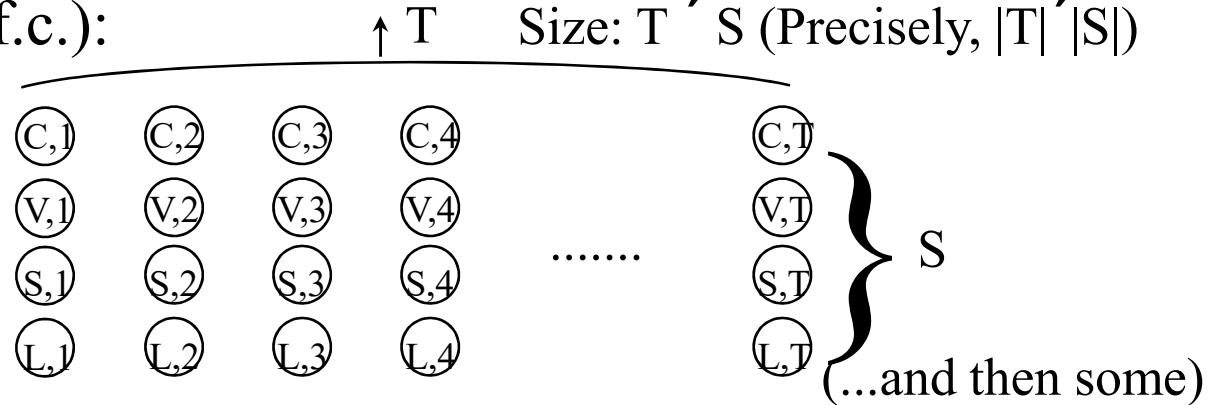
# Initialization

- For computing the initial expected “counts”
- Important part
  - EM guaranteed to find a *local* maximum only (albeit a good one in most cases)
- $P_Y$  initialization more important
  - fortunately, often easy to determine
    - **together with dictionary  $\leftrightarrow$  vocabulary mapping, get counts, then MLE**
- $P_S$  initialization less important
  - e.g. uniform distribution for each  $p(.|s)$

# Data Structures

- Will need storage for:
    - The predetermined structure of the HMM  
(unless fully connected → need not to keep it!)
    - The parameters to be estimated ( $P_S$ ,  $P_Y$ )
    - The expected counts (same size as  $P_S$ ,  $P_Y$ )
    - The training data  $T = \{y^i \in Y\}_{i=1..T}$
    - The trellis (if f.c.):  
↑ T      Size:  $T^S$  (Precisely,  $|T|^{|S|}$ )

Each trellis state:  
**two** [float] numbers  
(forward/backward)



# The Algorithm Part I

1. Initialize  $P_S, P_Y$
2. Compute “forward” probabilities:
  - follow the procedure for trellis (summing), compute  $\alpha(s,i)$
  - use the current values of  $P_S, P_Y$  ( $p(s'|s)$ ,  $p(y|s,s')$ ):
$$\alpha(s',i) = \sum_{s \rightarrow s'} \alpha(s,i-1) \times p(s'|s) \times p(y_i|s,s')$$
  - NB: do not throw away the previous stage!
3. Compute “backward” probabilities
  - start at all nodes of the last stage, proceed backwards,  $\beta(s,i)$
  - i.e., probability of the “tail” of data from stage  $i$  to the end of data
$$\beta(s',i) = \sum_{s \leftarrow s'} \beta(s,i+1) \times p(s|s') \times p(y_{i+1}|s',s)$$
  - also, keep the  $\beta(s,i)$  at all trellis states

# The Algorithm Part II

## 4. Collect counts:

- for each output/transition pair compute

$$c(y, s, s') = \sum_{i=0..k-1, y=y_{i+1}} \alpha(s, i) p(s'|s) \underbrace{p(y_{i+1}|s, s')}_{\text{this transition prob}} \beta(s', i+1)$$

one pass through data,  
only stop at (output)  $y$

prefix prob.

tail prob

output prob

$$c(s, s') = \sum_{y \in Y} c(y, s, s') \text{ (assuming all observed } y_i \text{ in } Y)$$

$$c(s) = \sum_{s' \in S} c(s, s')$$

5. Reestimate:  $p'(s'|s) = c(s, s')/c(s)$   $p'(y|s, s') = c(y, s, s')/c(s, s')$
6. Repeat 2-5 until desired convergence limit is reached.

# Baum-Welch: Tips & Tricks

- Normalization badly needed
  - long training data → extremely small probabilities
- Normalize  $\alpha, \beta$  using the same norm. factor:

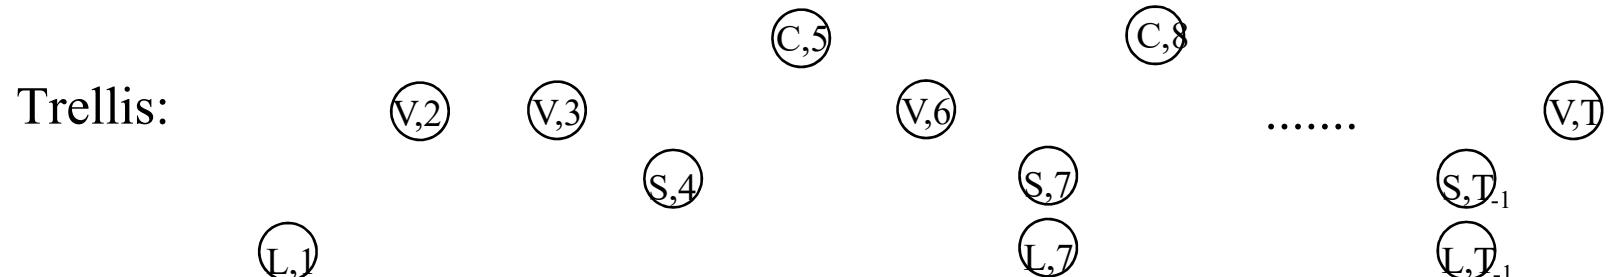
$$N(i) = \sum_{s \in S} \alpha(s,i)$$

as follows:

- compute  $\alpha(s,i)$  as usual (Step 2 of the algorithm), computing the sum  $N(i)$  at the given stage  $i$  as you go.
- at the end of each stage, recompute all  $\alpha$ s (for each state  $s$ ):
  - $\alpha^*(s,i) = \alpha(s,i) / N(i)$
- use the same  $N(i)$  for  $\beta$ s at the end of each backward (Step 3) stage:
  - $\beta^*(s,i) = \beta(s,i) / N(i)$

# Example

- Task: pronunciation of “the”
- Solution: build HMM, fully connected, 4 states:
  - S - short article, L - long article, C,V - starting w/consonant, vowel
  - thus, only “the” is ambiguous (a, an, the - not members of C,V)
- Output from states only ( $p(w|s,s') = p(w|s')$ )
- Data Y: an egg and a piece of the big .... the end



# Example: Initialization

- Output probabilities:  
 $p_{\text{init}}(w|c) = c(c,w) / c(c)$ ; where  $c(S,\text{the}) = c(L,\text{the}) = c(\text{the})/2$   
(other than that, everything is deterministic)
- Transition probabilities:
  - $p_{\text{init}}(c'|c) = 1/4$  (uniform)
- Don't forget:
  - about the space needed
  - initialize  $\alpha(X,0) = 1$  ( $X$  : the never-occurring front buffer st.)
  - initialize  $\beta(s,T) = 1$  for all  $s$  (except for  $s = X$ )

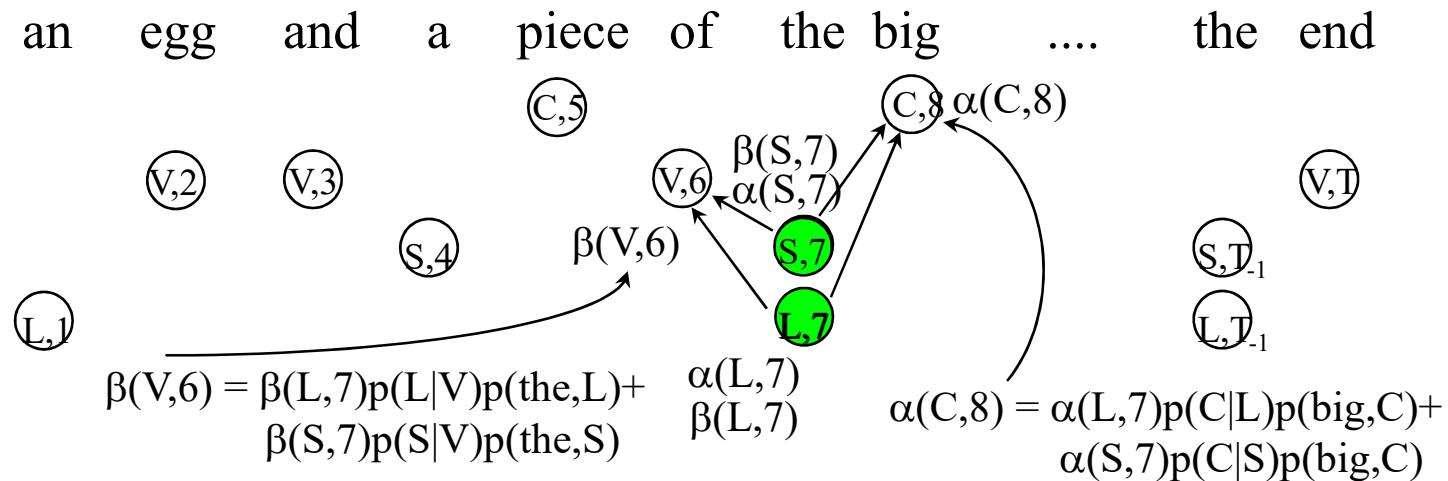
# Fill in alpha, beta

- Left to right, alpha:

$$\alpha(s', i) = \sum_{s \rightarrow s'} \alpha(s, i-1) \times p(s'|s) \times p(w_i|s')$$

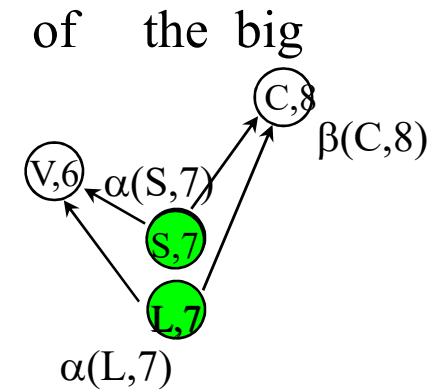
output from states

- Remember normalization ( $N(i)$ ).
- Similarly, beta (on the way back from the end).



# Counts & Reestimation

- One pass through data
- At each position  $i$ , go through all pairs  $(s_i, s_{i+1})$
- Increment appropriate counters by frac. counts (Step 4):
  - $\text{inc}(y_{i+1}, s_i, s_{i+1}) = a(s_i, i) p(s_{i+1}|s_i) p(y_{i+1}|s_{i+1}) b(s_{i+1}, i+1)$
  - $c(y, s_i, s_{i+1}) += \text{inc}$  (for  $y$  at pos  $i+1$ )
  - $c(s_i, s_{i+1}) += \text{inc}$  (always)
  - $c(s_i) += \text{inc}$  (always)
- $\text{inc}(\text{big}, L, C) = \alpha(L, 7)p(C|L)p(\text{big}, C)\beta(C, 8)$
- $\text{inc}(\text{big}, S, C) = \alpha(S, 7)p(C|S)p(\text{big}, C)\beta(C, 8)$
- Reestimate  $p(s'|s)$ ,  $p(y|s)$ 
  - and hope for increase in  $p(C|S)$  and  $p(V|L)$ ...!!



# HMM: Final Remarks

- Parameter “tying”:
  - keep certain parameters same (~ just one “counter” for all of them)
  - any combination in principle possible
  - ex.: smoothing (just one set of lambdas)
- Real Numbers Output
  - $Y$  of infinite size ( $R$ ,  $R^n$ ):
    - **parametric (typically: few) distribution needed (e.g., “Gaussian”)**
- “Empty” transitions: do not generate output
  - **~ vertical arcs in trellis; do not use in “counting”**